

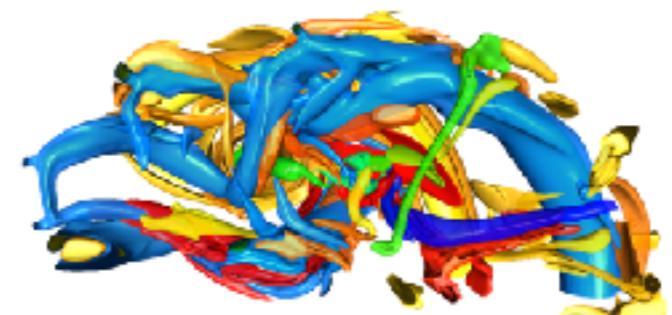
Scientific Visualization for Computational Science

Christoph Garth

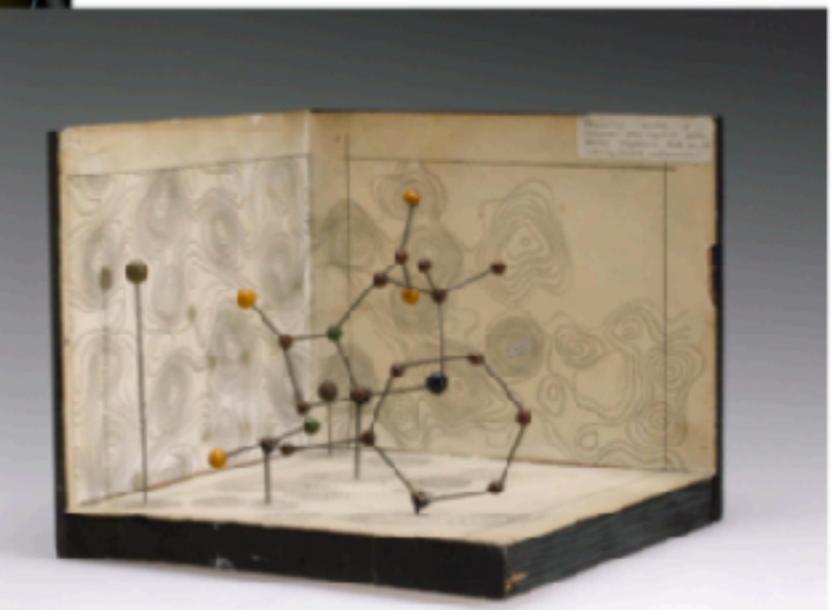
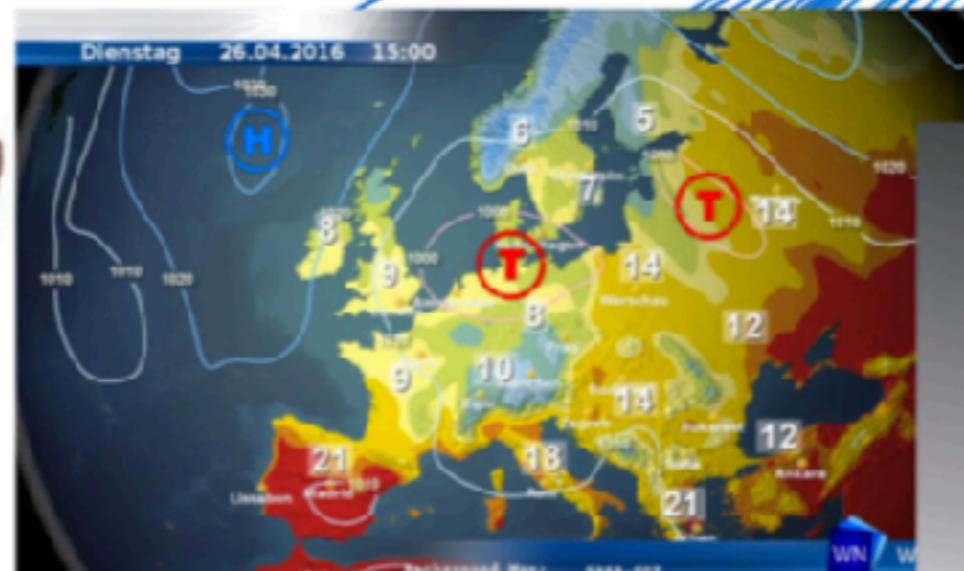
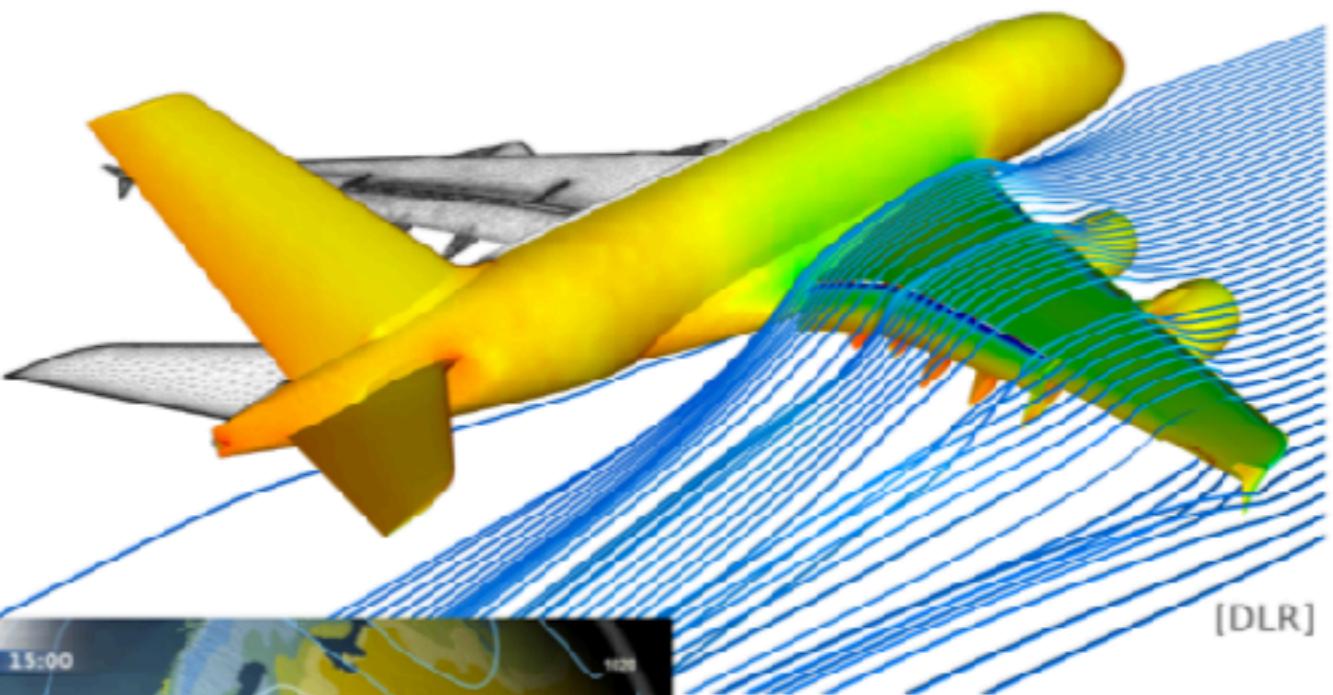
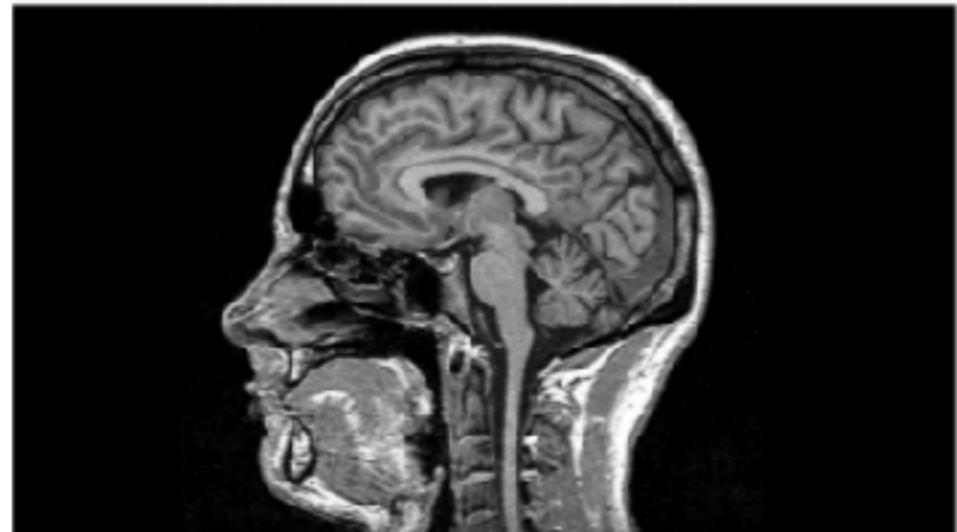
University of Kaiserslautern



International Summer Workshop on Visual Computing
Universidad de los Andes, Bogotá, June 2018



Scientific Visualization Examples



[science museum]

Visualization



Computer-based visualization systems provide visual representations of datasets designed to help people carry out tasks more efficiently.

Visualization is suitable when there is a need to augment human capabilities rather than replace people with computational decision-making methods.

[Munzner 2014]

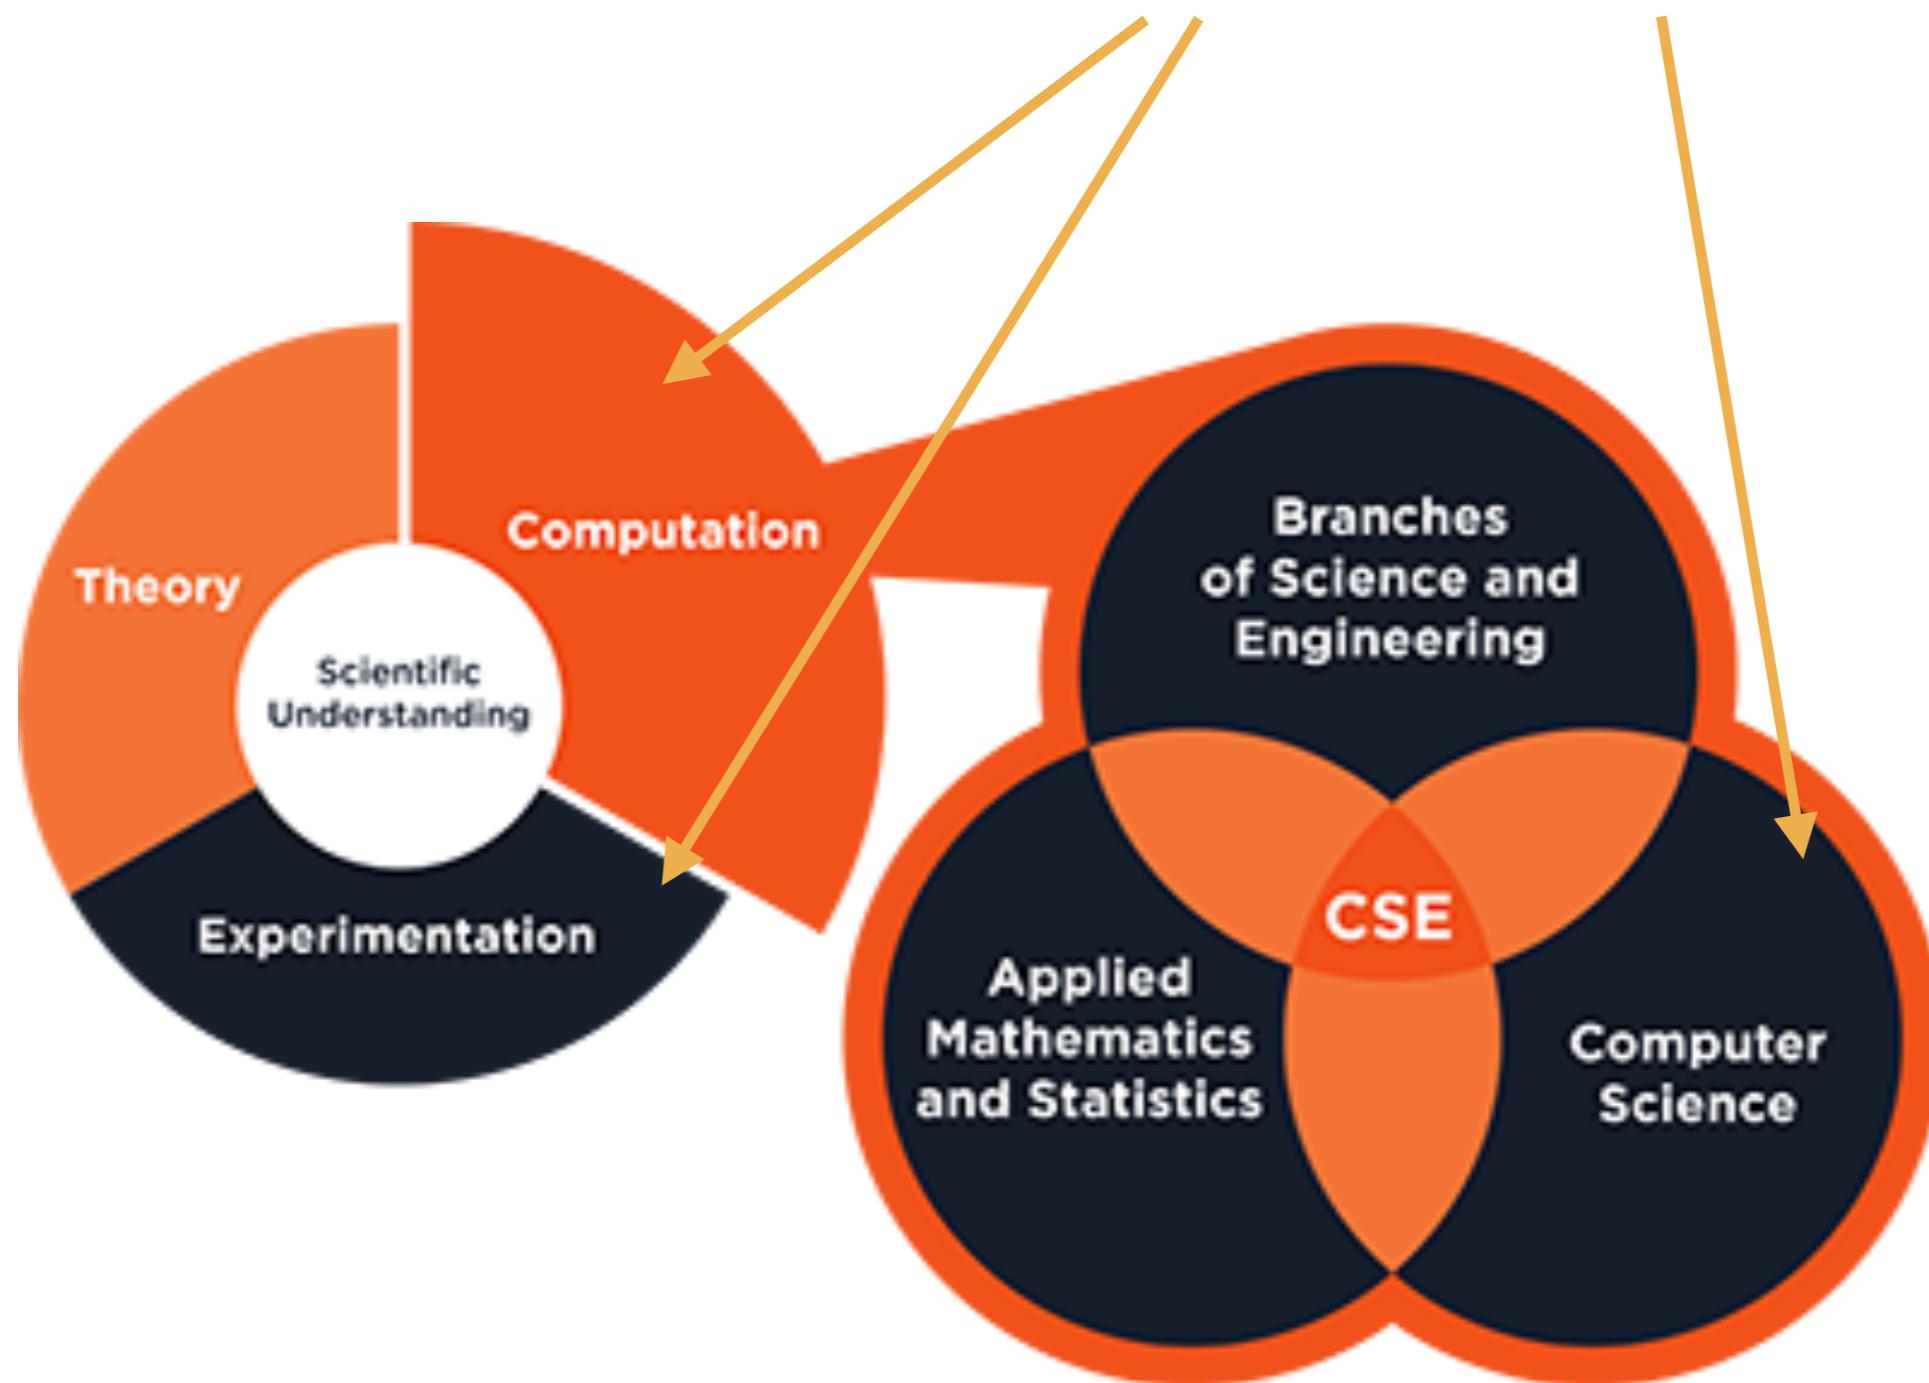
Key tasks:

- exploration and discovery
- hypothesis formation and confirmation
- communication of results, sharing of insight

Visualization in Computational Science



Visualization & Data Analysis



Large-Scale Computational Science



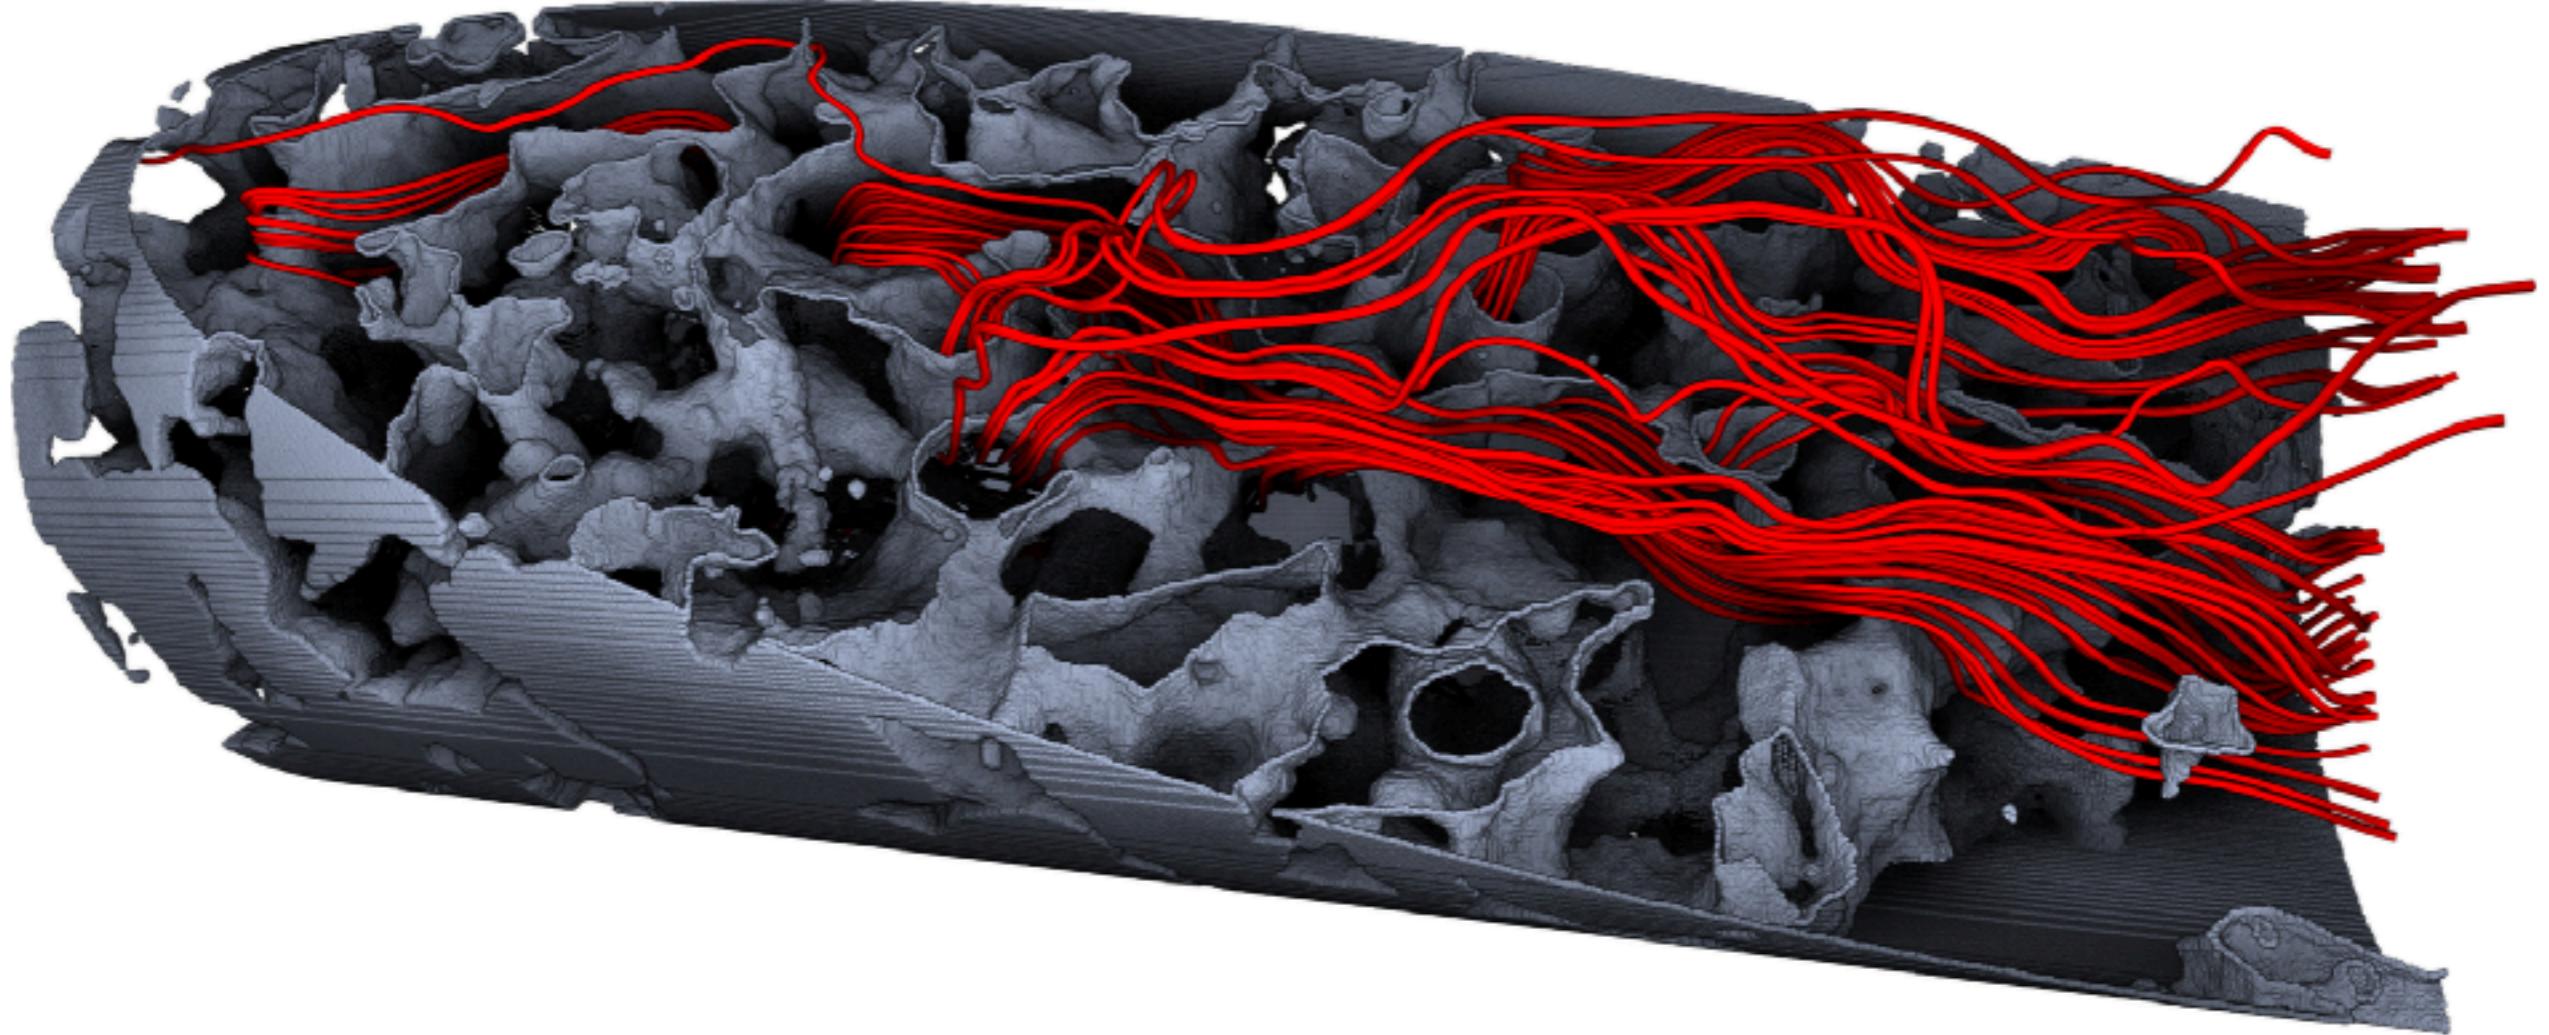
Modern workflows leverage massively parallel machines.

- clusters
(10s - 1000s of proc.)
- supercomputers
(1,000,000s of proc.)



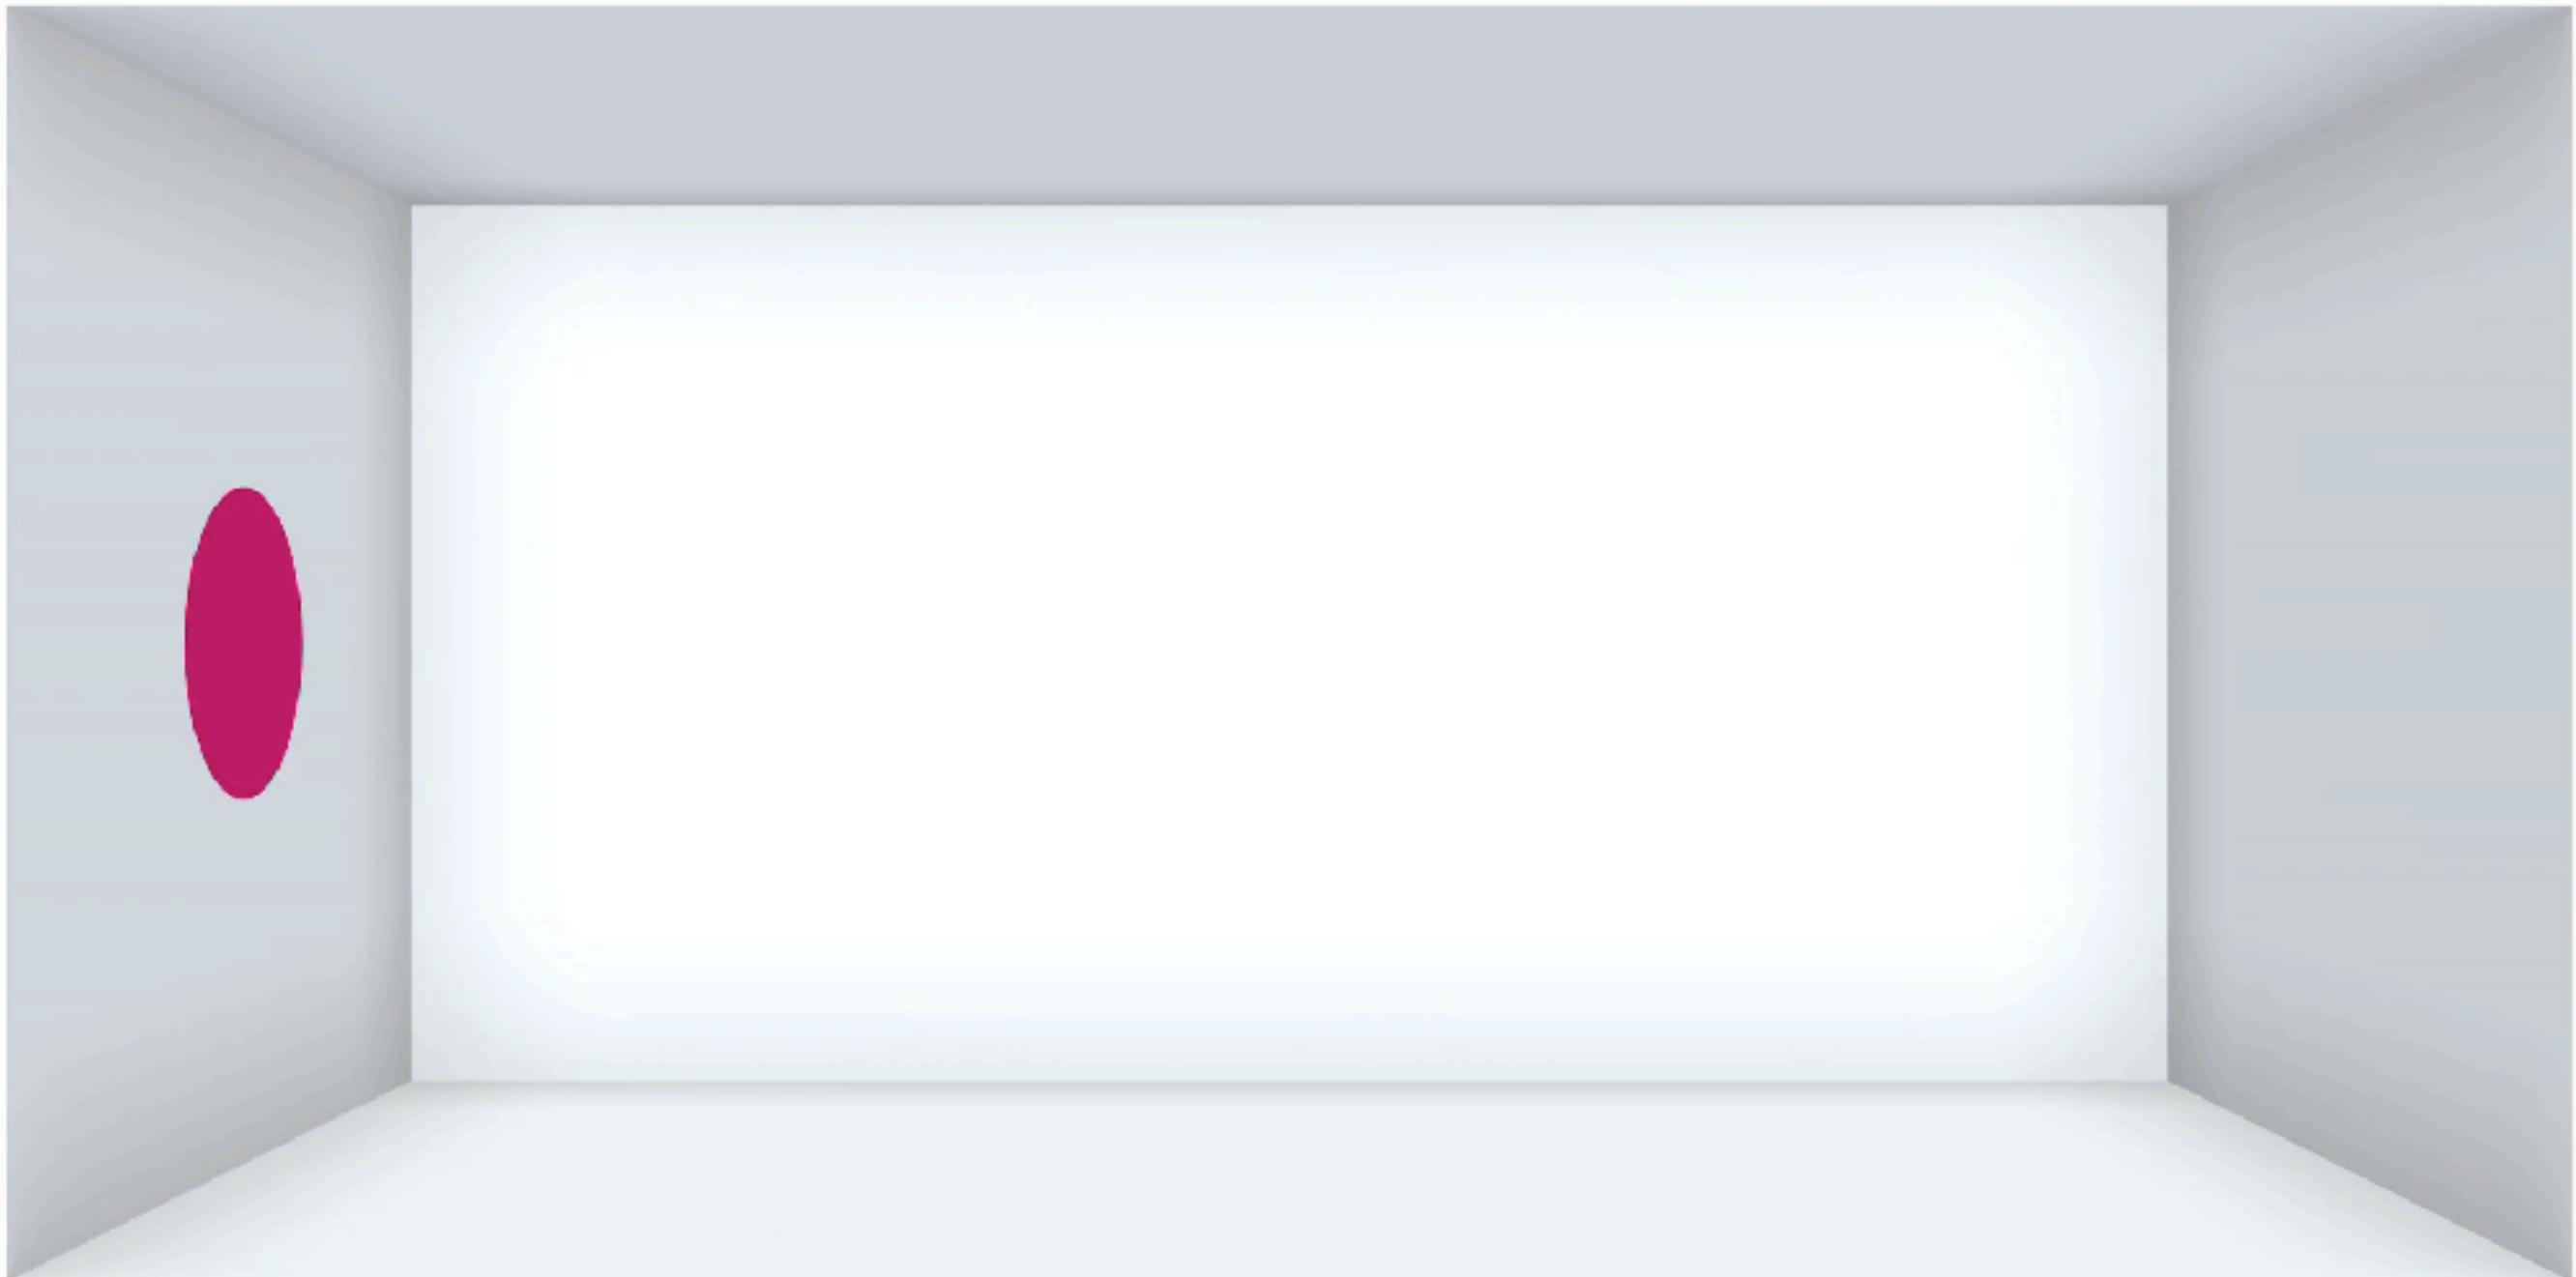
This allows to consider problems of almost arbitrary complexity, scale and resolution (Exascale era).

Scientific Visualization Examples



Flow of water through Karst limestone (data: TACC): 500 GB

Scientific Visualization Examples



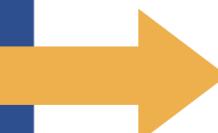
Jet flow (DNS, data: Scientific Computation Group @ TUK): 1TB

Modern Simulation Data



Modern Simulation Data

Complexity



State-of-the-art results from Computational Science

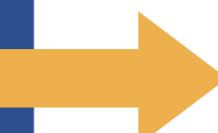
Structural richness of the data; simple feature definitions are not sufficient; quantification of relevance & simplification needed.

**High-Dimensional
Multivariate**



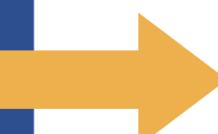
Traditional scalar, vector, tensor visualization techniques do not apply.

**Uncertainty
Quantification**



“Uncertainty” modelled along with data – lost in visualization unless explicitly cared for.

Large-Scale



Efficient algorithms are a must, or, alternative techniques are needed (\rightarrow *in situ* visualization).



A quick tour through some modern paradigms:

- Topology-based and Feature-based Techniques
- Parallel Visualization Algorithms
- In Situ Visualization

Necessarily incomplete.

Feature-Based Visualization

Features in Scientific Visualization



Feature (working definition): a subset of data that is interesting or relevant in the sense that it facilitates one or more of the three tasks of visualization.

For example, a region of a dataset where some functional relation holds

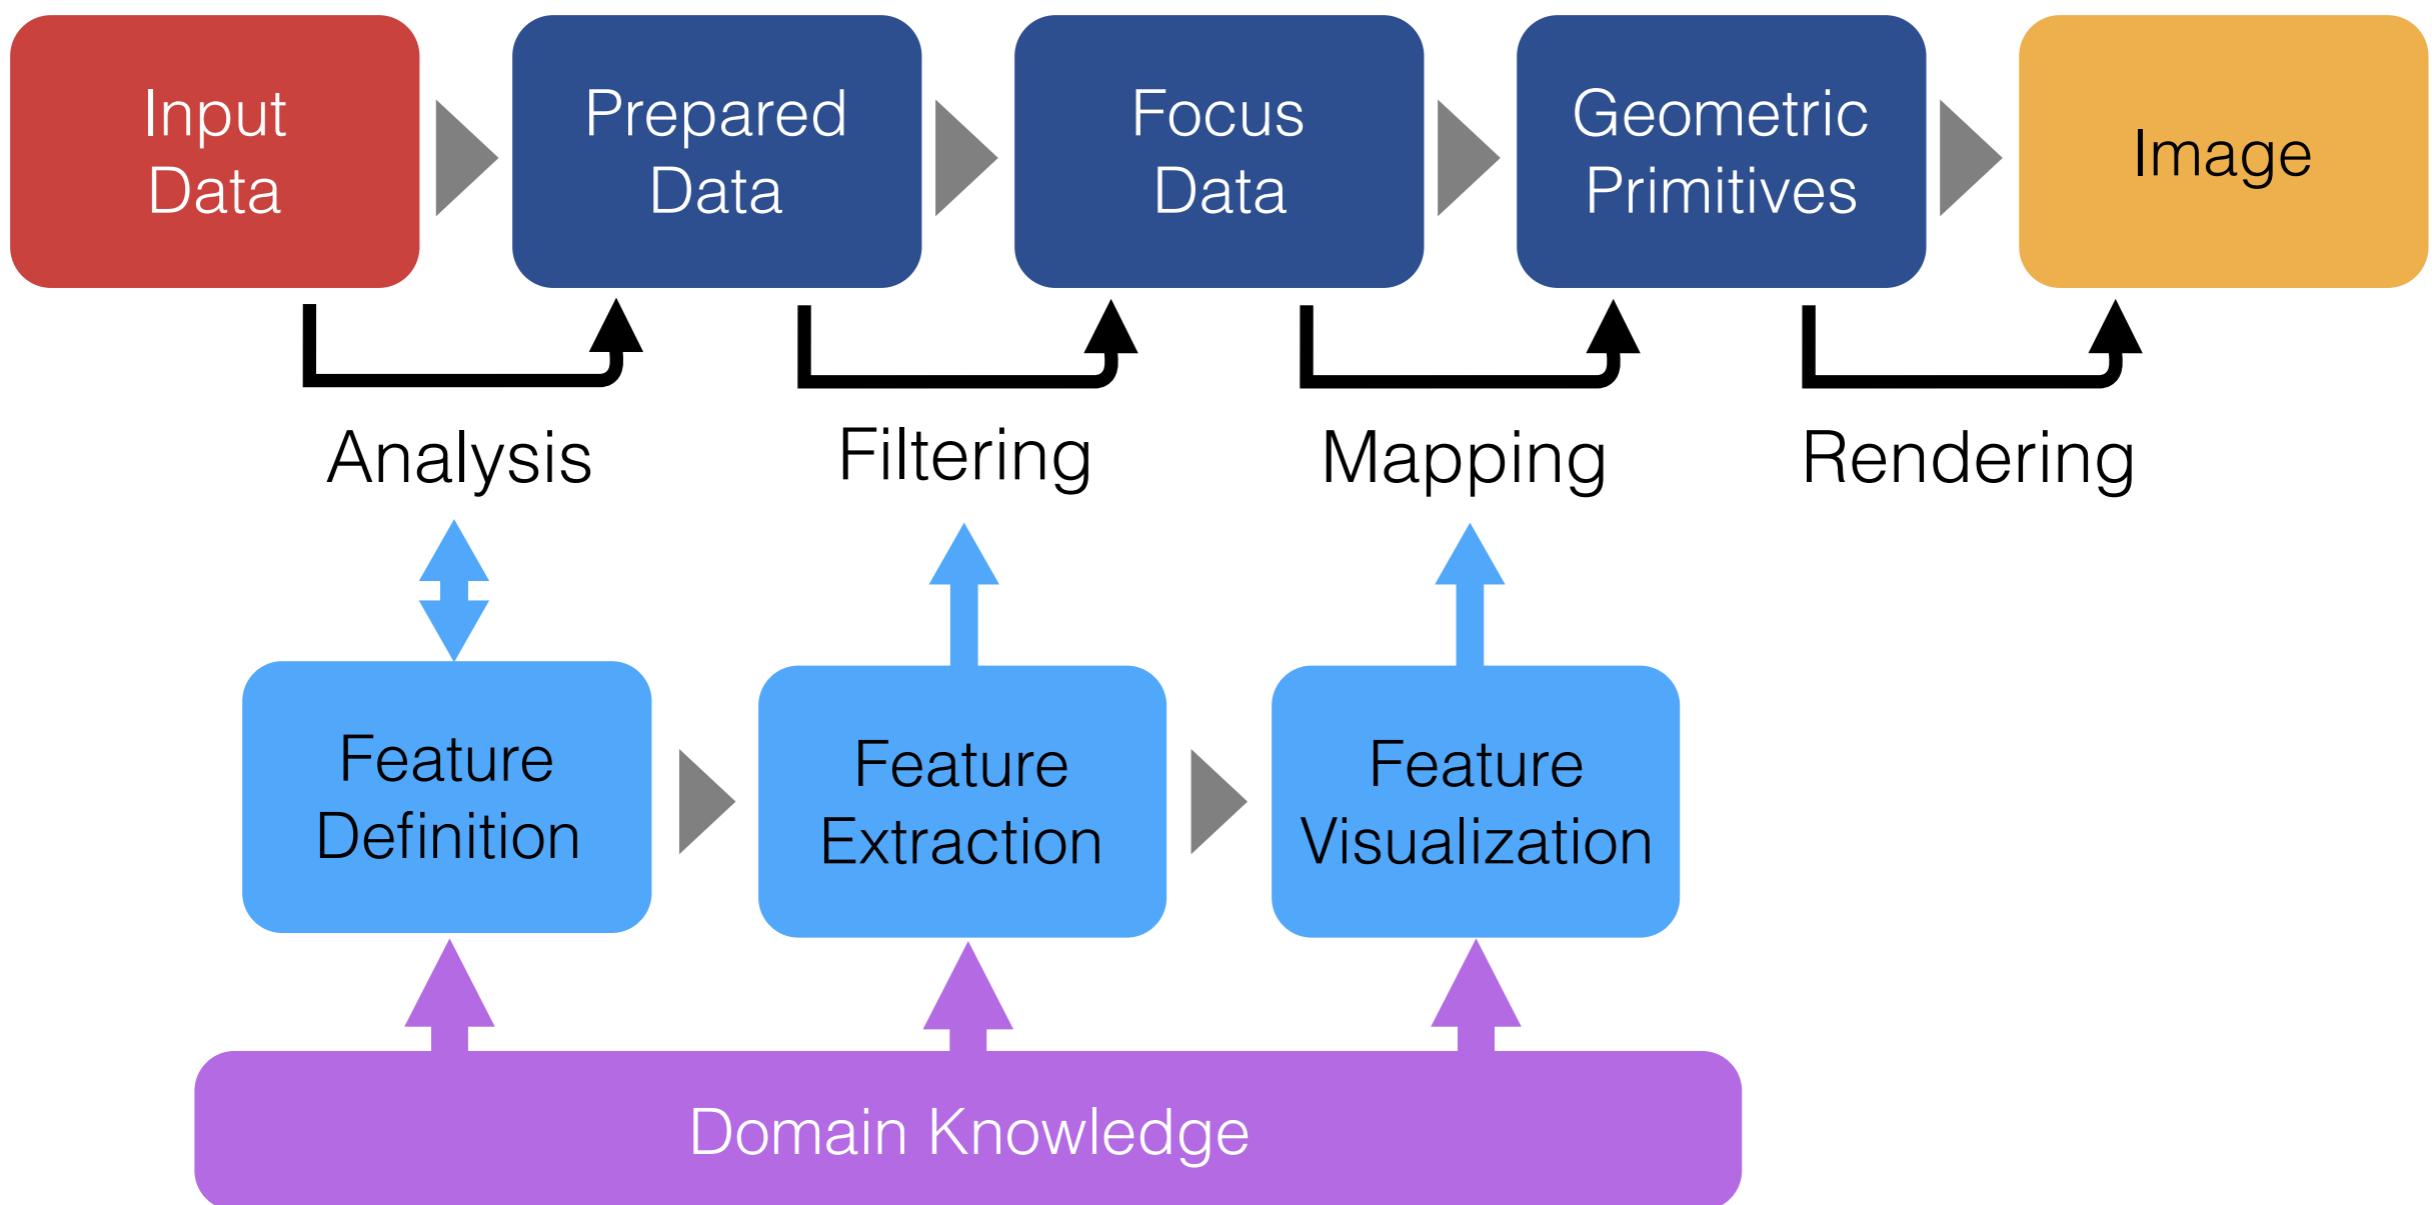
Defining relevant features specific problems requires describing an (efficiently) computable feature criterion.

Feature-based visualization is a particular mental framework to reason about visualization.

Features in the Visualization Pipeline



A model for the visualization process: the visualization pipeline.



Feature-Based Visualization



A first example: wind turbine arrangement



(image: leanfiberopticdesign.com)

Feature-Based Visualization



A first example: wind turbine arrangement



first turbine

downstream →

second turbine

data from R. Linn, E. Koo, Los Alamos National Laboratory

Vortices from first turbine hit second turbine rotor blades and induce heightened stresses, leading to maintenance overhead and failure.

S. Shafii, H. Obermaier, R. Linn, E. Koo, M. Hlawitschka, CG, B. Hamann, K. I. Joy: *Visualization and Analysis of Vortex-Turbine Intersections in Wind Farms*. IEEE TVCG 19(9):1579-1591, 2013

Feature-Based Visualization



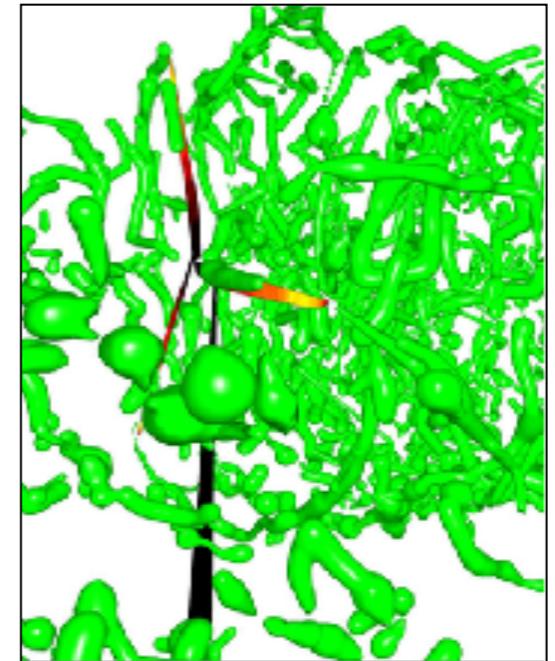
Feature definition:

Vortical regions transported along wind vector field from upstream to downstream turbine.



Feature extraction:

- extraction of vortical regions
- computation of particle paths along wind field
- filter for particles connecting upstream and downstream turbines, contained in vortical regions



S. Shafii, H. Obermaier, R. Linn, E. Koo, M. Hlawitschka, CG, B. Hamann, K. I. Joy: *Visualization and Analysis of Vortex-Turbine Intersections in Wind Farms*. IEEE TVCG 19(9):1579-1591, 2013

Feature-Based Visualization



Feature visualization: improved mapping

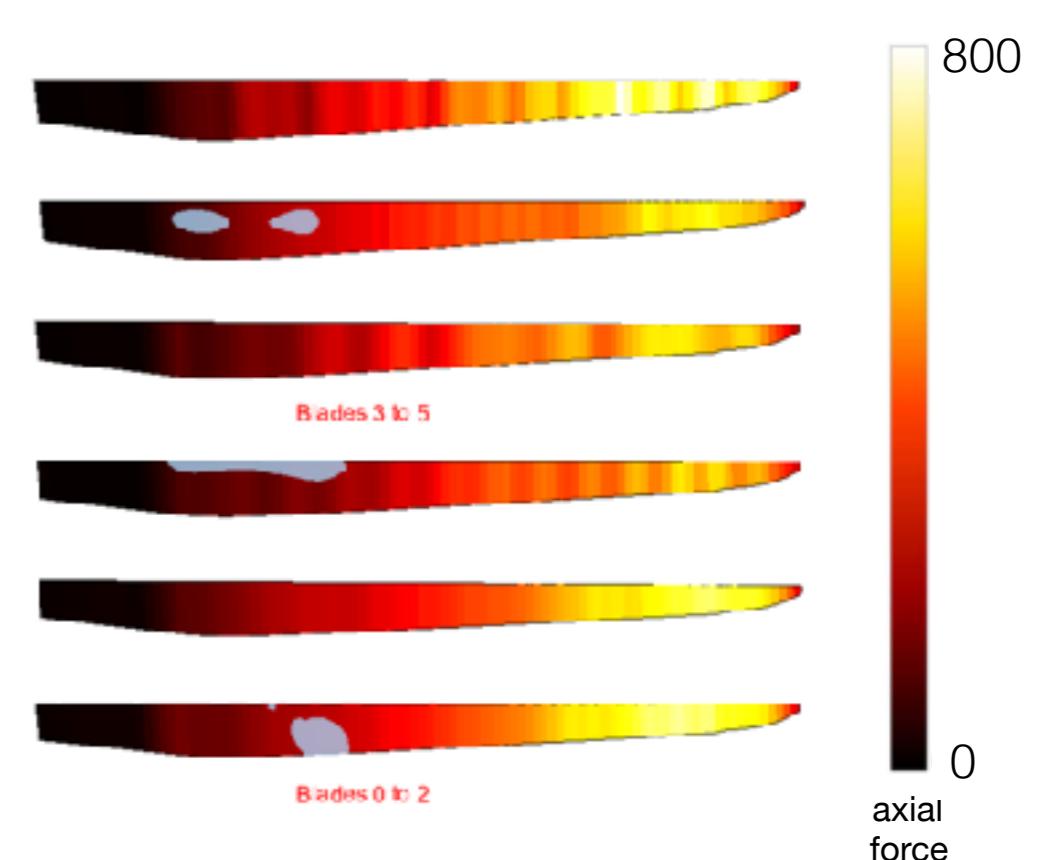
Intersections on (moving) blades are hard to see.
Where do vortices hit? How much force do they induce?



naïve (3D)



intersections (3D)



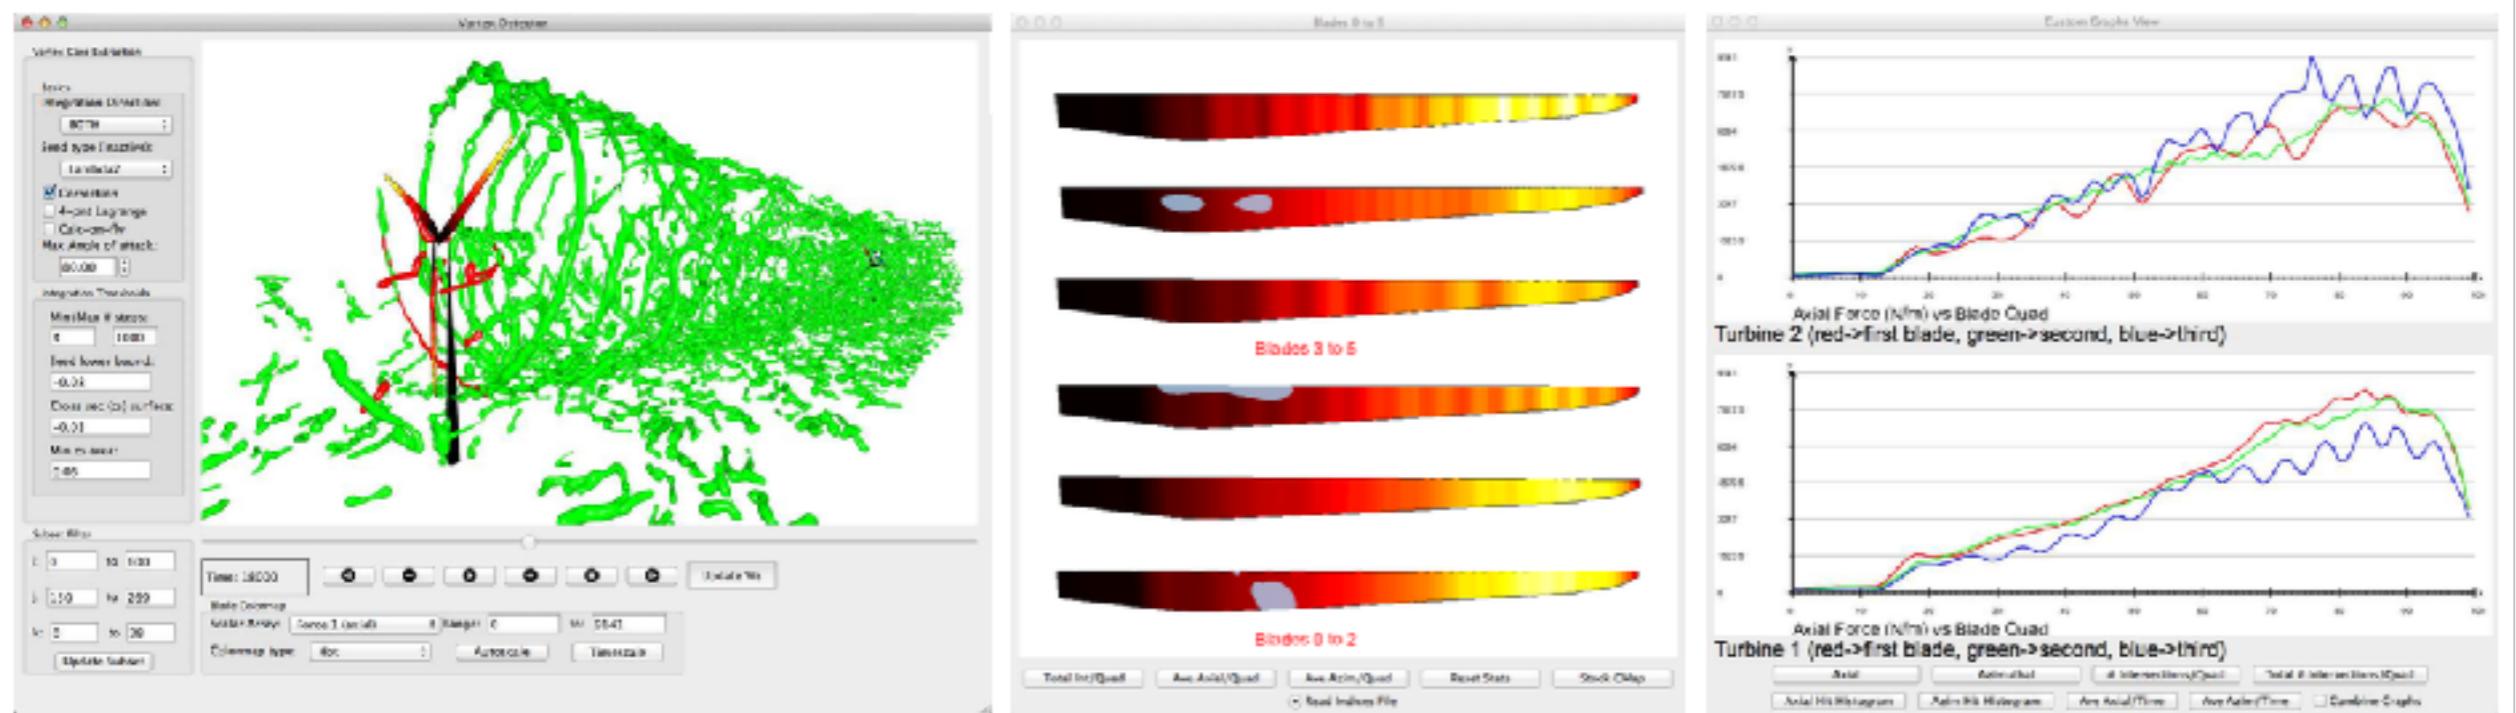
intersections (2D)

S. Shafii, H. Obermaier, R. Linn, E. Koo, M. Hlawitschka, CG, B. Hamann, K. I. Joy: *Visualization and Analysis of Vortex-Turbine Intersections in Wind Farms*. IEEE TVCG 19(9):1579-1591, 2013

Feature-Based Visualization



Feature-based visualization system:
coupling of feature definition, extraction, and visualization.



In this case, augmented with statistical feature information.

S. Shafii, H. Obermaier, R. Linn, E. Koo, M. Hlawitschka, CG, B. Hamann, K. I. Joy: *Visualization and Analysis of Vortex-Turbine Intersections in Wind Farms*. IEEE TVCG 19(9):1579-1591, 2013

Topological Techniques in Visualization



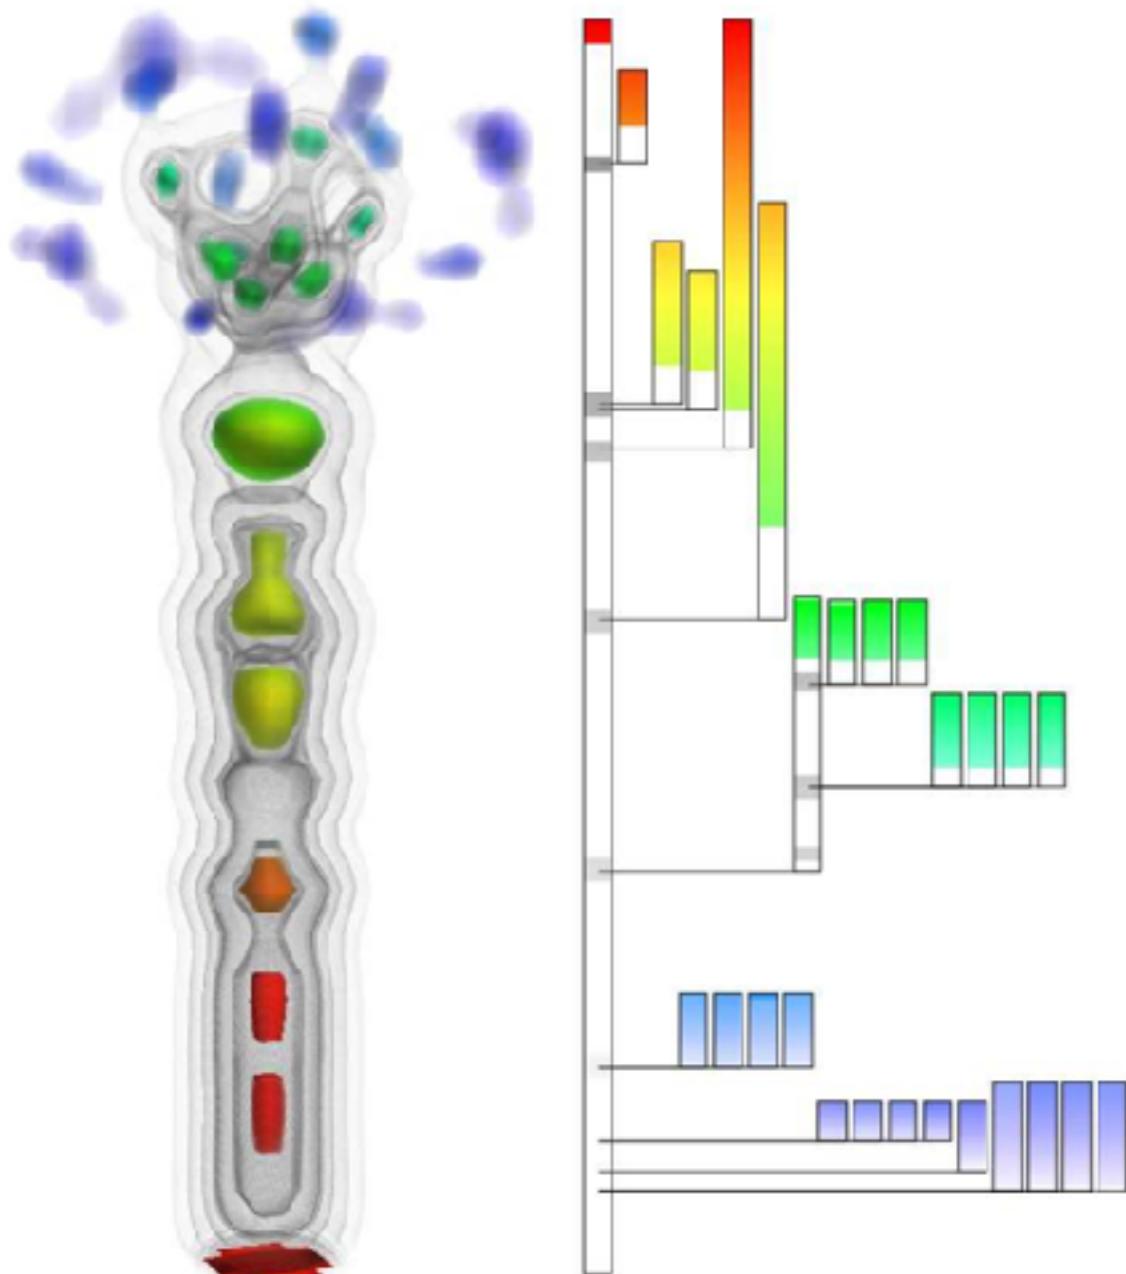
Topology is used to describe similarity up to smooth deformation.

Fundamentally, the mathematical notion of **topological equivalence** is applied to treat a complex dataset as a simpler, equivalent one; hence,

- to **abstract** complex datasets into comprehensible representation, and
- to **simplify** the structure in a meaningful way.

Fits well with the feature-based visualization mental model.

Contour Tree: extrema and basins



[Weber et al.: *Topology-controlled volume rendering*. IEEE TVCG 13(2):330-41, 2007.]

Topology of Area-Preserving Maps



Nuclear fusion in a Tokamak-type reactor:
Visualization of plasma confinement through closed fieldlines.

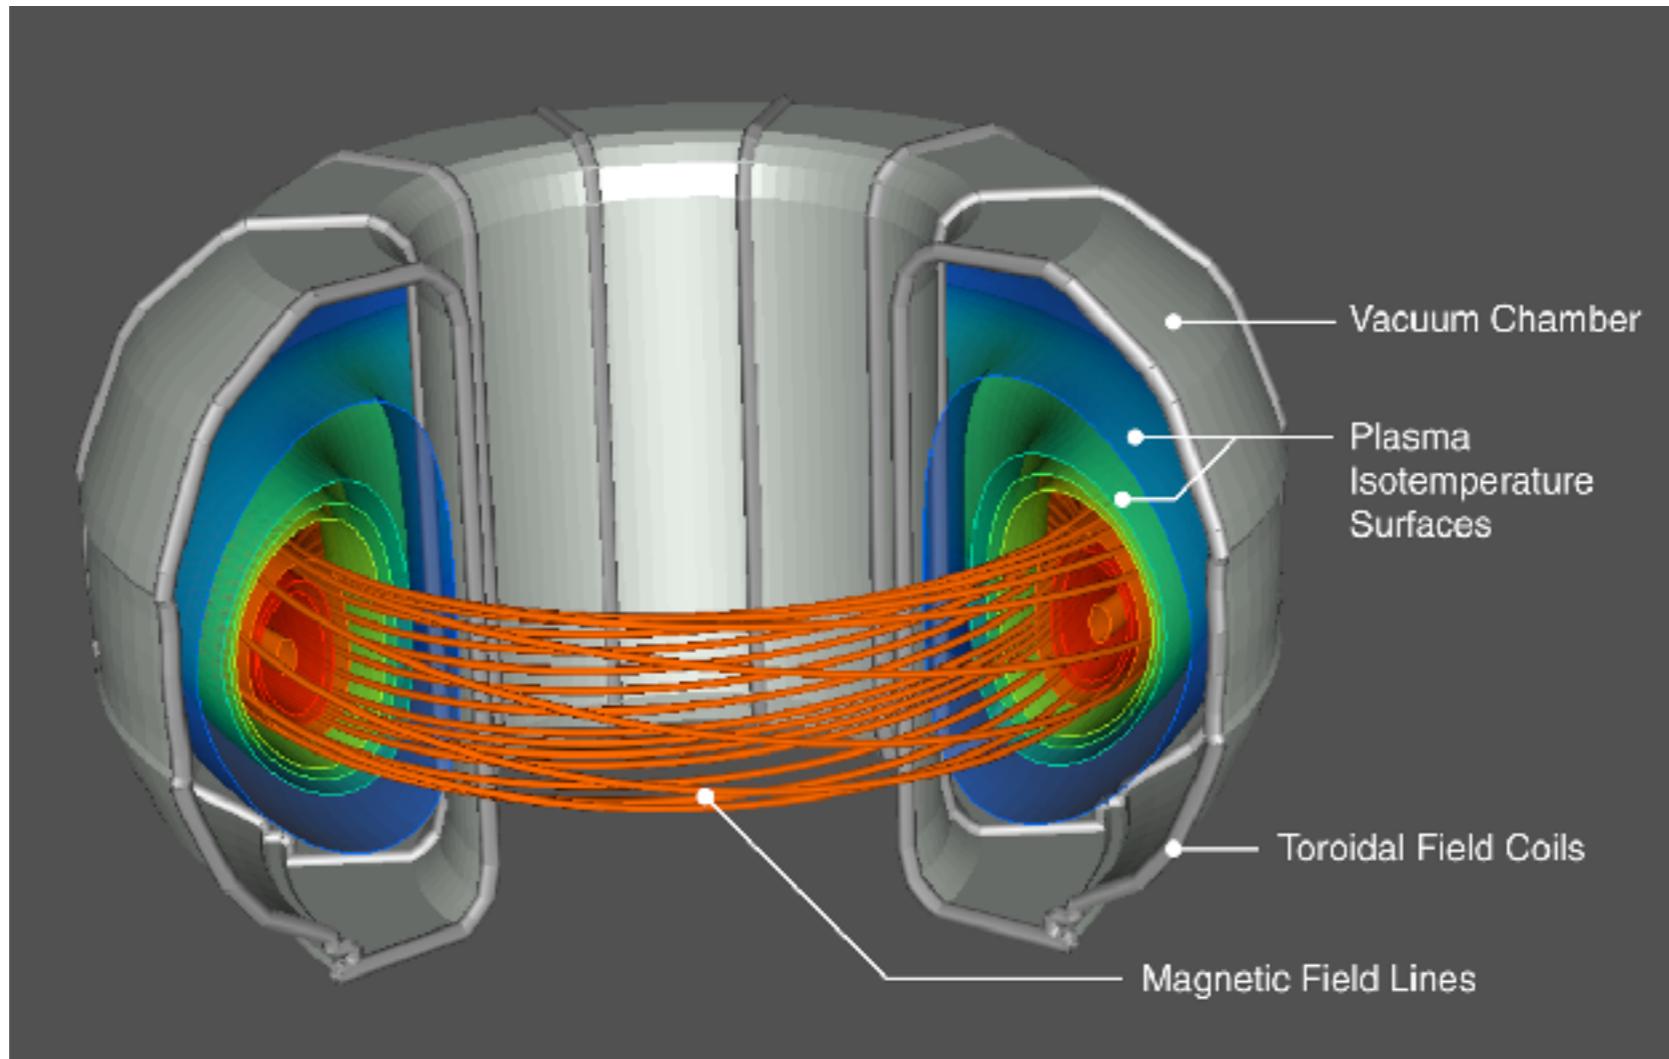


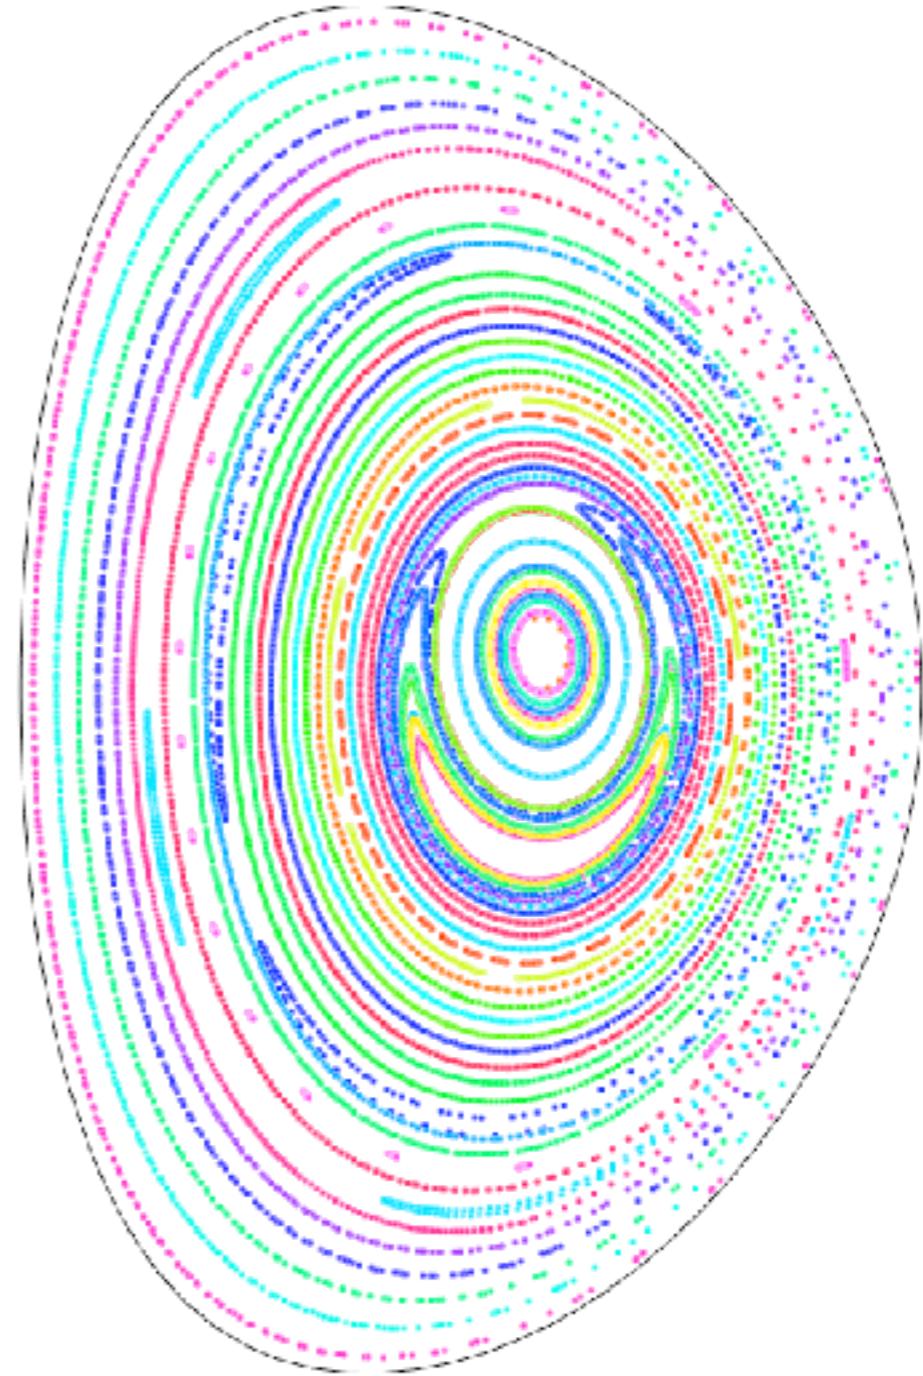
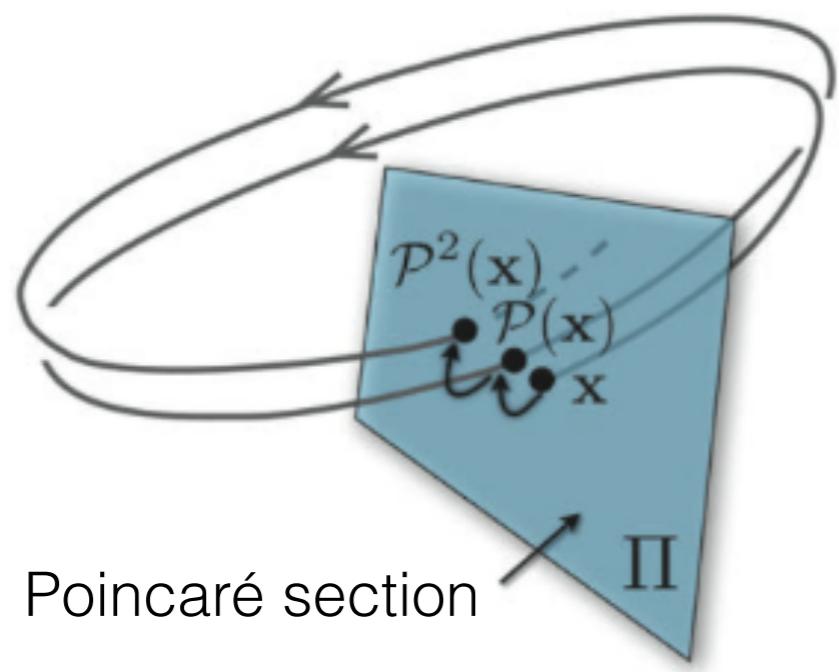
Image courtesy A. Sanderson, U. of Utah

Topology of Area-Preserving Maps



The magnetic field is described by a vector field; however, classical vector field topology is not applicable.

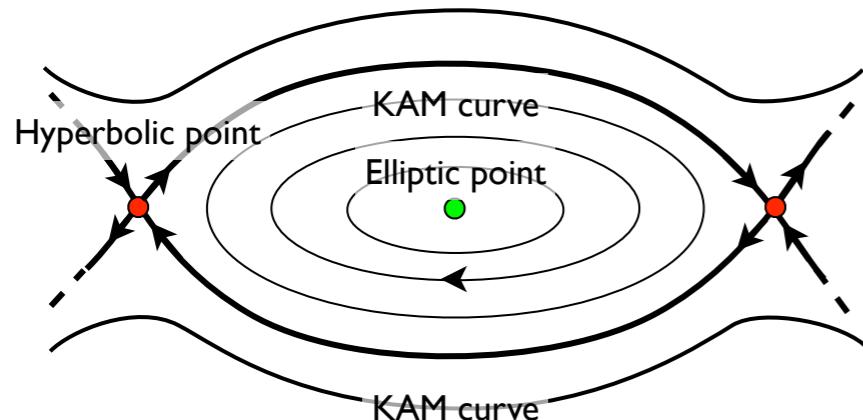
Typical visualization:
puncture plot (Poincaré map)



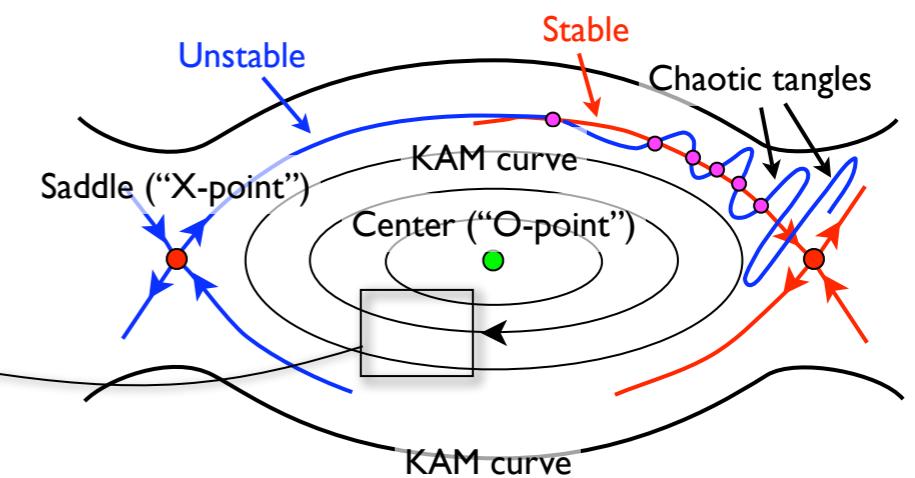
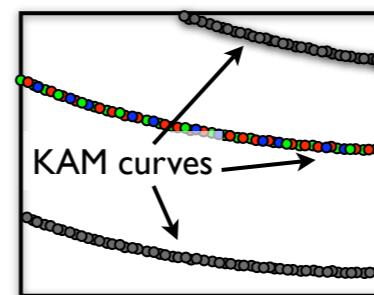
Topology of Area-Preserving Maps



Around periodic orbits, the return map exhibits a distinct topological structure:



non-chaotic



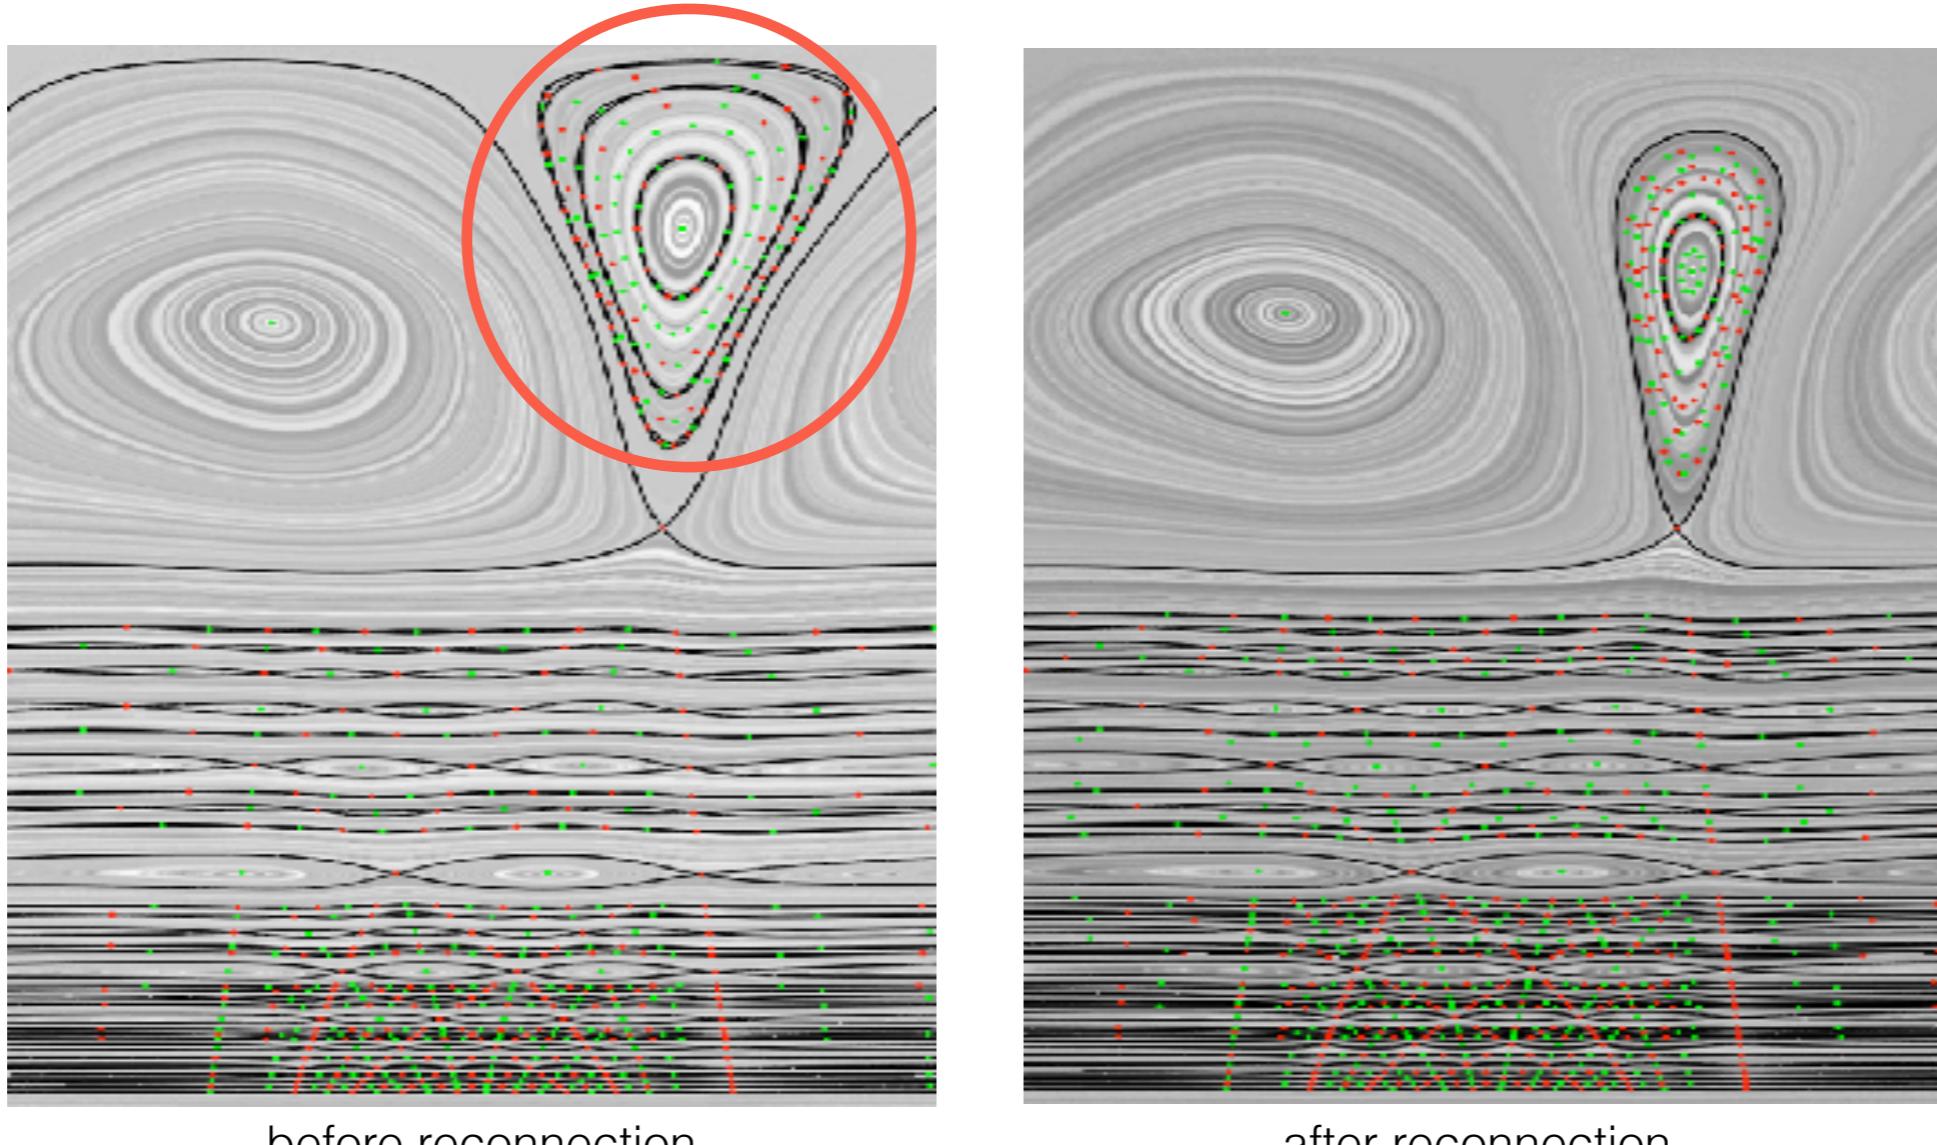
transition to chaos

This can be taken advantage of to enforce topological consistency for numerically inconsistent and challenging data.

Topology of Area-Preserving Maps



Visualization result: magnetic island reconnection



Poincaré map topology + Context (orbital averaging)

data: computed using NIMROD code, courtesy S. Kruger, Tech-X

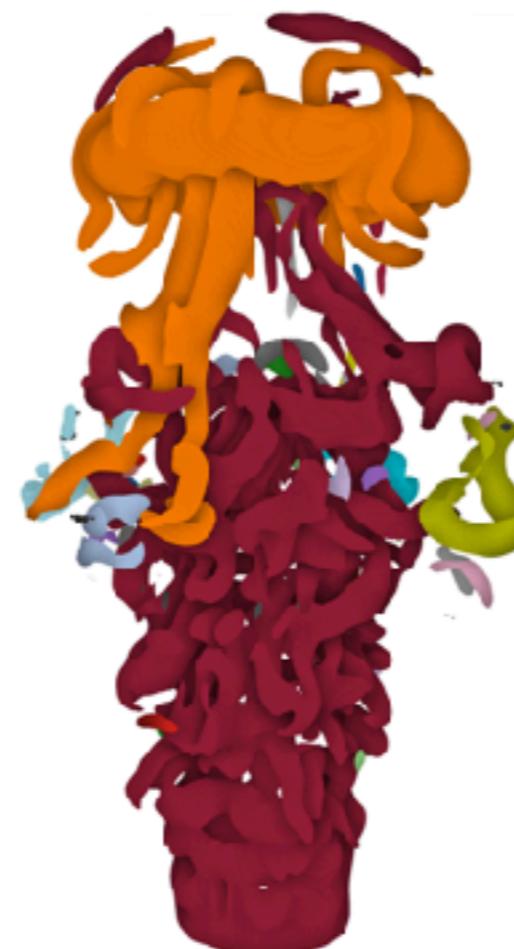
Topological Features



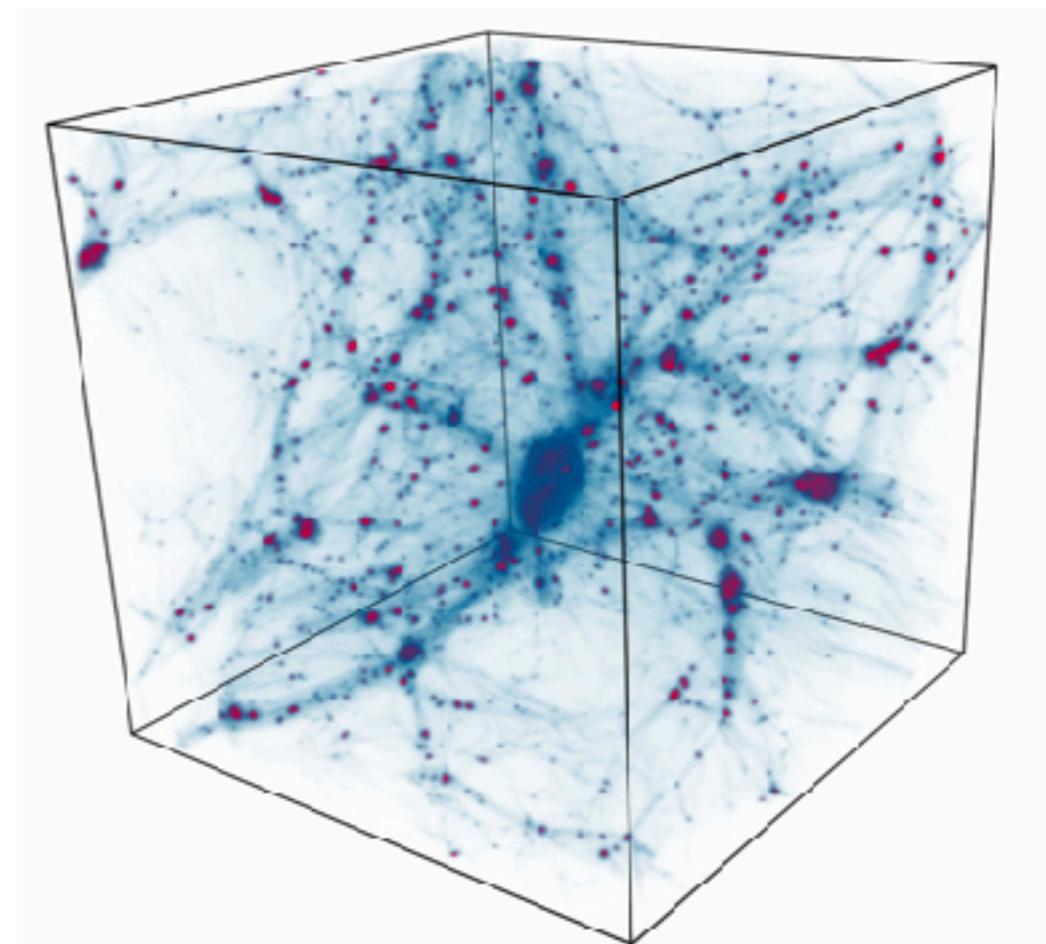
In many applications, features can be described as superlevel set components.



Viscous Fingering



Fluid Simulations

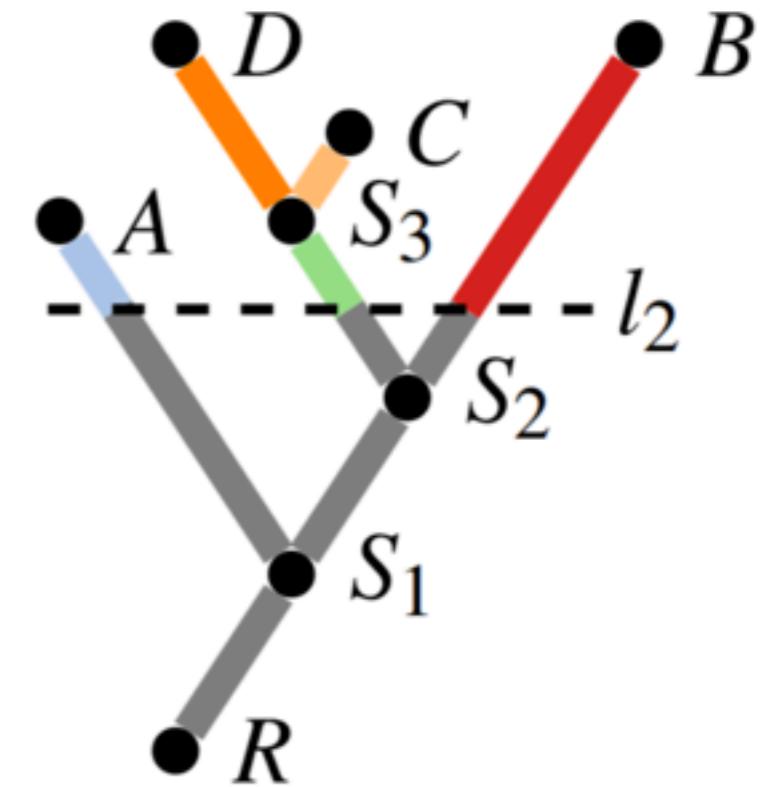
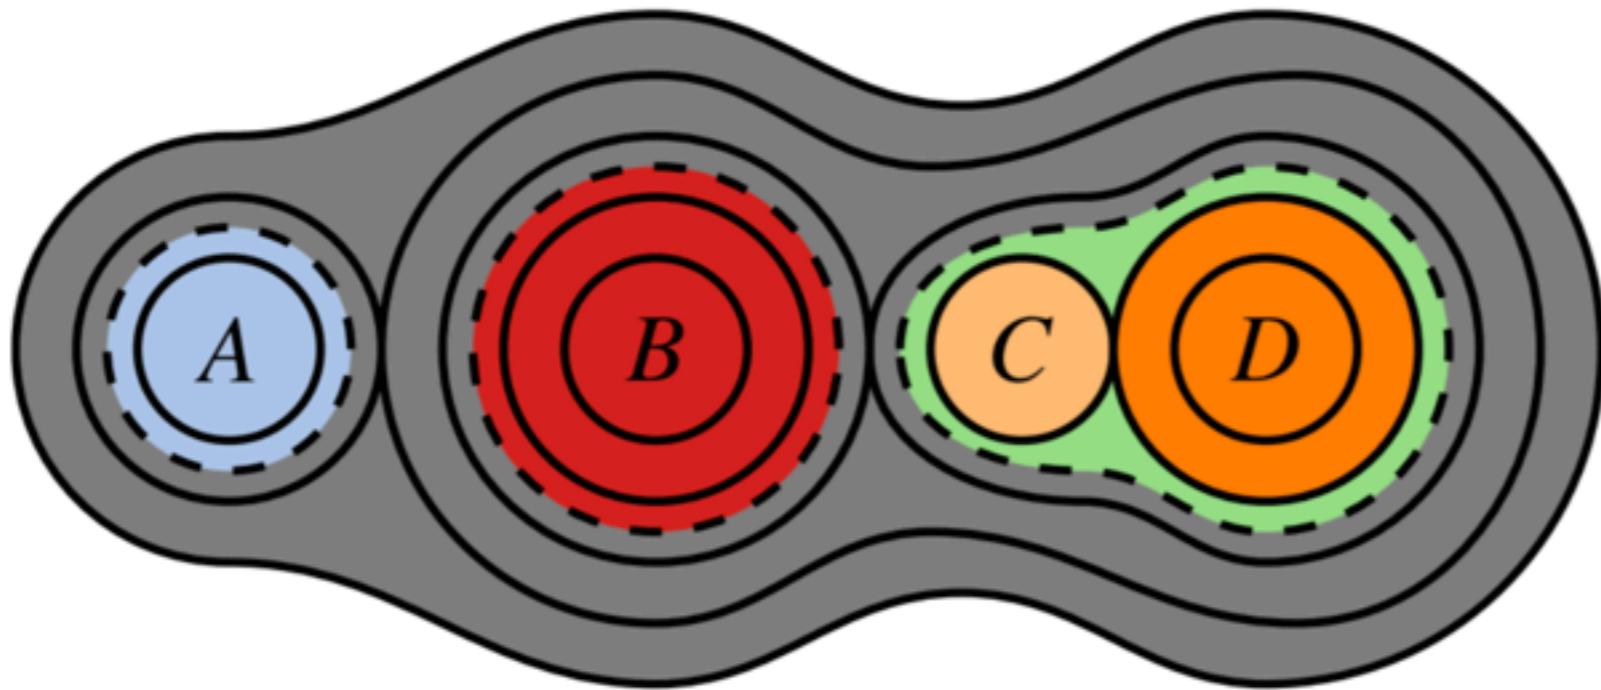


Dark Matter Halos

Topological Features



Simple example: Merge tree analysis

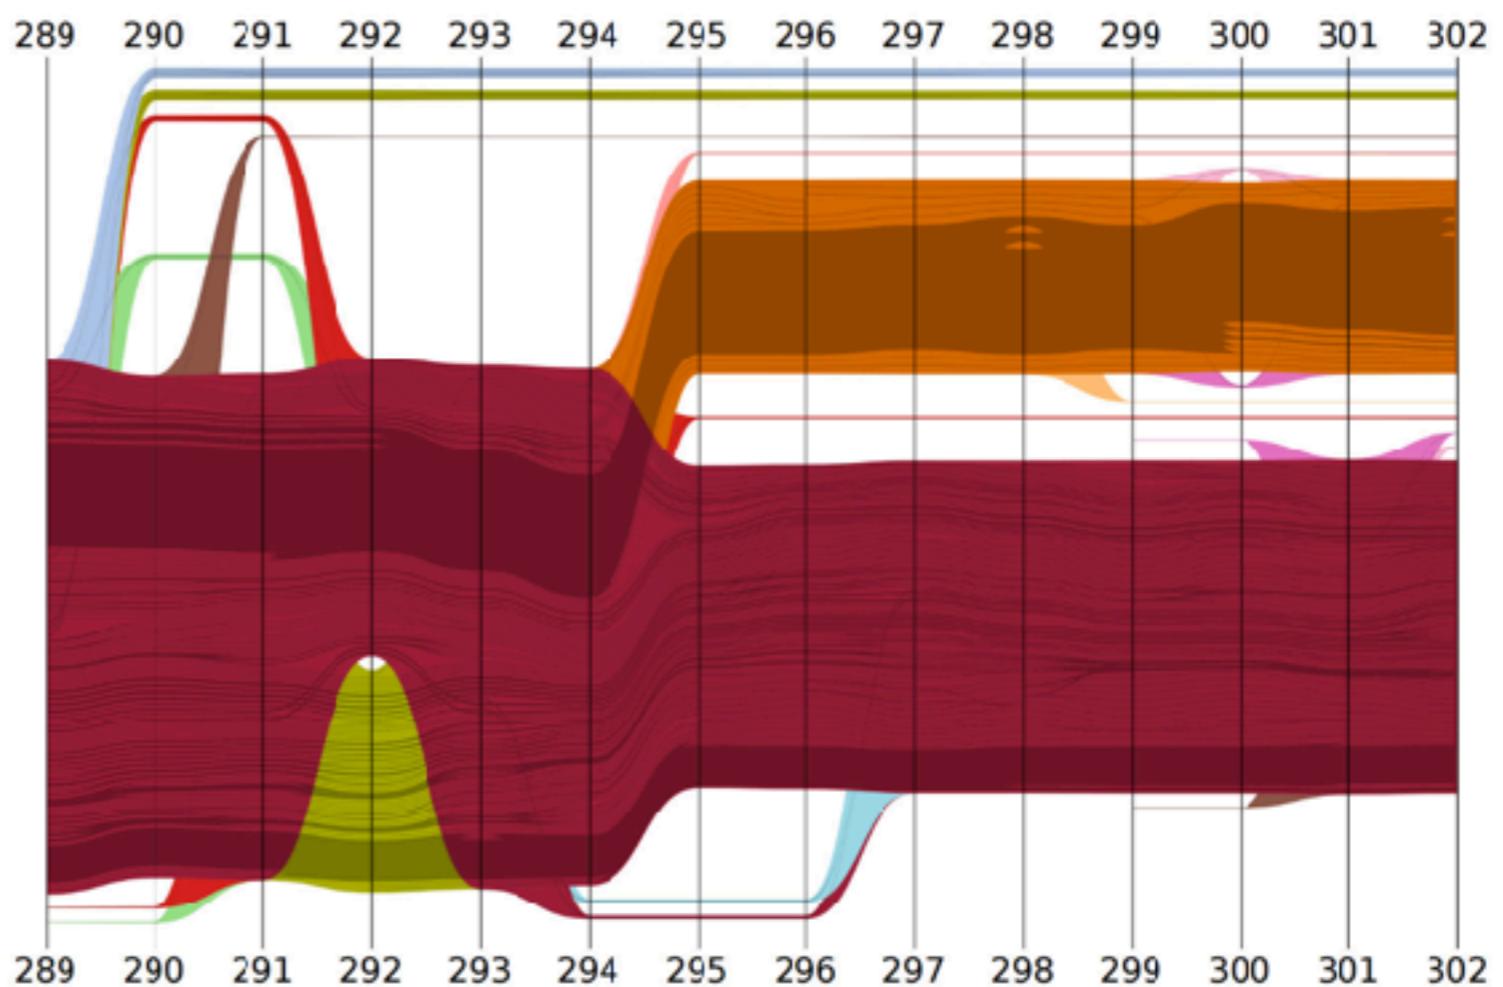
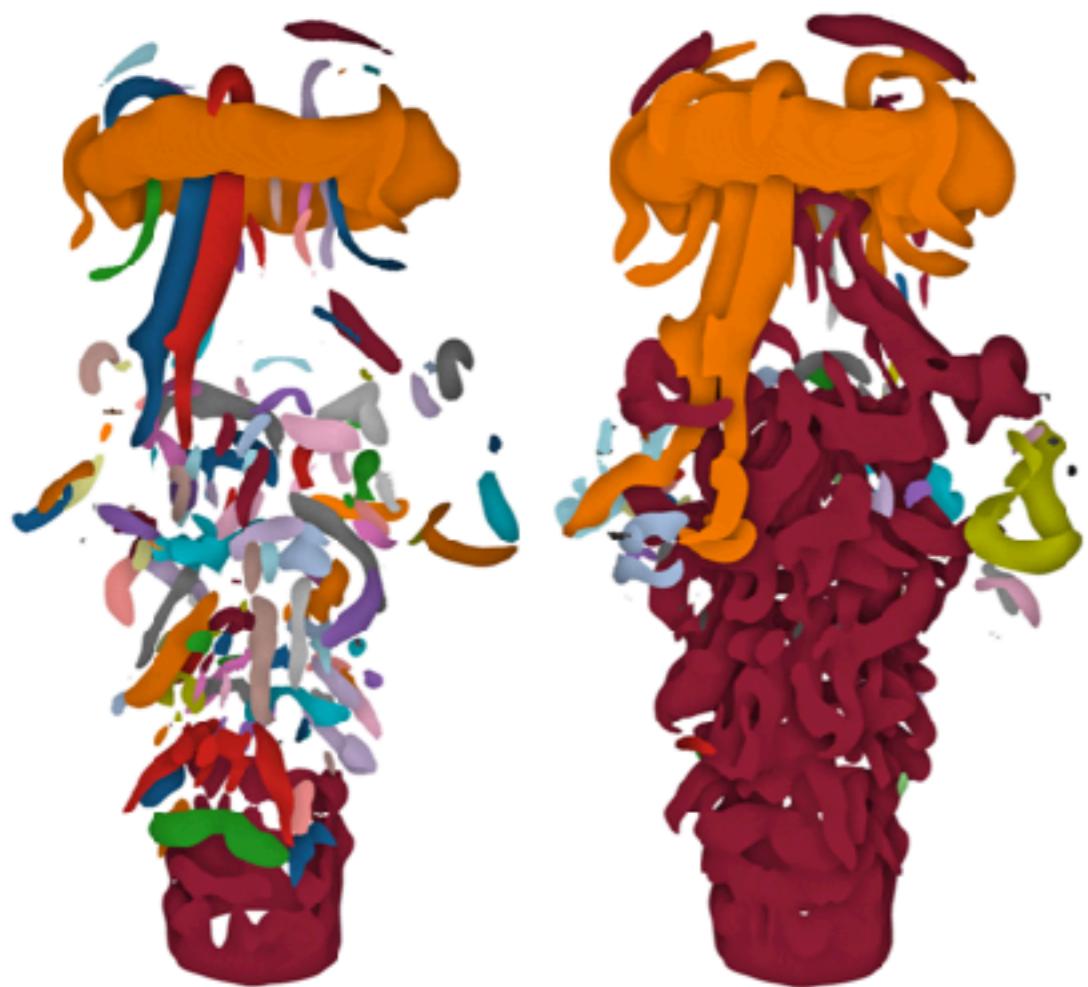


Evolution over time?

Tracking of Topological Features



Use topological abstraction to draw a nested graph (y-axis) over time (x-axis).

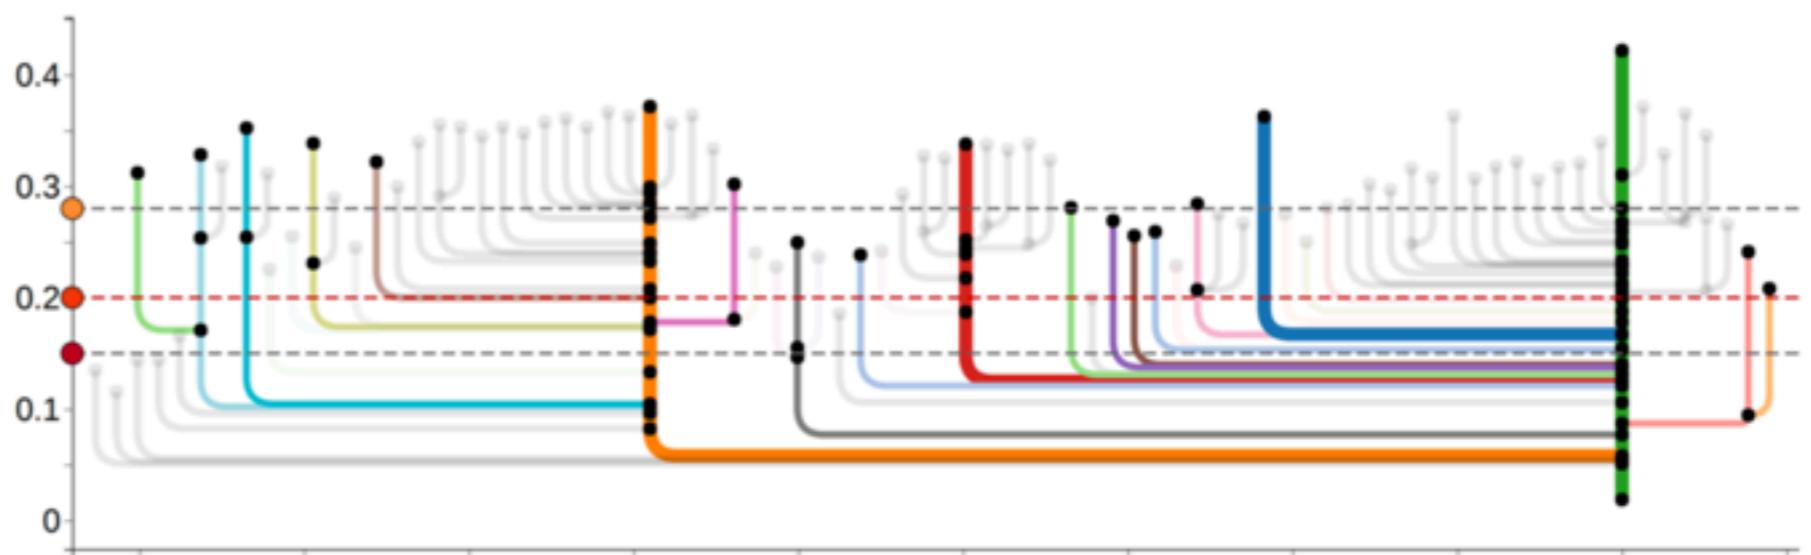
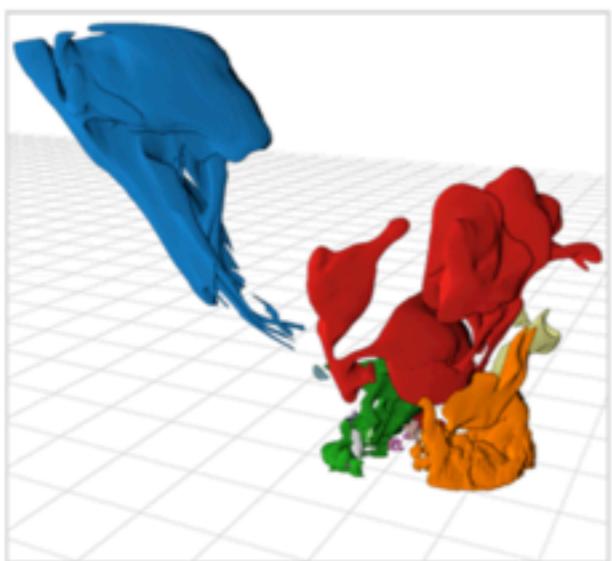
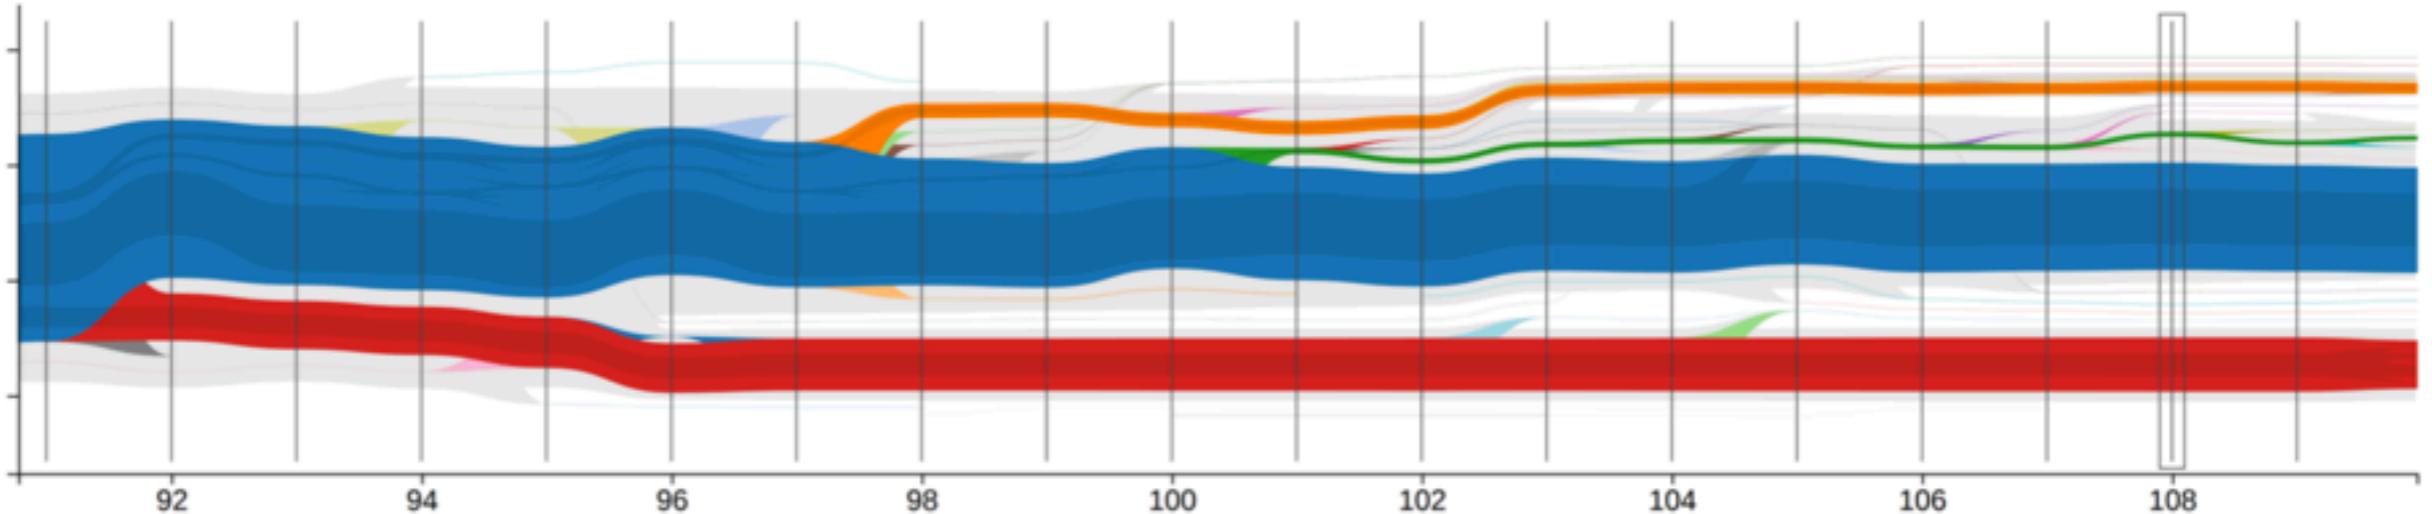


[Lukasczyk et al., Nested Tracking Graphs, EuroVis 2017]

Tracking of Topological Features



Case study: Asteroid Ocean Impact



[Lukasczyk et al., Nested Tracking Graphs, EuroVis 2017]

Parallel Algorithms for Large-Scale Data

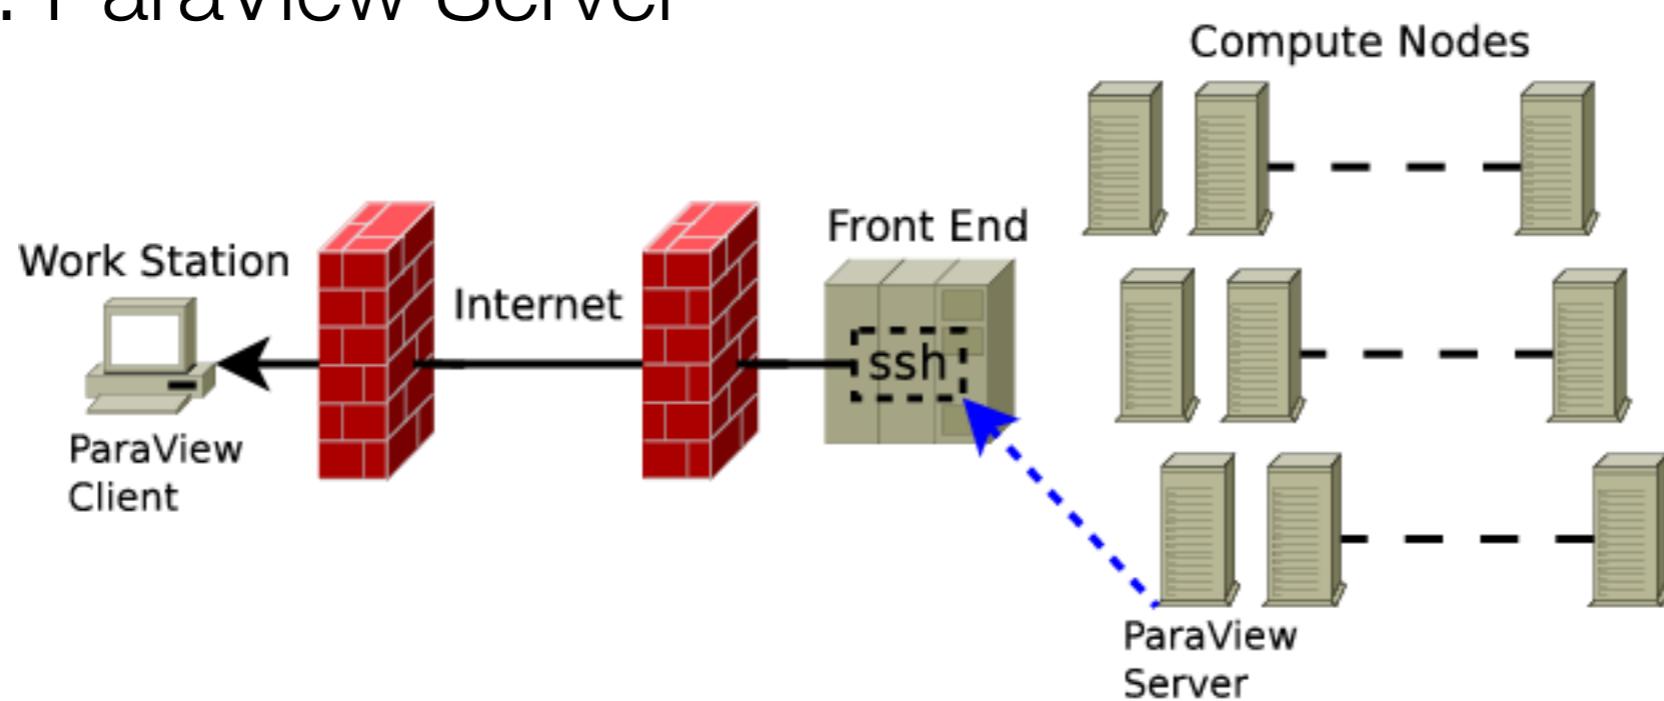
The Data Size Problem



Data is too large to be analyzed on a commodity workstation.

- Analyze / visualize the data where it is stored, directly on the cluster / supercomputer
- **remote visualization**

Example: ParaView Server



Scalable Algorithms



Large-scale data: want to use supercomputing resources also for visualization, i.e. run visualization on many processors.

This requires **scalable, parallel** algorithms.

Scalable algorithms, in a nutshell (p = number of processors):

Time on one processor: $T_1(n)$

Time on p processors: $T_p(n)$

Overhead: $p * T_p(n) - T_1(n)$

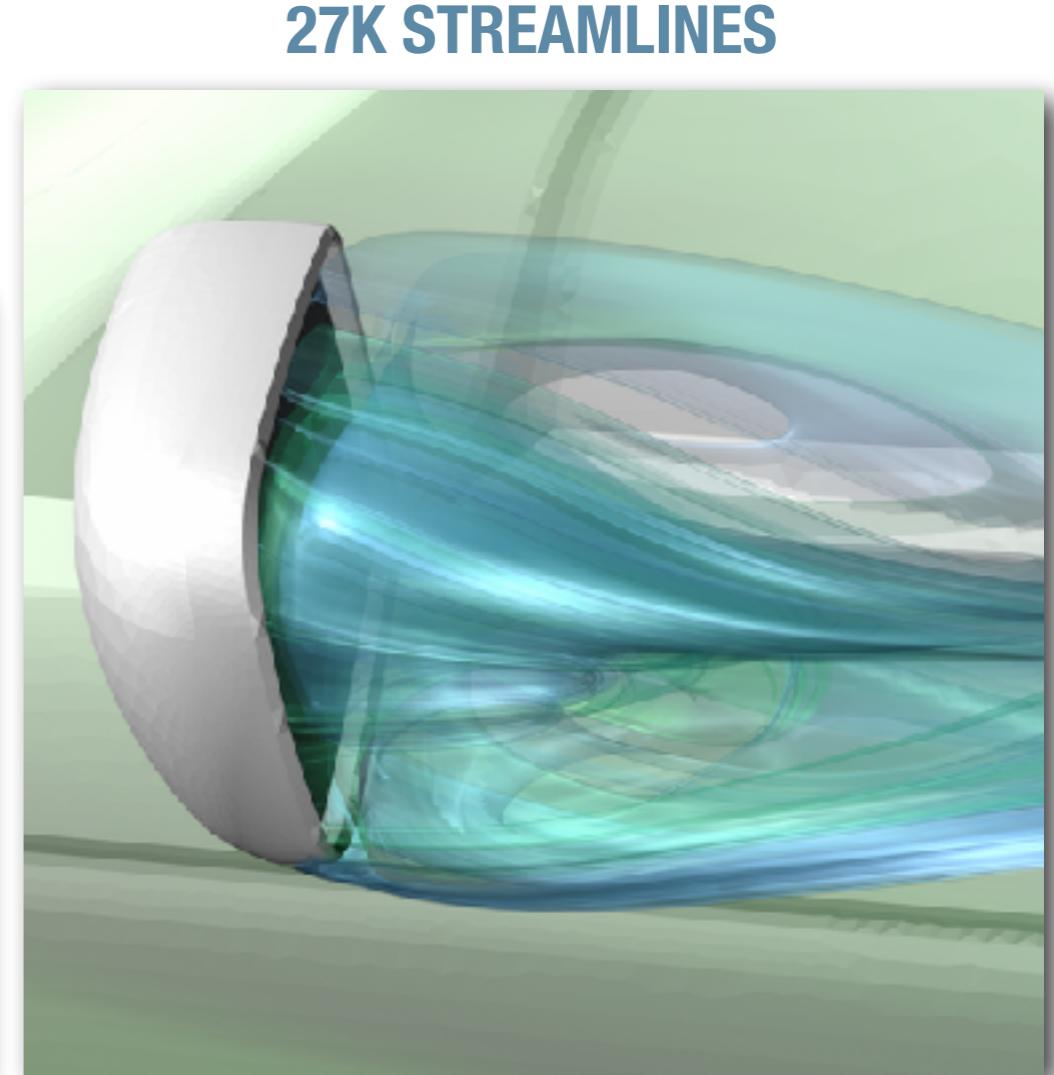
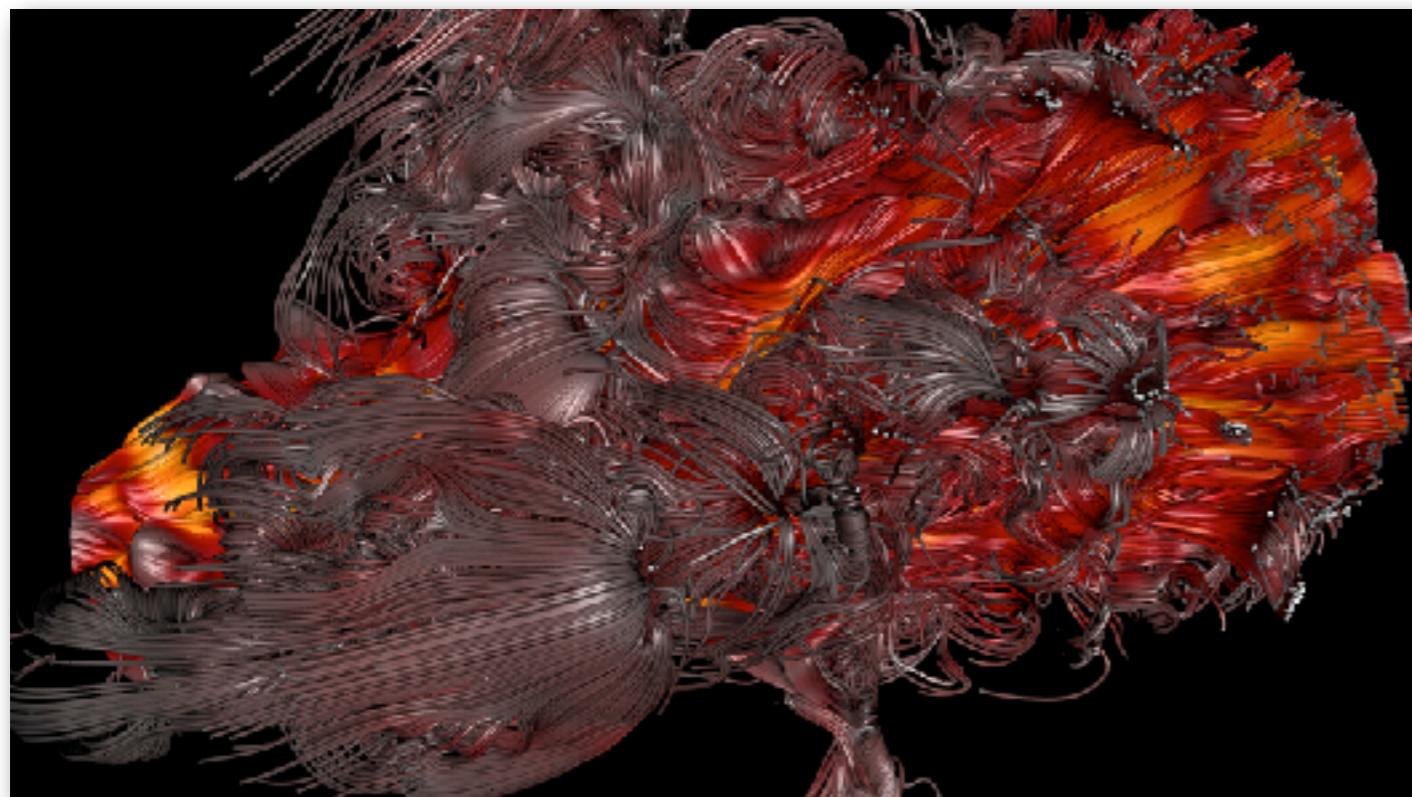
Scalable: overhead grows “slowly” (sub-linear) as a function of p .

Example: Integration-Based Visualization



Integration-based Vector Field Visualization

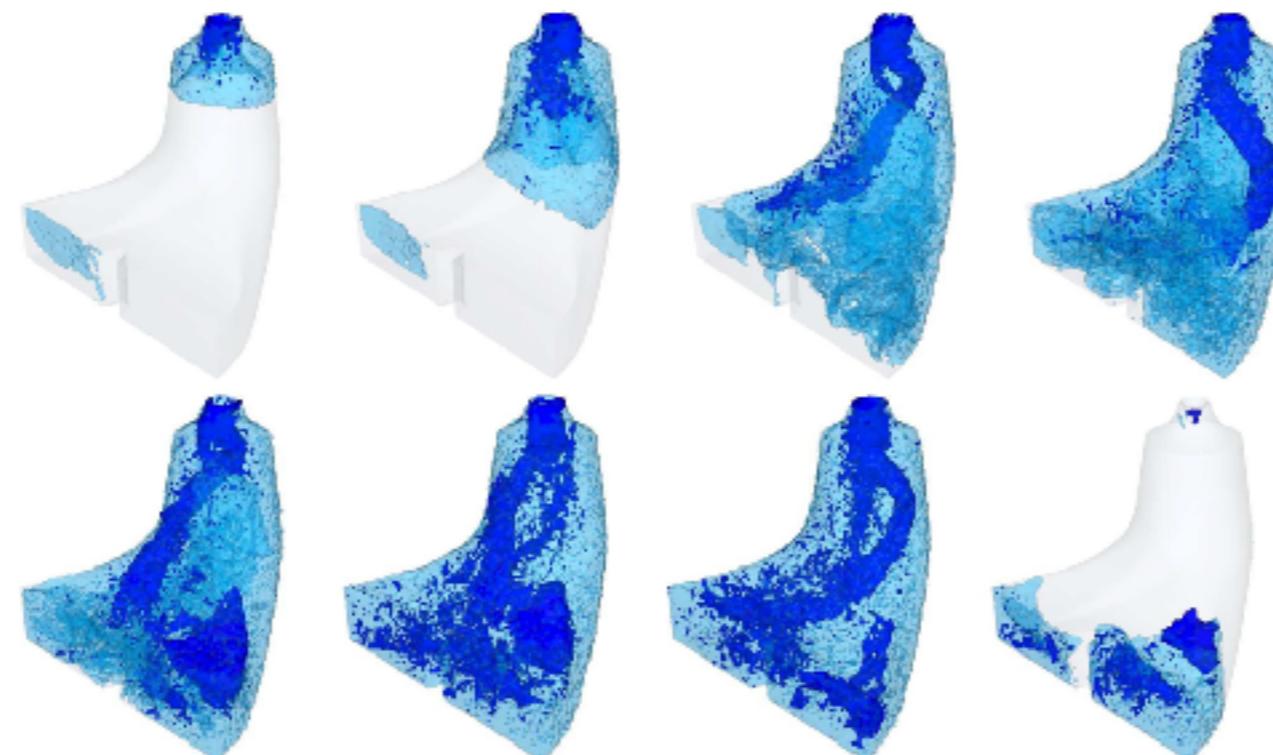
- Integral Curves and Surfaces



Example: Integration-Based Visualization



Pathline Predicates: integral curves in combination with Boolean predicates and set operations



pathline predicates over 250,000 pathlines

Definition and extraction of features as relevant sets of curves.

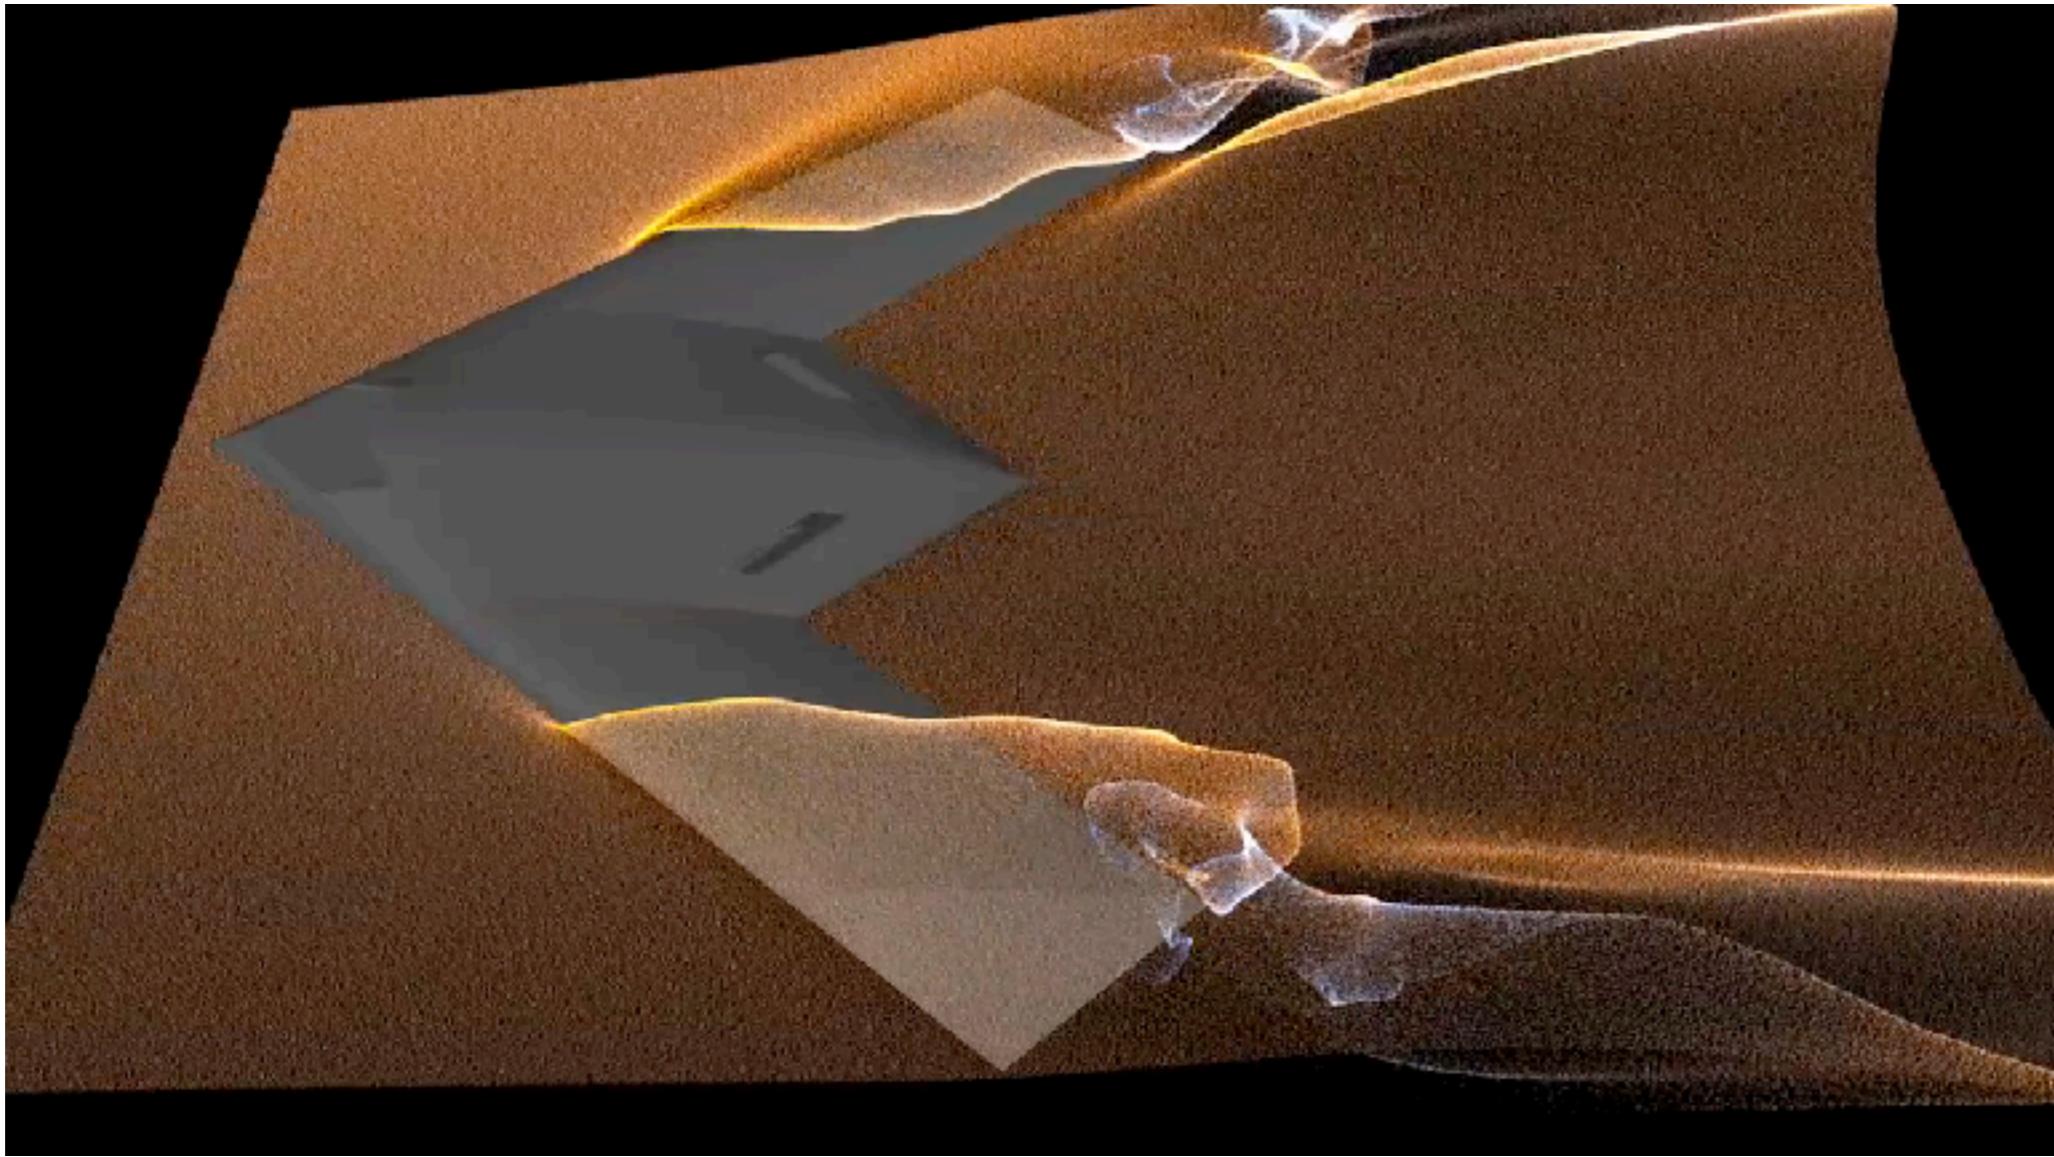
T. Salzbrunn, CG, G. Scheuermann, J. Meyer. *Pathline predicates and unsteady flow structures.* The Visual Computer 24(12):1039–1051, 2008

Turbine dataset courtesy R. Peikert, ETH Zürich

Example: Integration-Based Visualization



Typical (naïve) flow visualization: 1M particles moving with a flow.



Real-time due to GPU parallelization (all data resides in GPU memory).

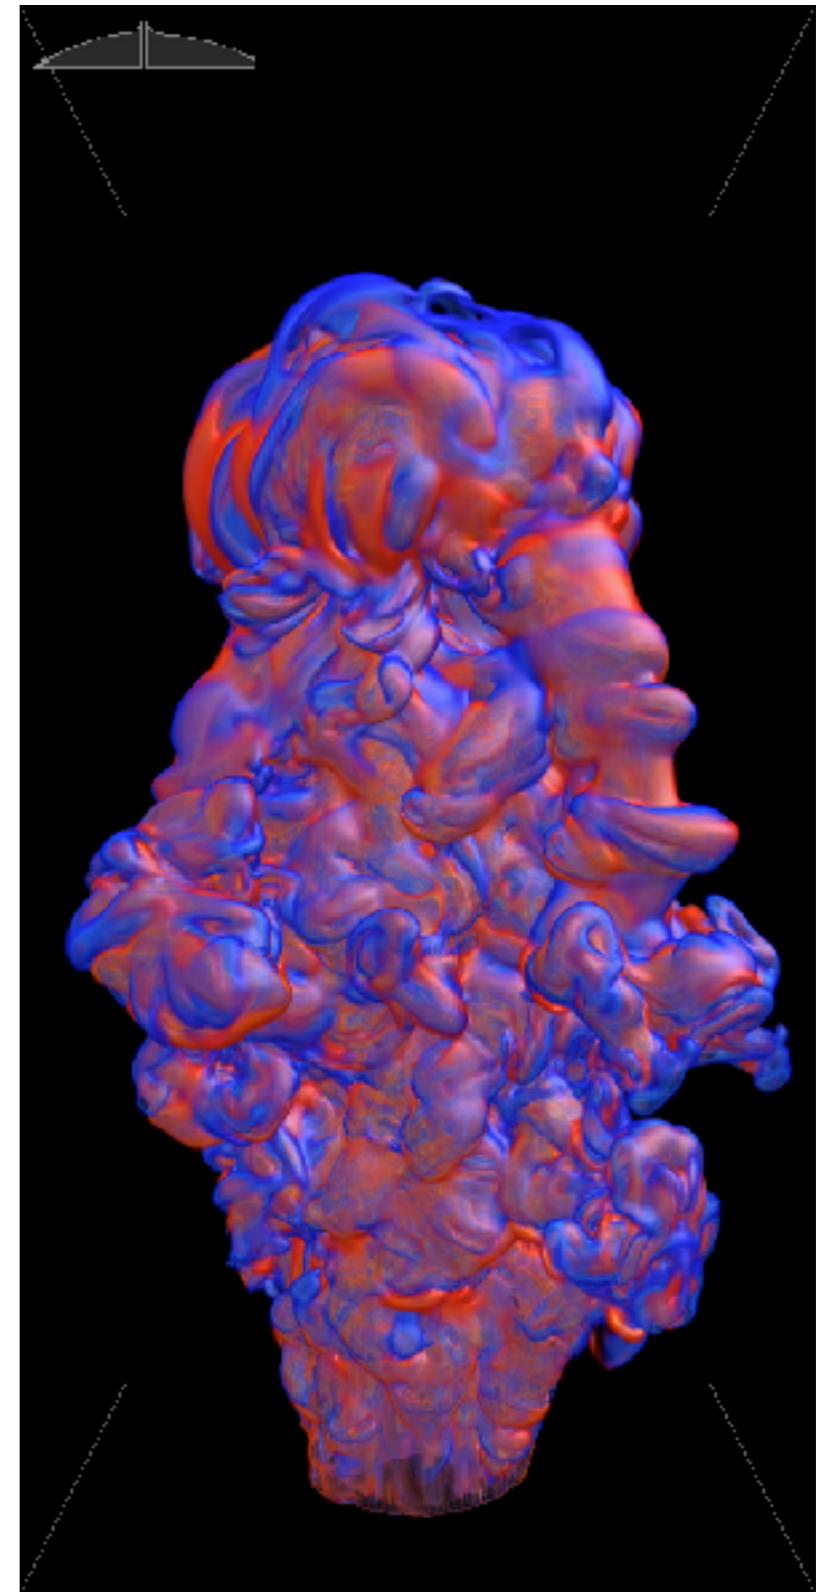
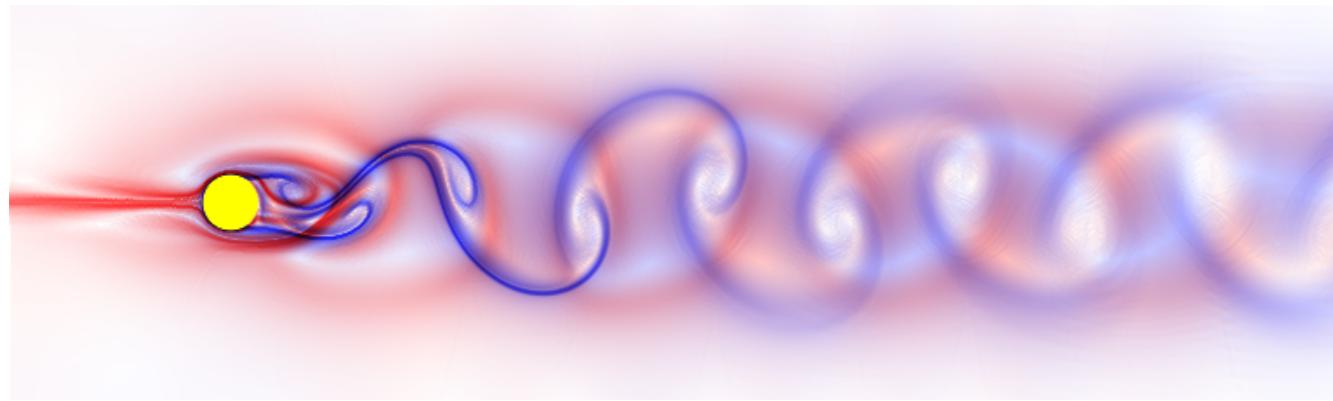
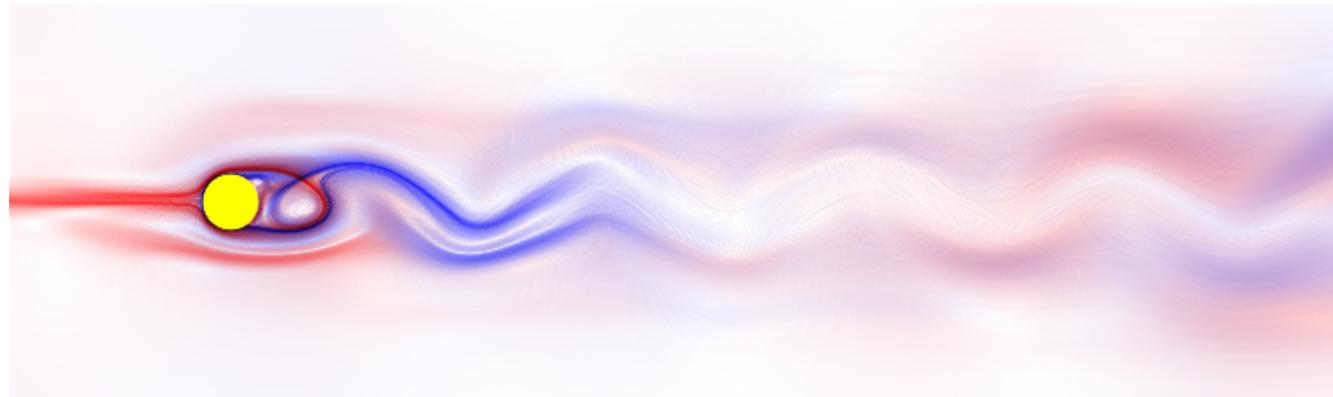
CG, K. I. Joy: *Fast, Memory-Efficient Cell Location in Unstructured Grids for Visualization*. IEEE TVCG 16(6):1541-1550.

Example: Integration-Based Visualization



Integration-based Vector Field Visualization

- Lagrangian Methods



5 billion pathlines

Parallel Algorithms



Integral Curves / Particles: conceptually simple

- Trace particle through vector field,
choice of numerical method and interpolation scheme
- Curves are independently computed,
parallelization is straightforward for in-core data

Large data:

- Data in blocks, all blocks do not fit in core of single machine
- Particles need to come together with correct
data blocks to allow propagation.

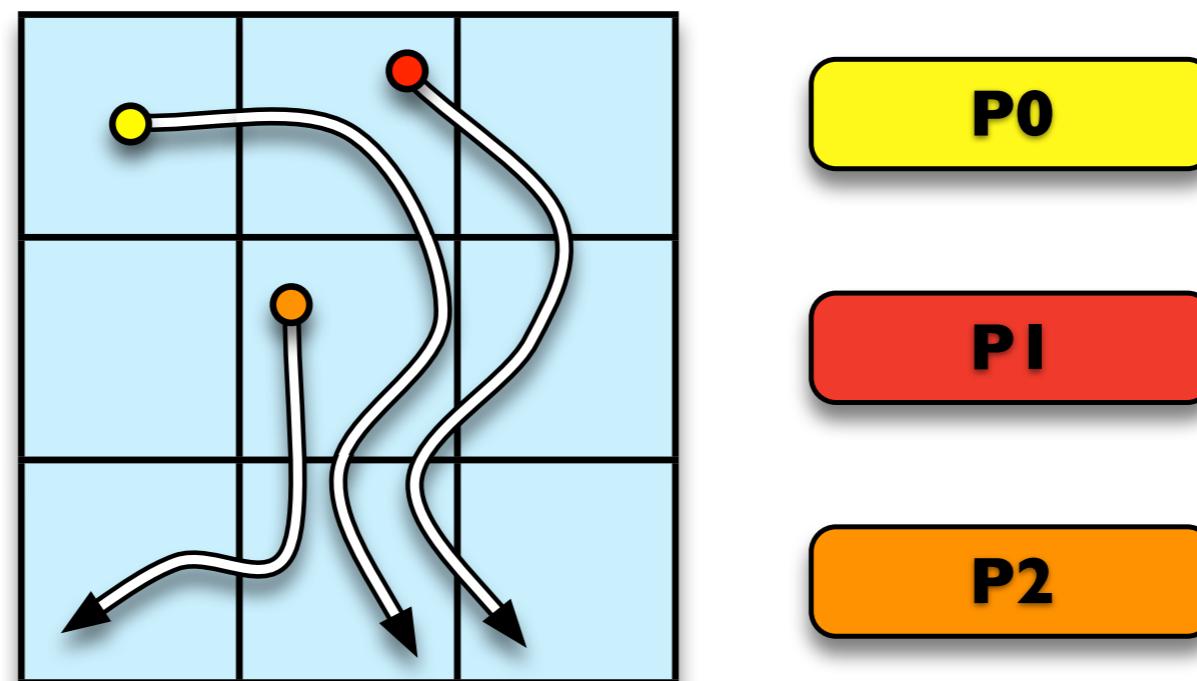
Parallel / Distributed Integration



Two basic parallelization/distribution schemes:

- **Parallelize over Seeds** (== Curves)

distributed computation, data on demand



Parallel / Distributed Integration

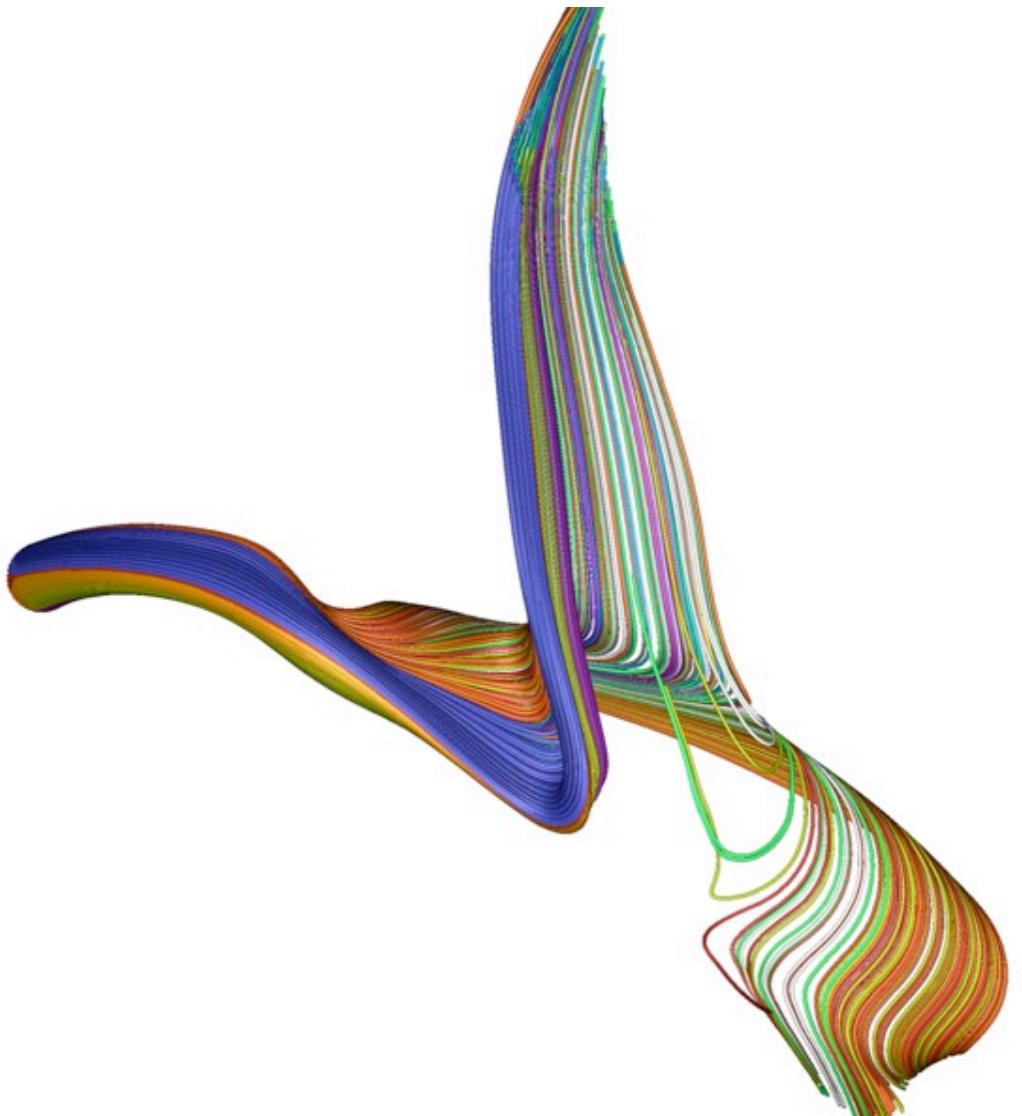


Example:

Supernova Core Collapse (10^9 cells in 4096 blocks)

Parallelize over Seeds

- lots of redundant I/O since same subset of data touched by many curves



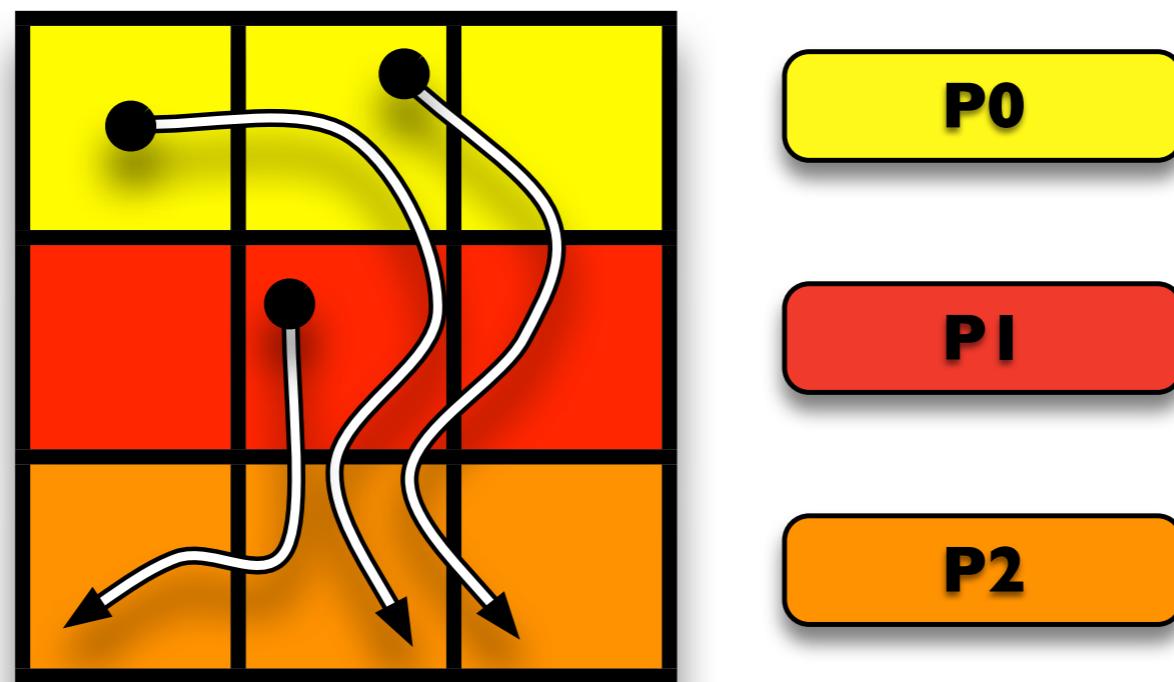
Parallel / Distributed Integration



Two basic parallelization/distribution schemes:

- **Parallelize over Blocks**

distributed data, computation on demand



Parallel / Distributed Integration

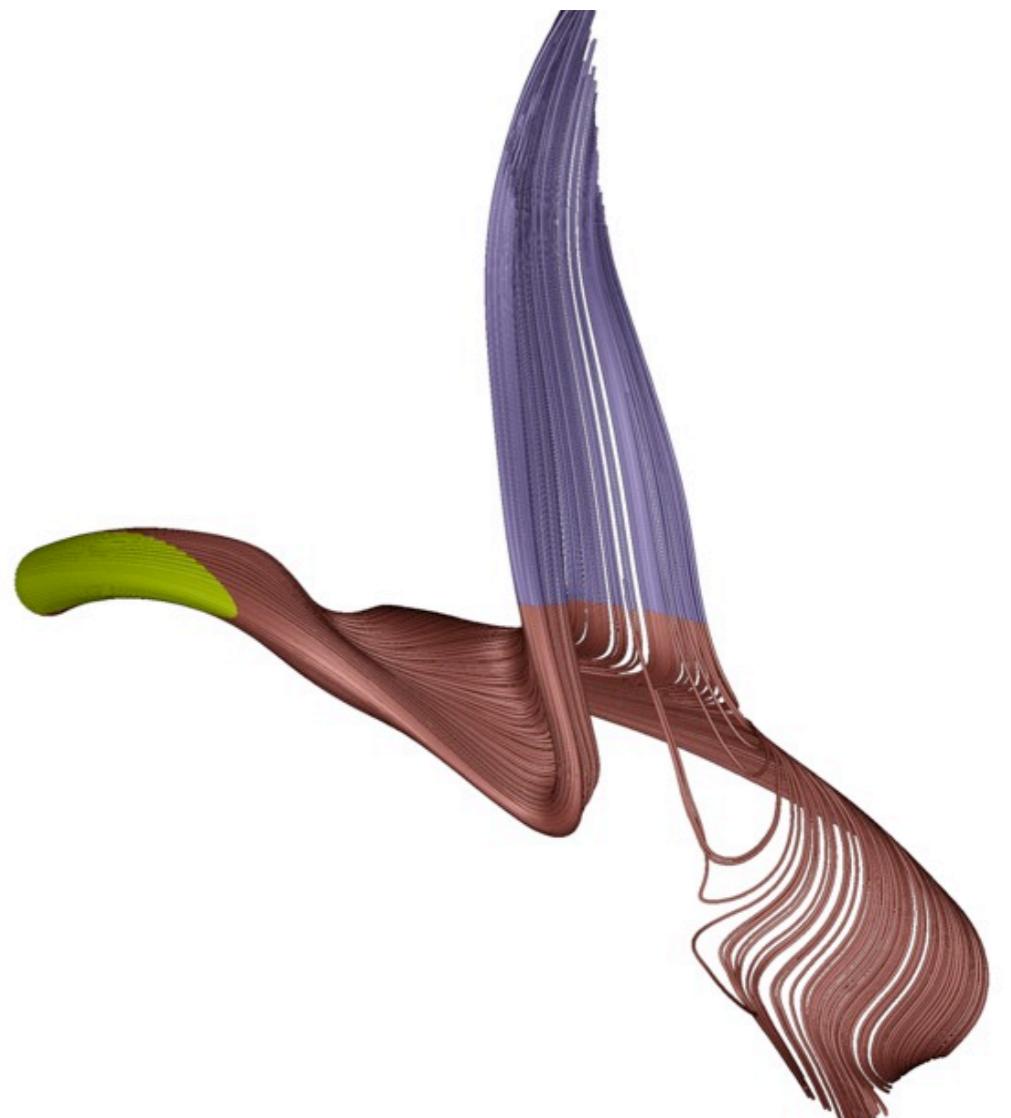


Example:

Supernova Core Collapse (10^9 cells in 4096 blocks)

Parallelize over Blocks

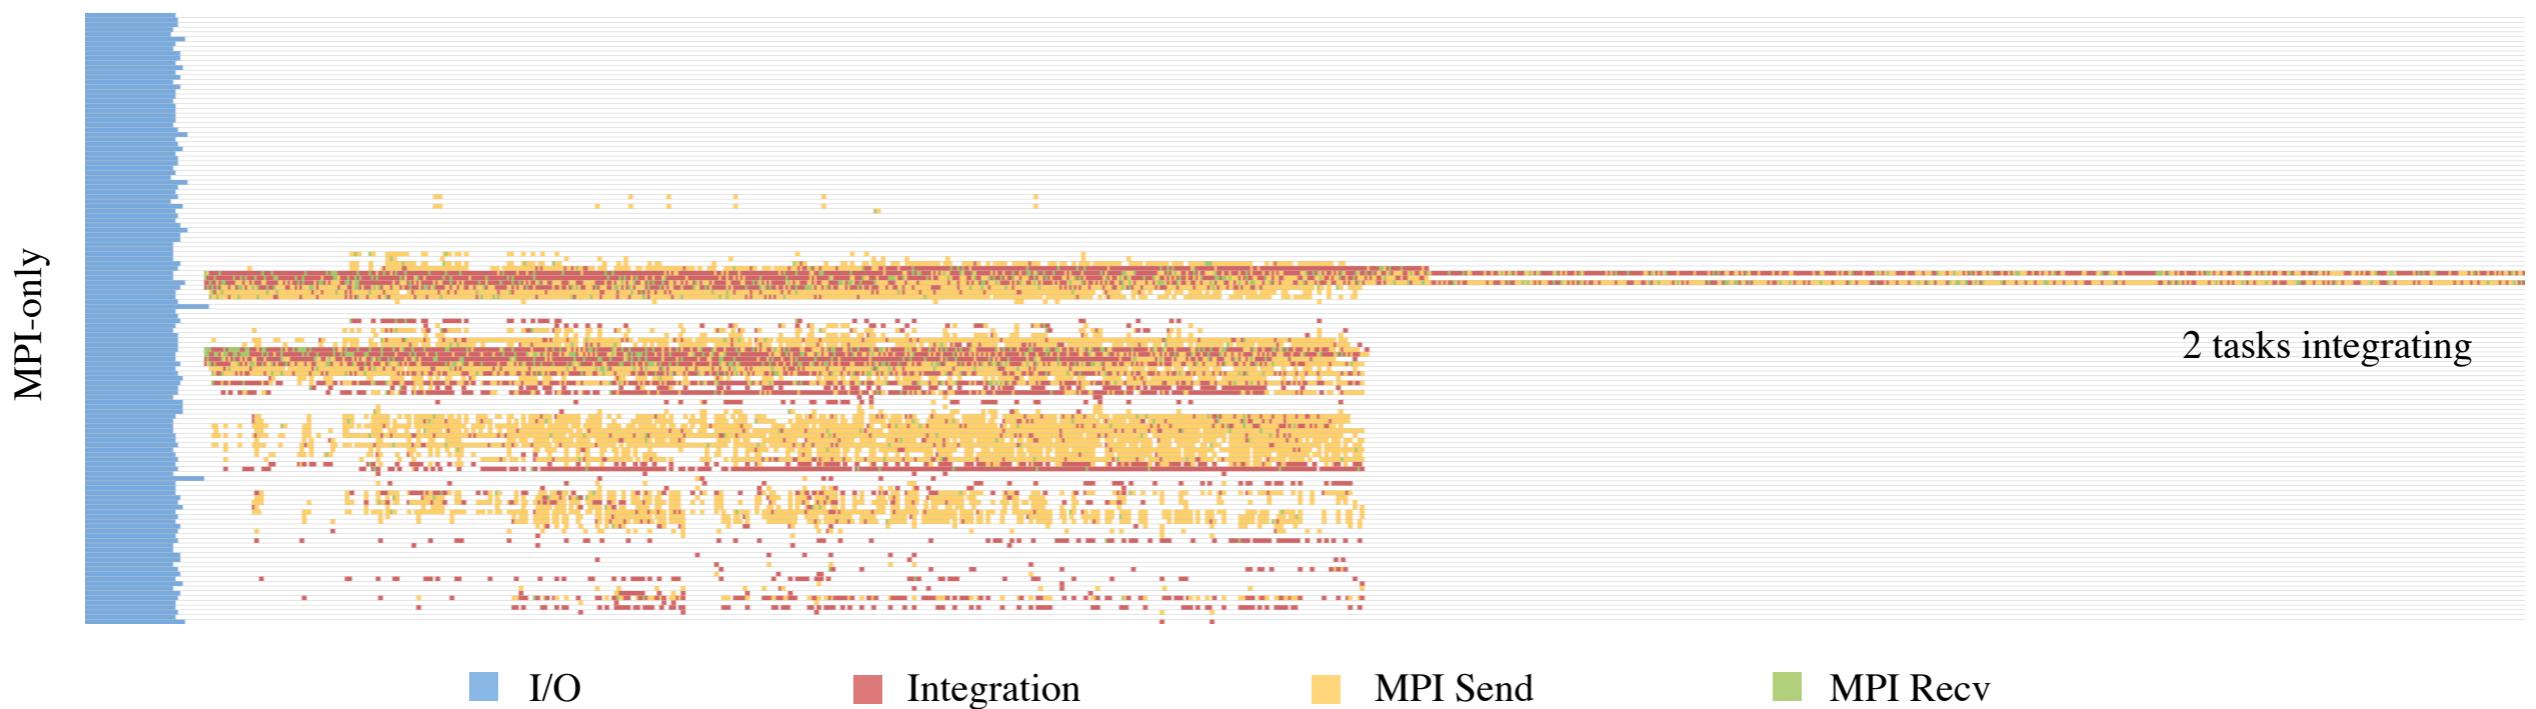
→ only 3 processors used



Parallel / Distributed Integration



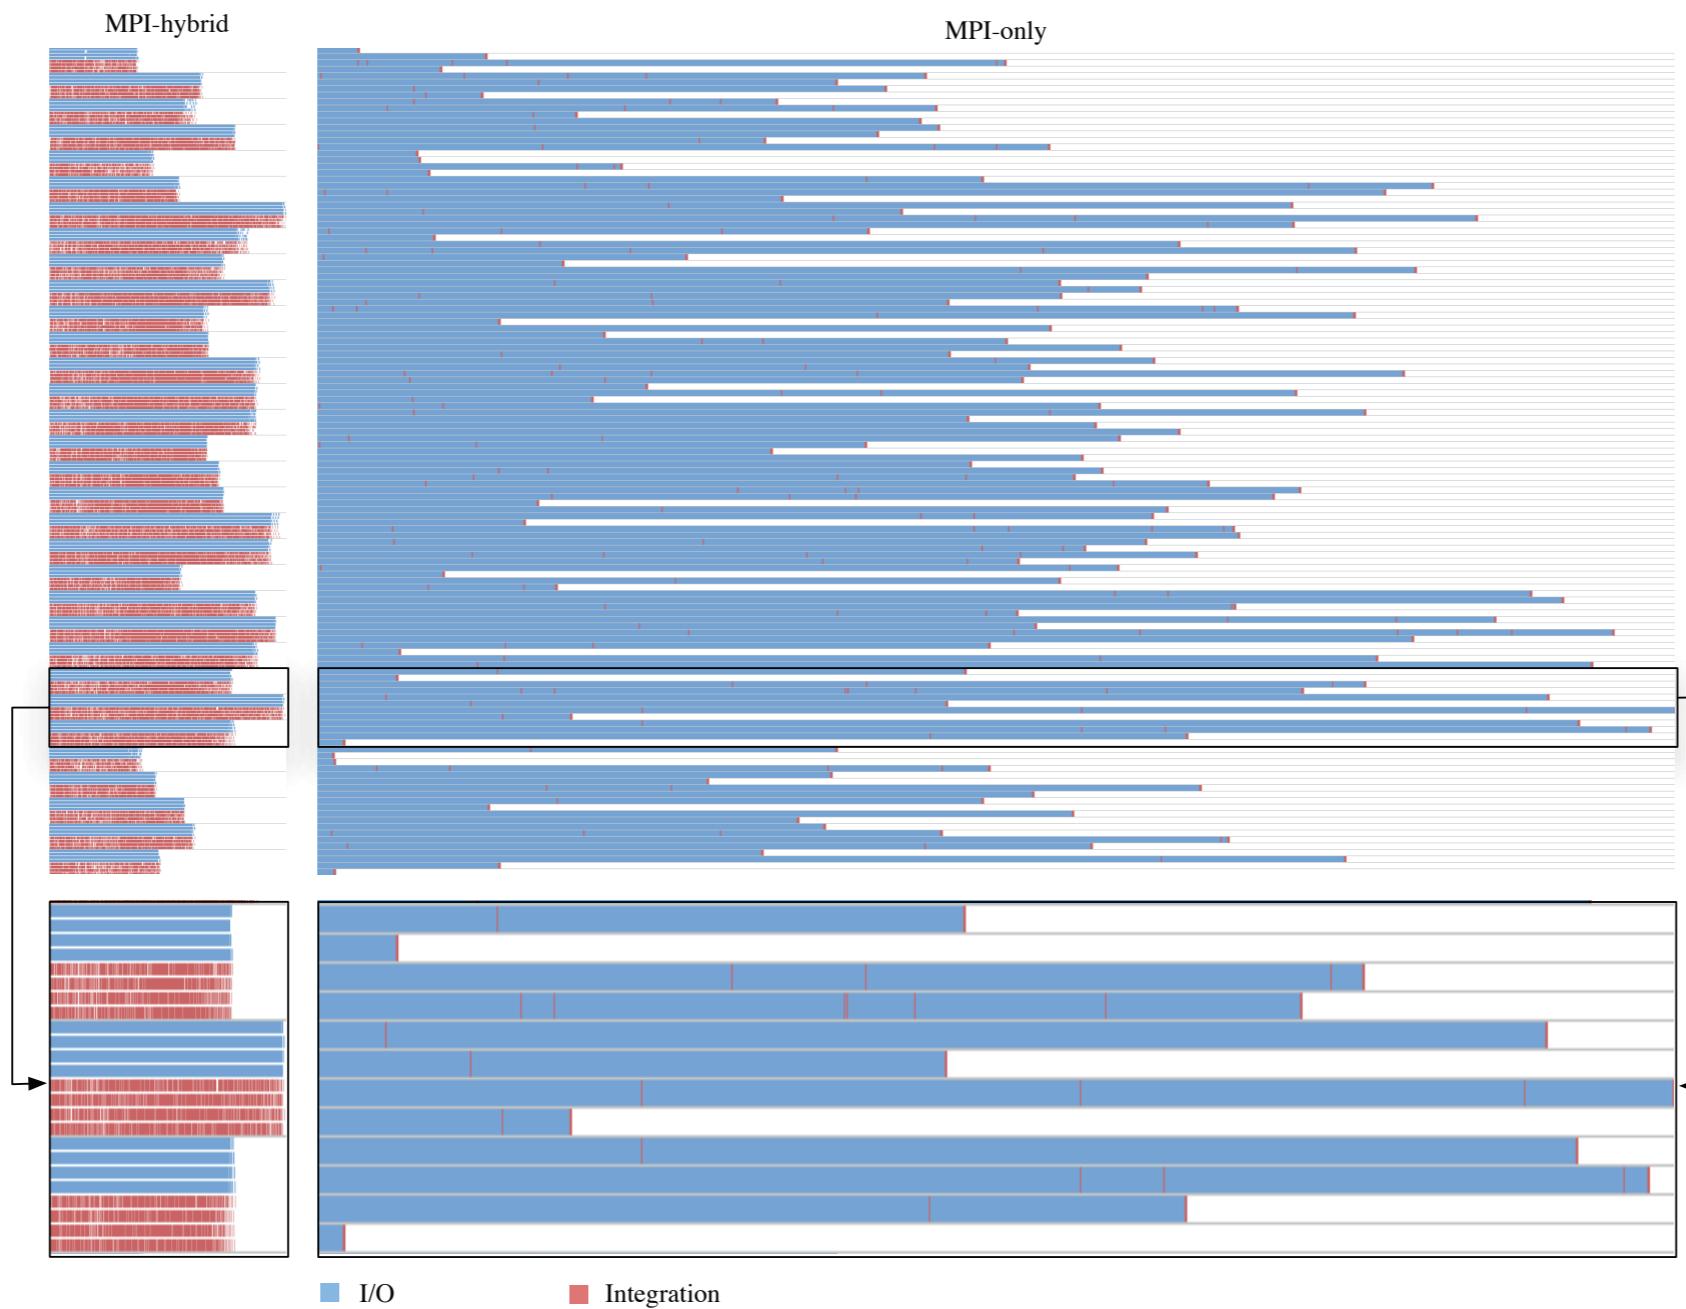
Degenerate case Gantt chart



Better Parallel Integration



Adaptive Load Balancing and Hybrid Parallelization



Pugmire, Garth, Childs, Joy, Bethel: *Scalable Computation of Integral Curves*. Proc. ACM Supercomputing, 2009.

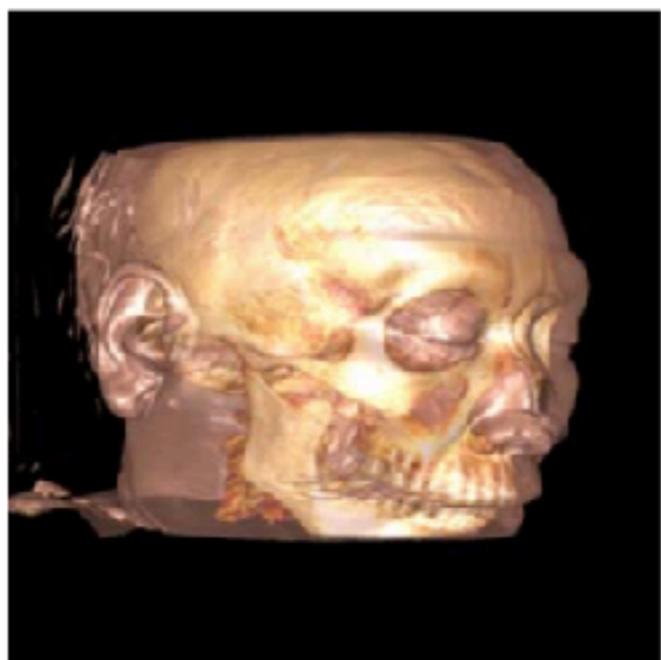
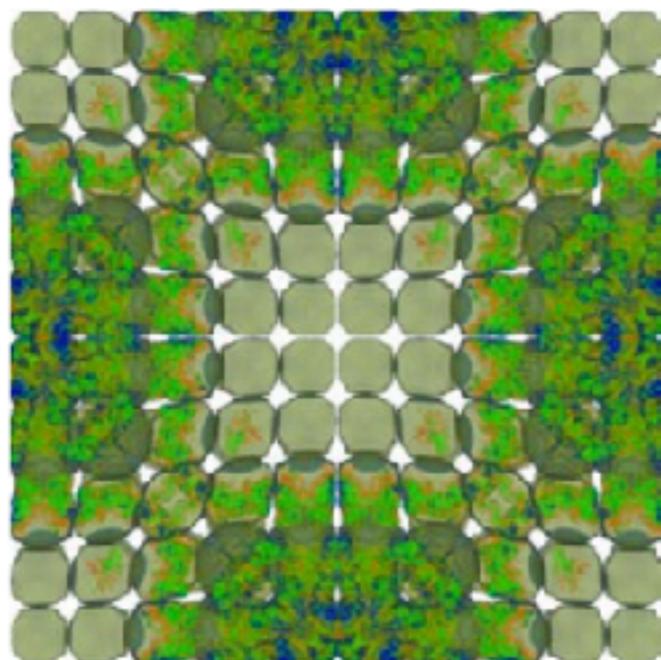
Parallel Algorithms



Parallelizing visualization algorithms can be **very** hard.

Interesting examples:

- Map-Reduce implementation of visualization algorithms (isosurface, volume rendering, ...)



PPM Richtmyer-Meshkov instability (7.6GB)			
#isovalue	CLuE time	Cluster time	File Written
	256M/256R	480M/480R	
1	2min 55s	57s	1.78MB
2	2min 00s	57s	1.81MB
4	9min 35s	2min38s	35.26GB
8	10min 26s	2min37s	103.92GB
16	28min 29s	5min27s	252.38GB

SKULL (256MB)			
#isovalue	CLuE time	Cluster time	File Written
	256M/256R	480M/480R	
1	2min 22s	54s	1.94MB
2	2min 00s	56s	15.83MB
4	2min 02s	56s	31.22MB
8	2min 11s	58s	3.35GB
16	6min 06s	1min27s	18.74GB

[Vo et al., Parallel Visualization on Large Clusters using MapReduce, Proc. LDAV 2011]

Parallel Algorithms



Interesting examples:

- Parallel Contour Tree Computation

[Carr et al., *Hybrid Data-Parallel Contour Tree Computation*, [dx.doi.org/10.2312/cgvc.20161299](https://doi.org/10.2312/cgvc.20161299)]

- Feature-Tracking in Large Simulations

[Widanagamaachi et al., Tracking features in embedded surfaces: Understanding extinction in turbulent combustion, doi.org/10.1109/LDAV.2015.7348066]

- ...

Community Effort: [vtk-m](#)



- “data-parallel” high-performance GPU implementations

Parallel Algorithms



Open-source tools with good general-purpose parallel implementations:

- Paraview

paraview.org



- VisIt

<https://wci.llnl.gov/simulation/computer-codes/visit/>

In Situ Techniques

I/O Bottleneck



The ability to generate data has increased tremendously;
storing simulation output is becoming too expensive

	2010	2018	Factor Change
System Peak	2 Pf/s	1 Ef/s	500
Power	6 MW	20 MW	3
System Memory	0.3 PB	10 PB	33
Node Performance	0.125 Tf/s	10 Tf/s	80
Node Memory BW	25 GB/s	400 GB/s	16
Node Concurrency	12 CPUs	1000 CPUs	83
Interconnect BW	1.5 GB/s	50 GB/s	33
System Size (nodes)	20 K nodes	1 M nodes	50
Total Concurrency	225 K	1 B	4444
Storage	15 PB	300 PB	20
Input/Output Bandwidth	0.2 TB/s	20 TB/s	100

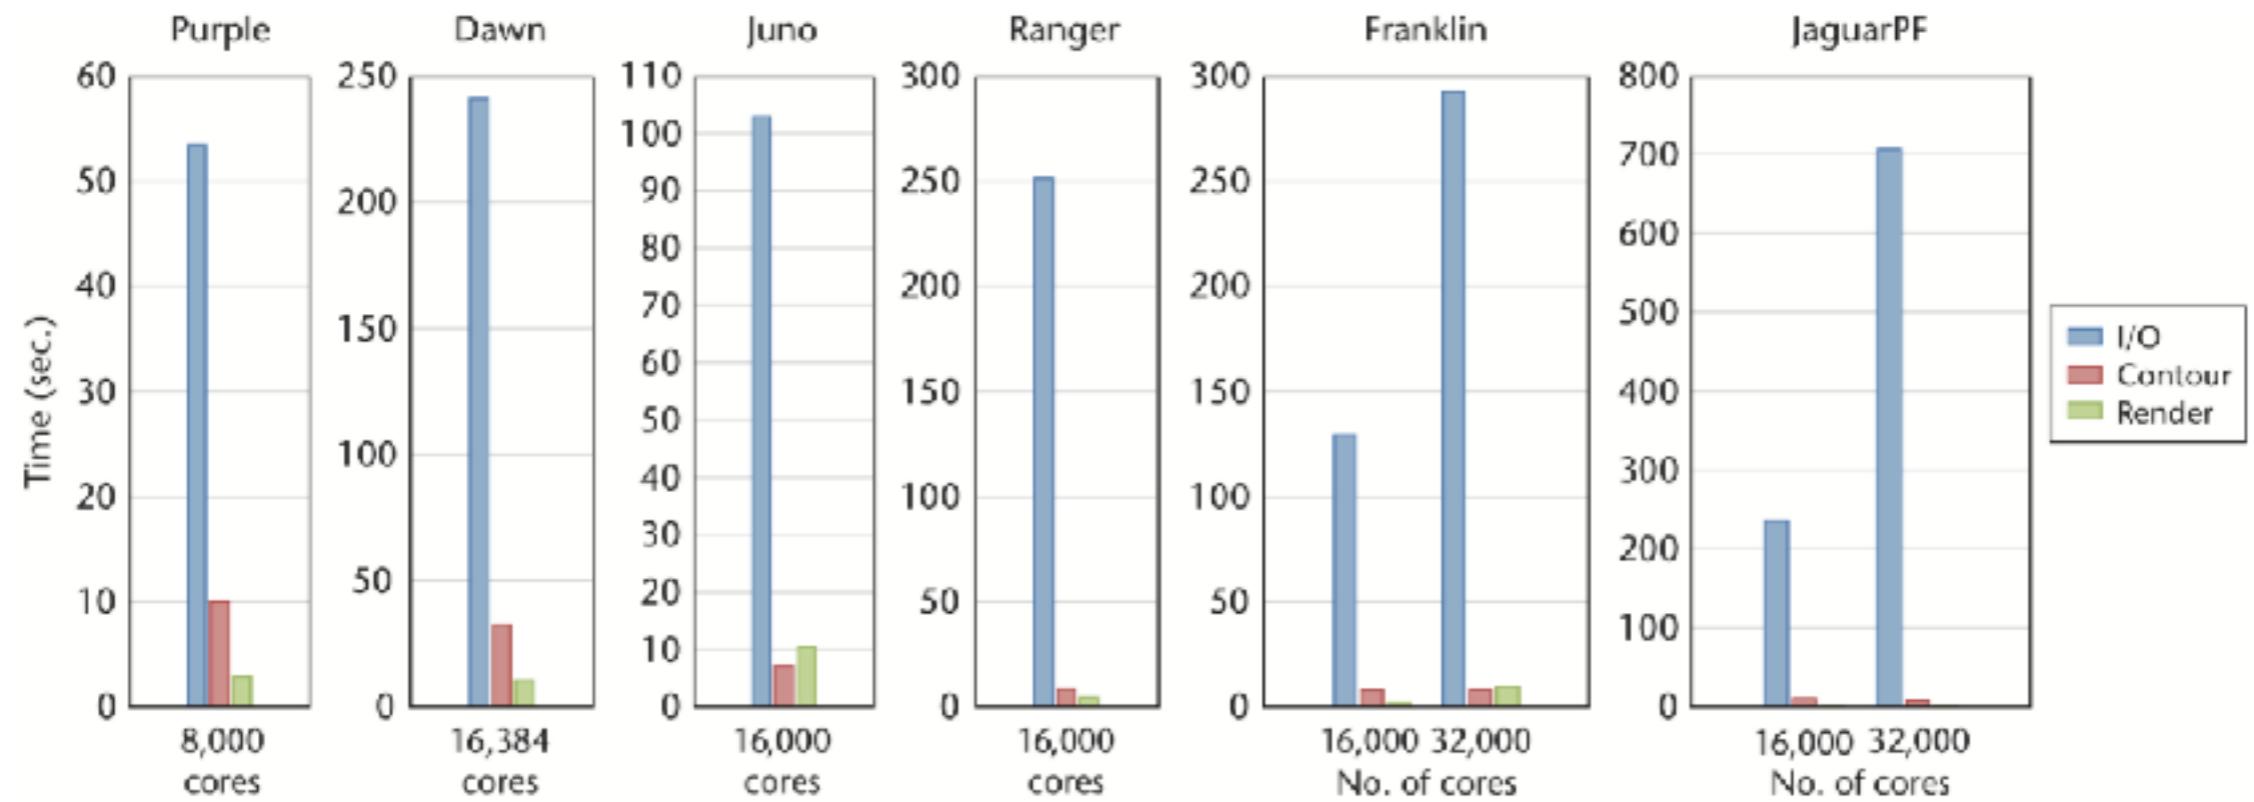
from Vetter et al. "The key to Exascale Computing." HPC workshop, 2008

Visualization as pure **post-processing** is no longer practical.

In Situ Visualization



Re-thinking visualization as **post-processing**: for many practical problems, storing simulation output has become too expensive.



[from Childs et al.: *Extreme scaling of production visualization software on diverse architectures*. IEEE CG&A 30(3):22-31, 2010]

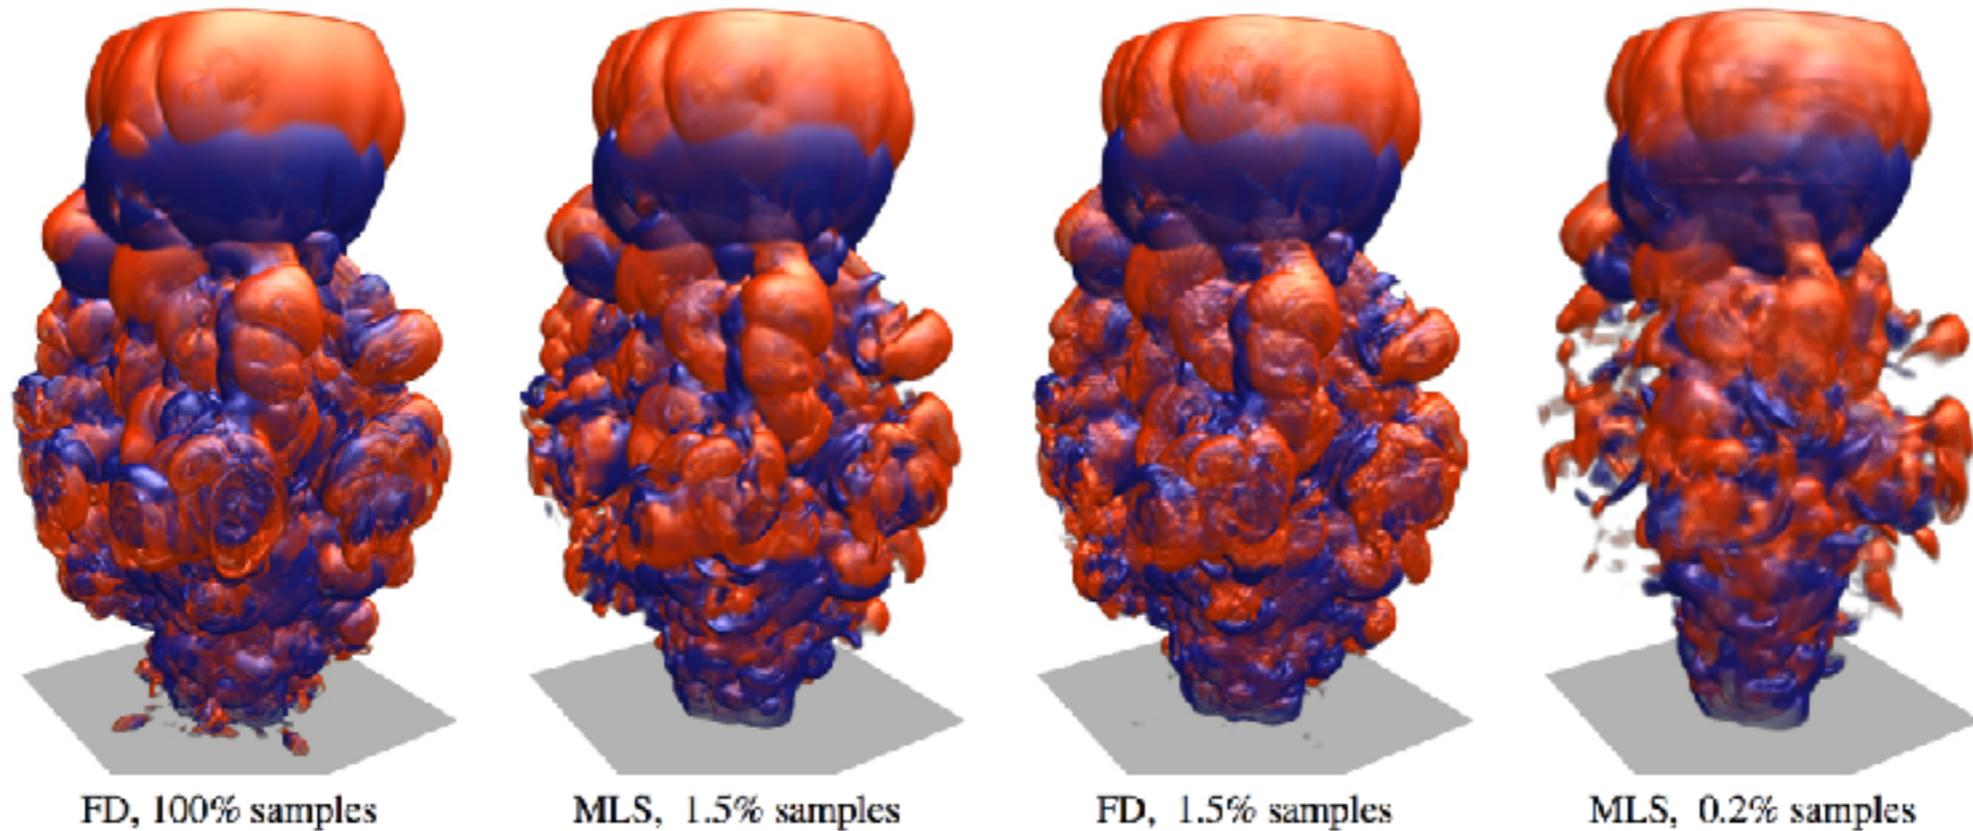
I/O bandwidth does not scale with CPU & memory performance.

In Situ Visualization



This necessitates smart data reduction at simulation time – *in situ*.

Example: representation of flow maps (computed in situ) using a representative set of sparse particles.



A. Agranovsky, CG, K. I. Joy: *Extracting Flow Structures using Sparse Particles*. In Proc. Vision, Modeling, Visualization, 2011.

A. Agranovsky, D. Camp, CG, E. W. Bethel, K. I. Joy, H. Childs: *Improved Post Hoc Flow Analysis Via Lagrangian Representations*. In Proc. Large-Data Analysis and Visualization, 2014 (best paper)

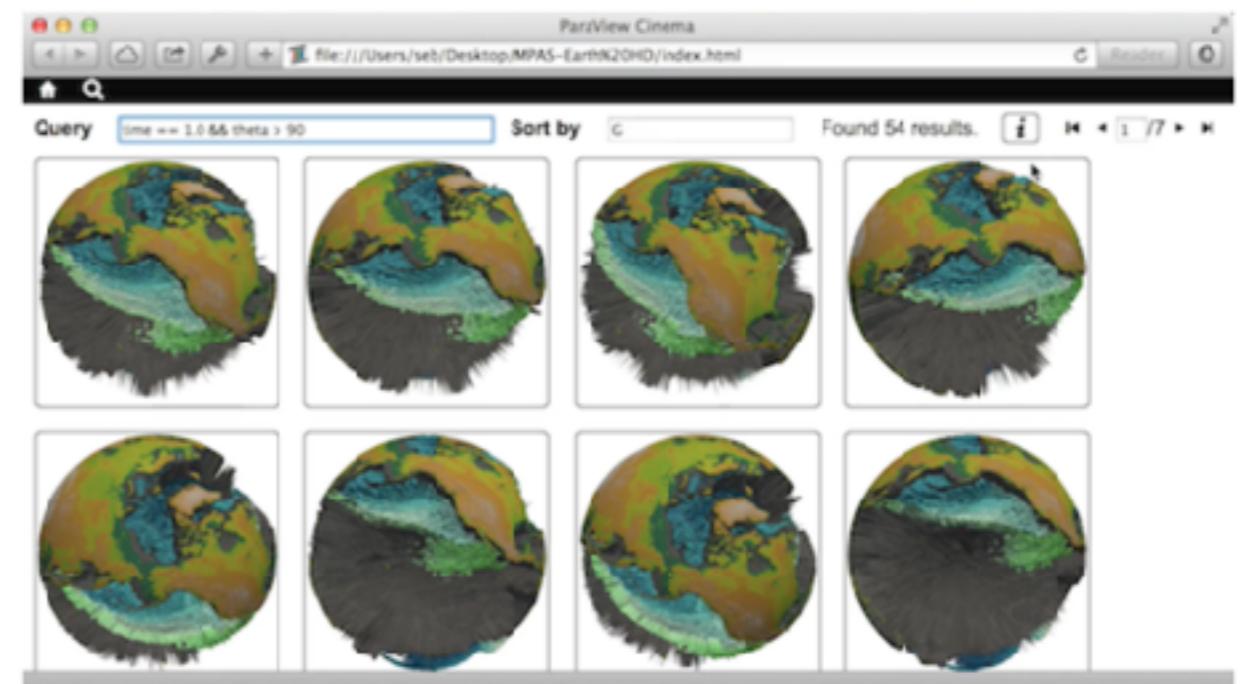
In Situ Visualization



Even more extreme: execute visualization algorithms (e.g. feature extraction) during simulation.

Example: Paraview Cinema
(Ahrens et al. 2014)

- pre-specify all visualization parameters and store only images
- preserve limited flexibility by storing many different combinations of visualizations and viewpoints



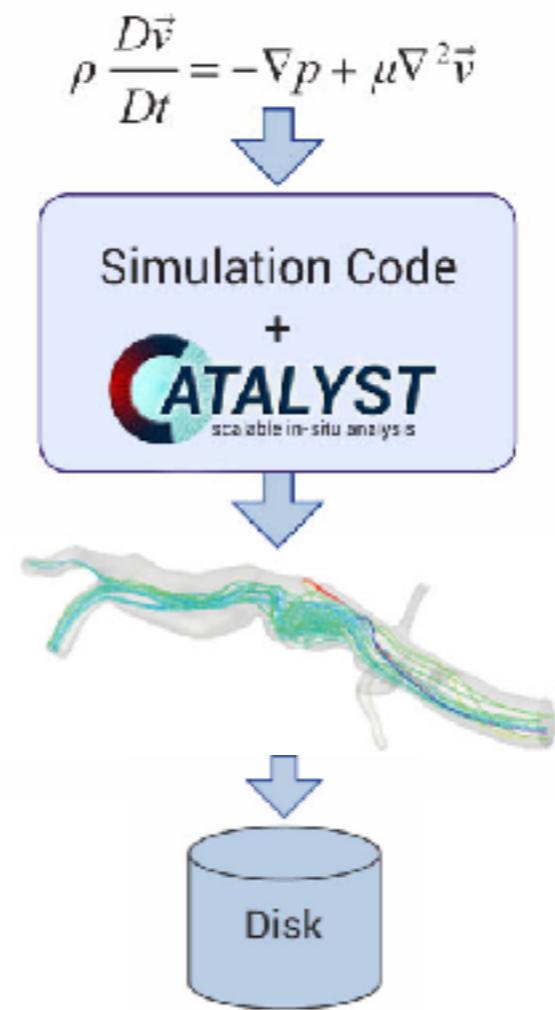
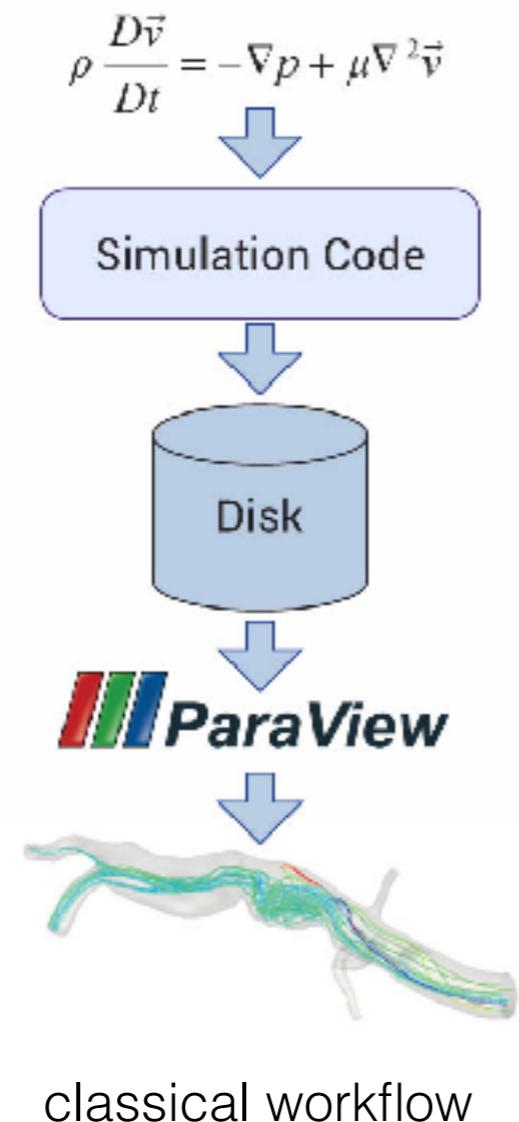
from Ahrens et al.: In Situ MPAS-Ocean Image-based Visualization. In Proc ACM Supercomputing (SC'14).

[Video](#)

In Situ Visualization



In situ visualization: compute visualization while simulating.



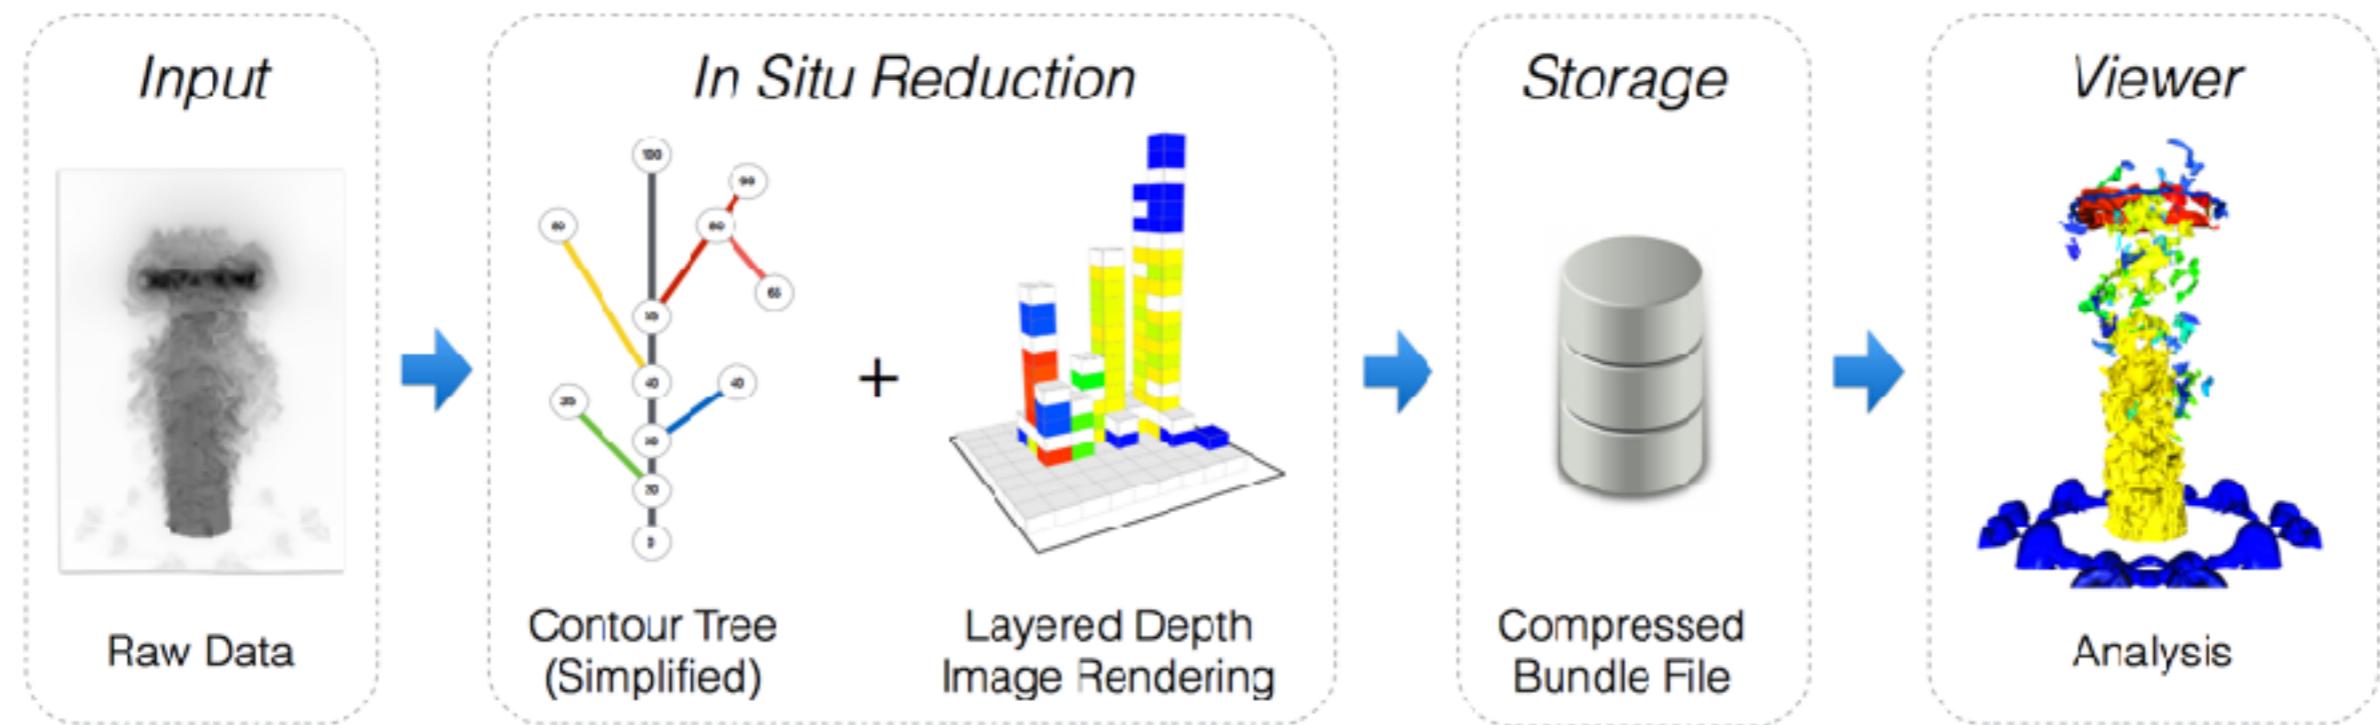
in-situ workflow

Problem: very limited interaction.

In Situ Visualization



Feature-based visualization example: combine **image-based approach** with **topological feature extraction**.



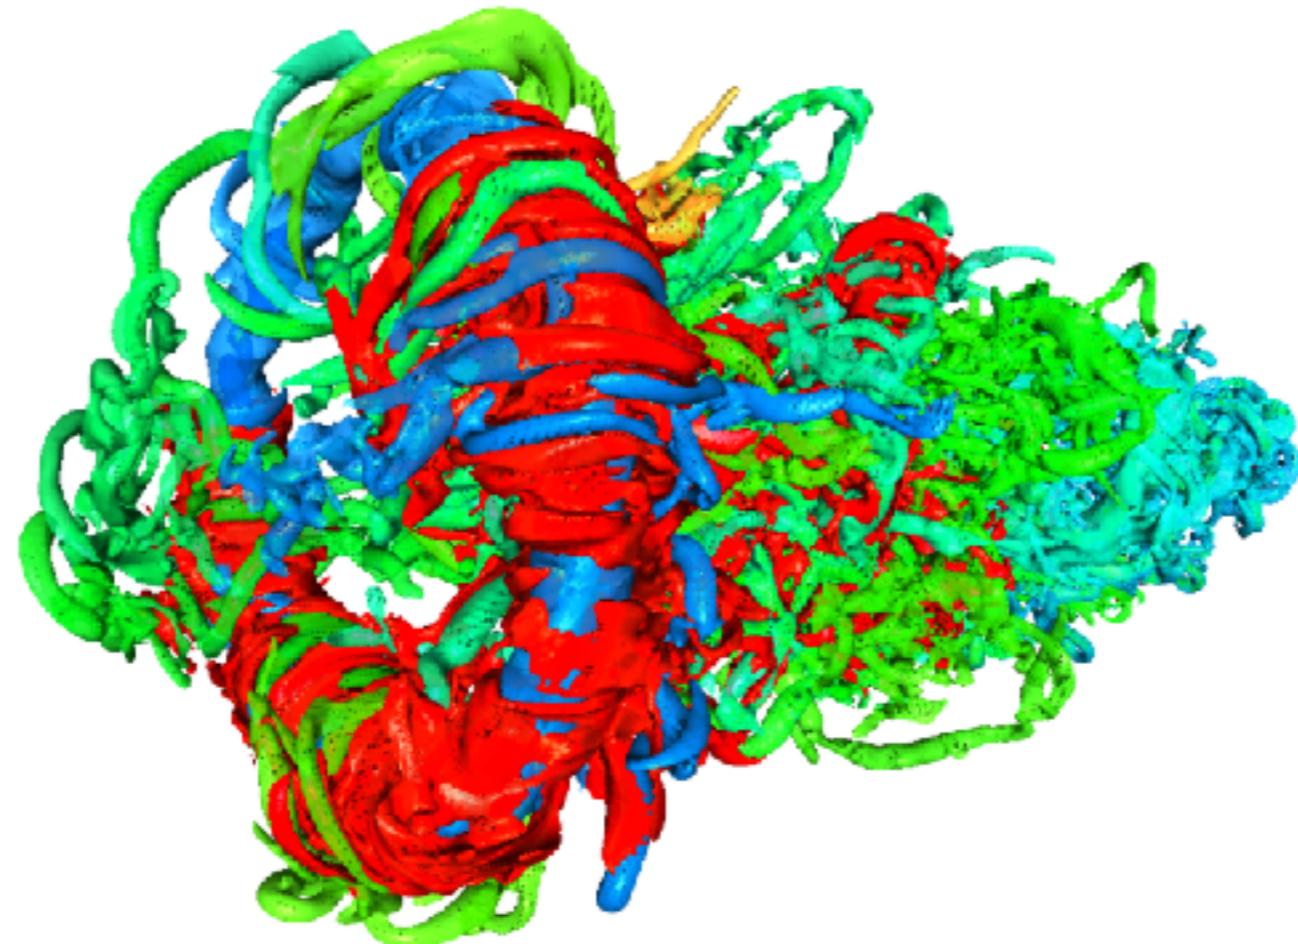
Raw data is converted into a topological representation, simplified, and rendered into a **layered depth image**.

T. Biedert, CG: Contour Tree Depth Images For Large Data Visualization. In Proc. EGPGV, 2015.

Contour Tree Depth Images



The depth image stores only a fixed number of the most important topological features, after simplification.



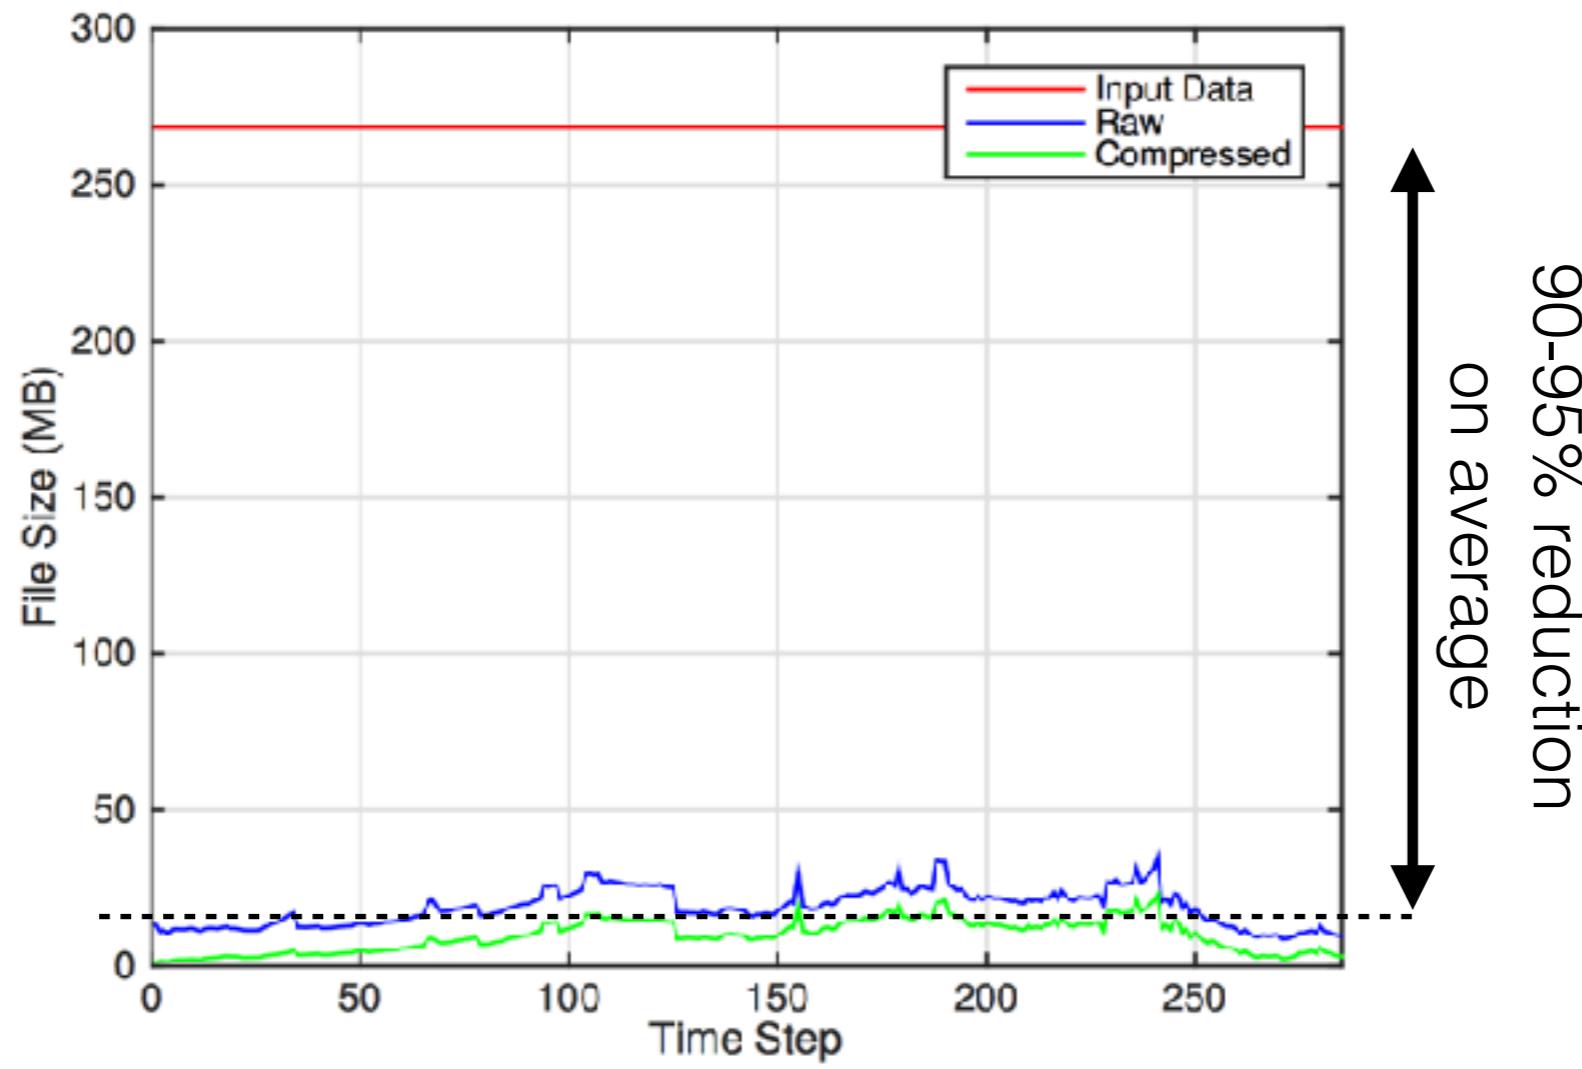
Vortices in a jet flow characterized via λ_2 -criterion
(100 branches stored, **1.5M** branches before simplification)

T. Biedert, CG: Contour Tree Depth Images For Large Data Visualization. In Proc. EGPGV, 2015.

Contour Tree Depth Images



Large storage savings are achievable.



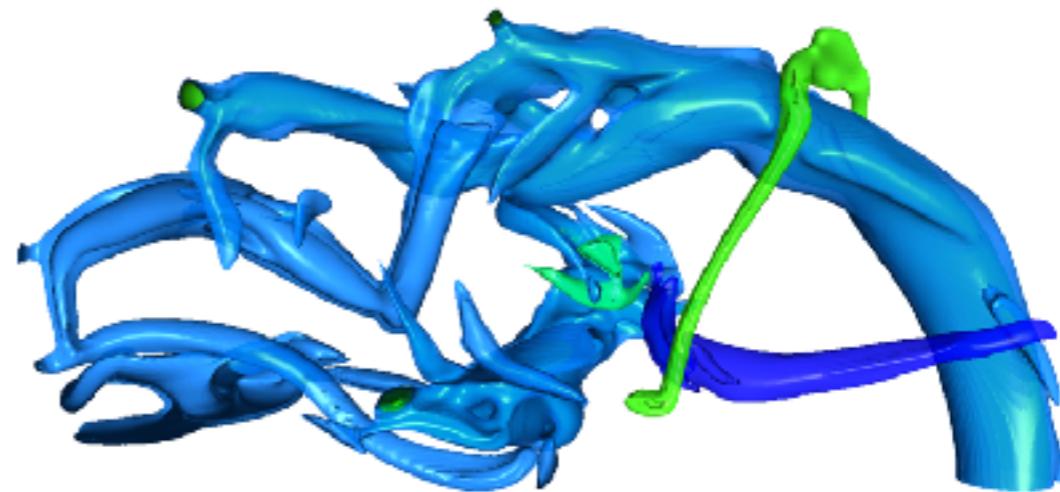
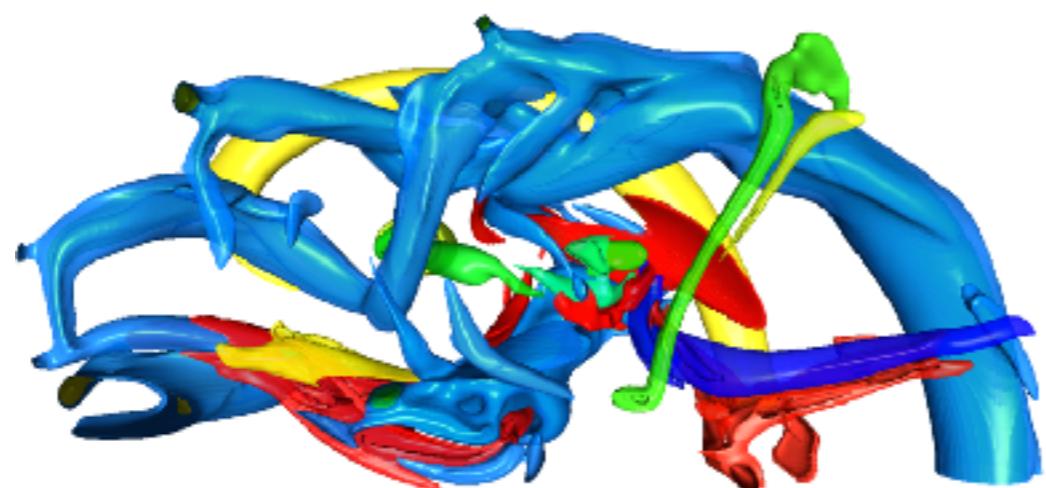
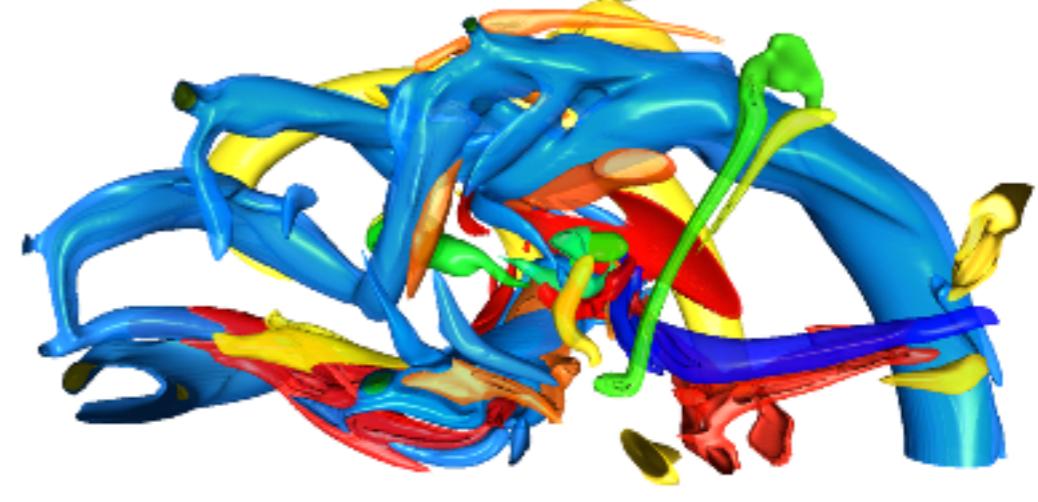
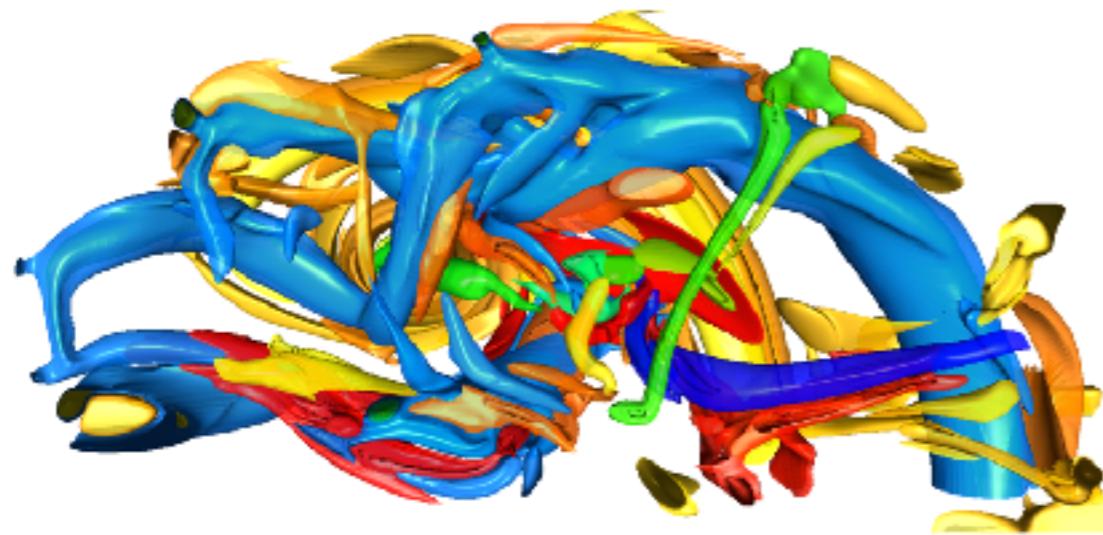
Data can be copied to and visualized on a standard PC.

T. Biedert, CG: Contour Tree Depth Images For Large Data Visualization. In Proc. EGPGV, 2015.

Contour Tree Depth Images



Flexibility for further filtering and simplification is preserved.



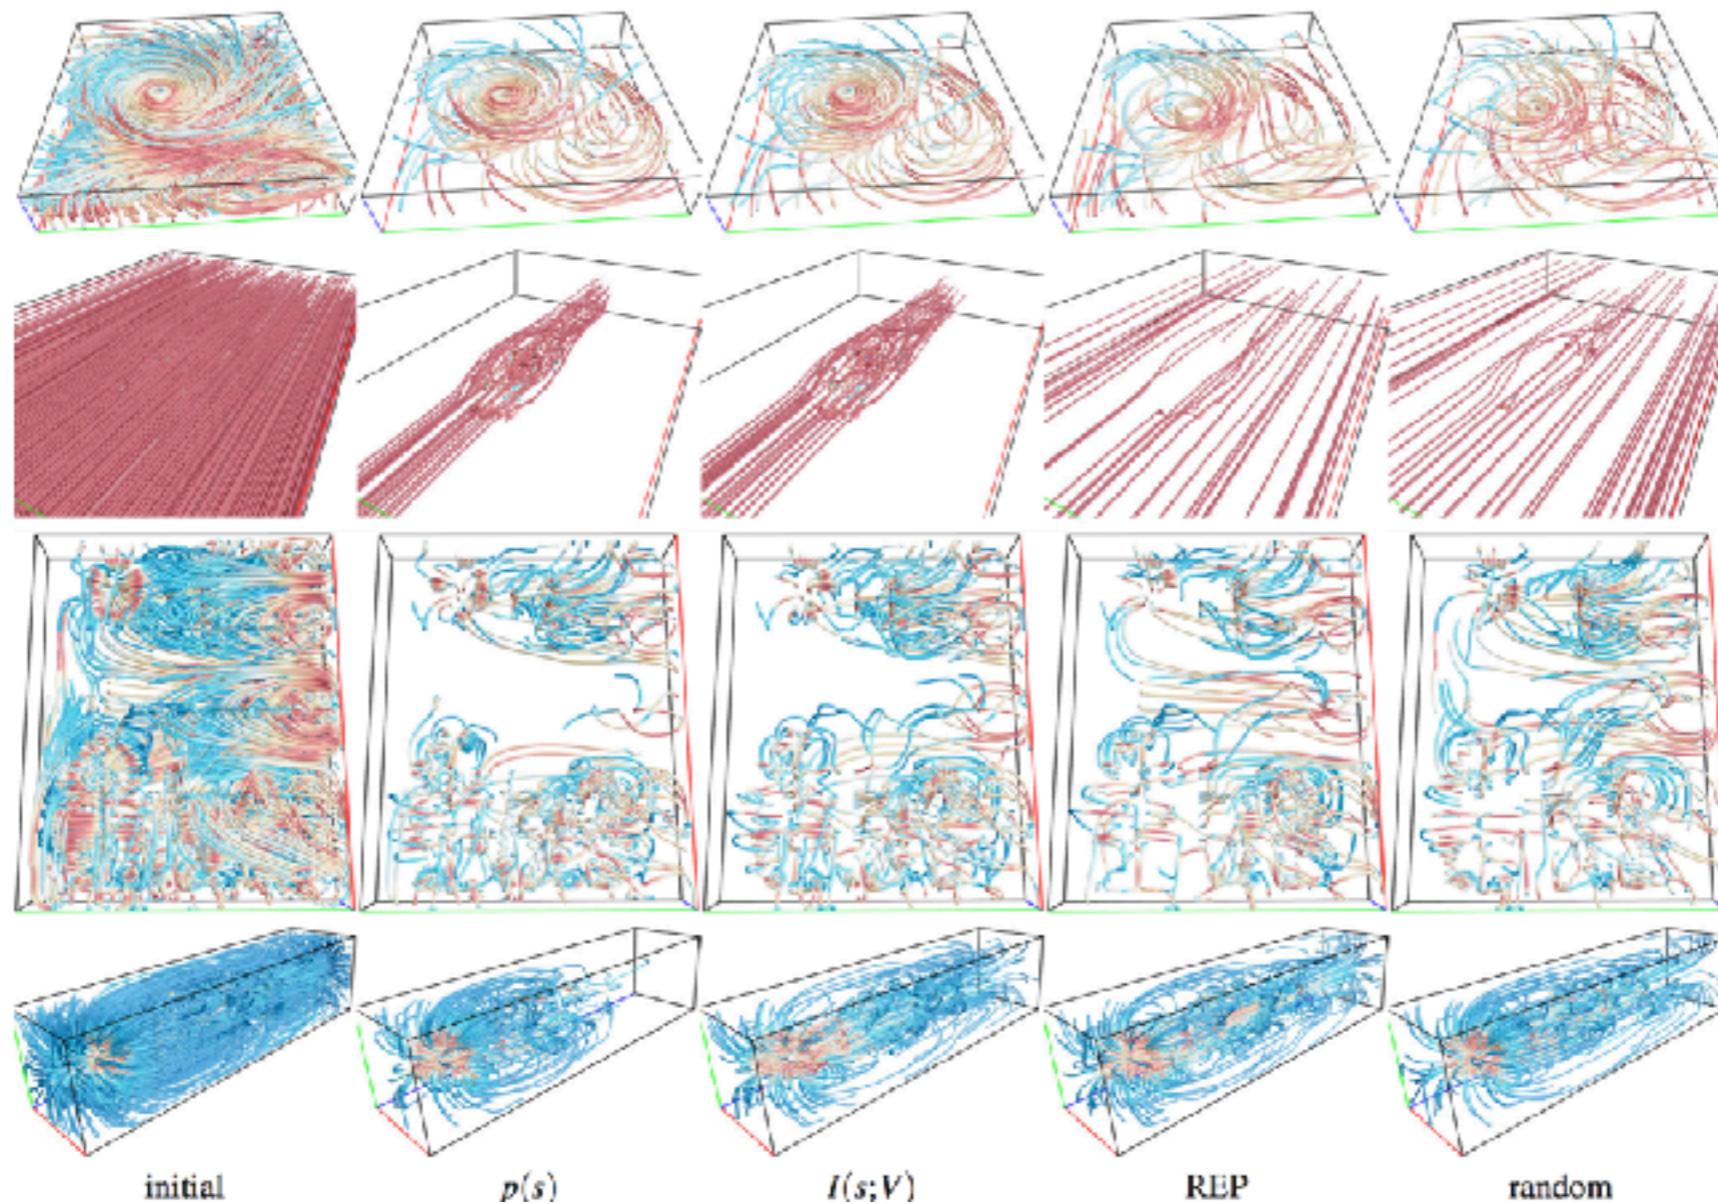
T. Biedert, CG: Contour Tree Depth Images For Large Data Visualization. In Proc. EGPGV, 2015.

In Situ Visualization



Various techniques to generate “good” images.

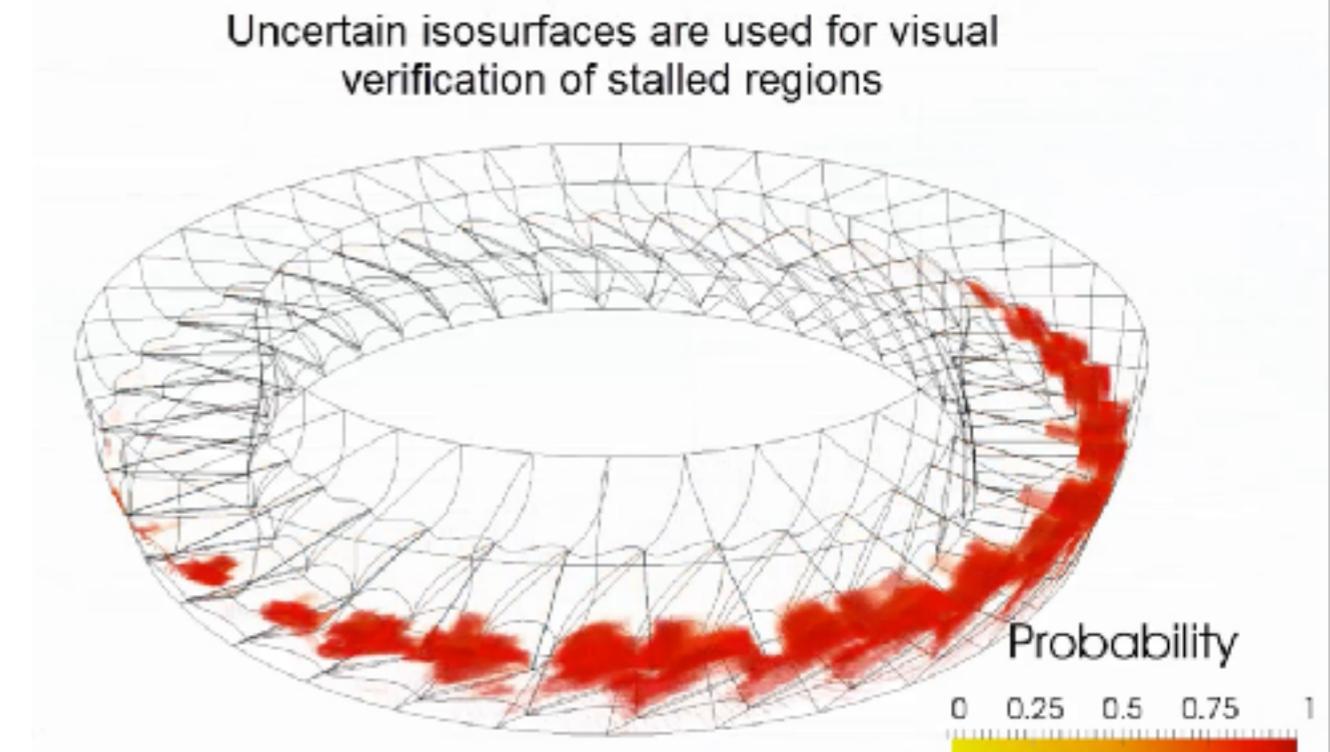
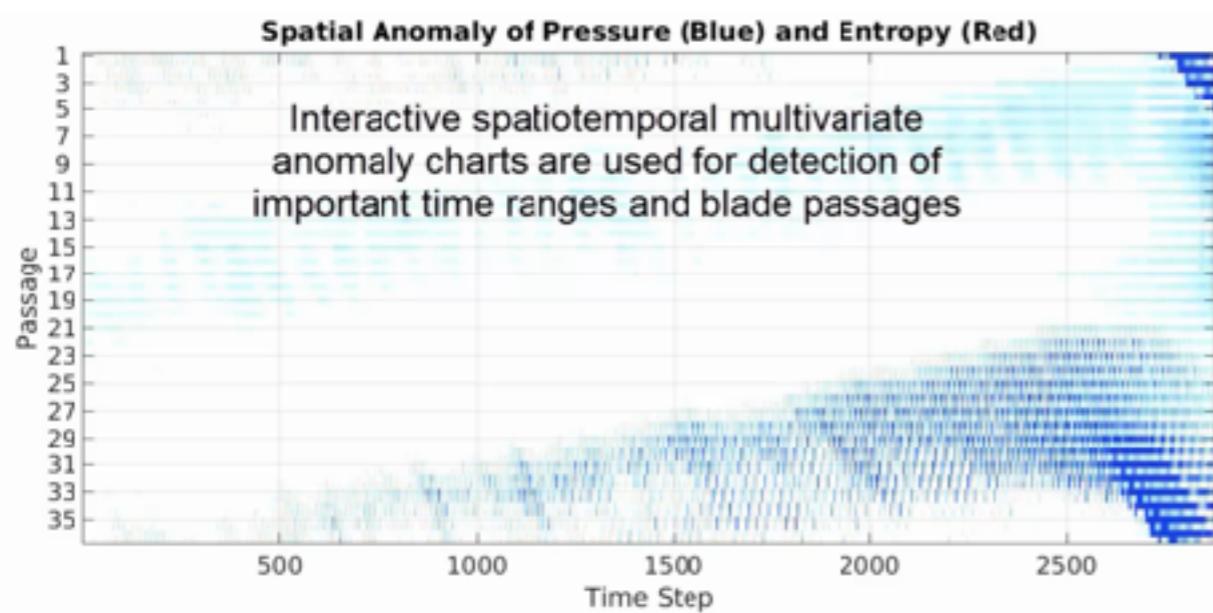
E.g. select “interesting” streamlines using information theory



In Situ Visualization



E.g. using distribution analysis to select interesting data subsets and store only these (data compression).



Dutta et al., In Situ Distribution Guided Analysis and Visualization of Transonic Jet Engine Simulations, IEEE TVCG 23(1):1077-2626, 2016

And many many more techniques for very specific scenarios.

Conclusion

Conclusion



Visualization of modern datasets (in the sense of this talk) is not a straightforward undertaking.

Research is ongoing on which techniques are general.

Axes to consider are:

- Visual scalability
- Cognitive scalability
- Computational scalability

Not addressed today: ensemble analysis, uncertainty visualization, high-dimensional point data, etc.